



PAPER

OPEN ACCESS

RECEIVED

24 October 2025

REVISED

9 January 2026

ACCEPTED FOR PUBLICATION

26 February 2026

PUBLISHED

30 March 2026

Original content from this work may be used under the terms of the [Creative Commons Attribution 4.0 licence](#).

Any further distribution of this work must maintain attribution to the author(s) and the title of the work, journal citation and DOI.



MDCrow: automating molecular dynamics workflows with large language models

Quintina Campbell^{1,4} , Sam Cox^{1,3,4} , Jorge Medina^{1,4} , Brittany Watterson² and Andrew D White^{1,3,*} ¹ Department of Chemical Engineering, University of Rochester, Rochester, NY, United States of America² Department of Biomedical Engineering, University of Rochester, Rochester, NY, United States of America³ FutureHouse Inc., San Francisco, CA, United States of America⁴ These authors contributed equally to this work.

* Author to whom any correspondence should be addressed.

E-mail: andrew.white@rochester.edu**Keywords:** agent, agentic AI, computational biology, molecular dynamics, large language modelsSupplementary material for this article is available [online](#)

Abstract

Molecular dynamics (MD) simulations are essential for understanding biomolecular systems but remain challenging to automate. Recent advances in large language models (LLMs) have demonstrated success in automating complex scientific tasks using LLM-based agents. In this paper, we introduce MDCrow, an agentic LLM assistant capable of automating MD workflows for proteins. MDCrow uses chain-of-thought over 40 expert-designed tools for handling and processing files, setting up simulations, analyzing the simulation outputs, and retrieving relevant information from literature and databases. We assess MDCrow's performance across 25 common tasks of varying complexity, and we evaluate the agent's robustness to difficulty and prompt style. gpt-4o is able to complete increasingly complex tasks with low variance, followed closely by llama3-405b, a compelling open-source model. While prompt style does not influence the best models' performance, it has significant effects on smaller models.

1. Introduction

Molecular dynamics (MD) is a common method to understand dynamic and complex systems in chemistry and biology. While MD is now routine, its integration into and impact on scientific workflows has increased dramatically over the past few decades [1–3]. There are two main reasons for this: First, MD provides valuable insights. Through simulations, scientists can study structural and dynamic phenomena, perturbations, and dynamic processes in their chemical systems. Second, innovations in hardware and expert-designed software packages have made MD much more accessible to both experienced and novice users [3].

For a given protein simulation, parameter selection is nontrivial: the user must provide the input structure (such as a Protein Data Bank (PDB) [4] file), select a force field (e.g. CHARMM [5], AMBER [6]), and specify parameters such as temperature, integrator, simulation length, and equilibration protocols. Simulations also generally require pre- and post-processing steps, along with various analyses. For instance, a user may need to clean or trim a PDB file, add a solvent, or analyze the protein's structure. After simulation, they might examine protein conformation over time or assess its stability under different conditions. The choices for pre-processing, analysis, and simulation parameters are highly specific to any given use case and often require expert intuition. In practice, even with careful parameter selection, one often needs to run simulations, evaluate results, and adjust settings iteratively. Thus, automating this process is difficult but beneficial.

Several efforts have been made to automate MD workflows, including Admiral [7], CHAPERONg [8], FabSim [9], the Simulation Foundry [10], Gmx_qk [11], MolAr [12], ProFESSA [13], ProtoCaller

[14], SimStack [15], with most targeting specific application domains, such as RadonPy for polymer's simulations [16], or PyAutoFEP for proteins and small molecules for drug-screening [17]. Other approaches are constrained to a particular combination of simulation software and simulation (e.g. GROMACS and Free Energy Perturbation). Certainly, there has been significant community-driven improvement in automating and creating MD toolkits (EasyAmber [18], PACKMOL [19], MDAnalysis [20], MDTraj [21], OpenMM [22], GROMACS [23], LAMMPS [24]) and user-friendly interfaces and visualizations (MDANSE [25], QwikMD [26], MDWiZ [27], METAGUI [28], VMD [29], Gromita [30], PlayMolecule [31]). While these advances improve the capabilities and ease of use, the inherent variability of MD workflows still poses a great challenge for full automation.

Large language model (LLM) agent frameworks, such as Toolformer [32], MRKL [33], ReAct [34], Aviary [35], have gained popularity for their ability to automate technical tasks through scientific reasoning and tool usage. LLM agents even surpass domain-specialized LLMs (e.g. BioGPT [36], Med-PaLM [37]) when programmed for specialized roles [38]. These agents have demonstrated promising results in scientific tasks within a predefined toolspace like ChemCrow [39], Coscientist [40], and CACTUS [41], successfully automating complex workflows and novel design in chemical synthesis. Similarly, LLM-driven automation has been explored in material science and chemistry research [42], including catalyst design [43], inorganic synthesis [44], scientific literature review [45], and data aggregation [46]. LLM-driven automation is also explored in more sophisticated tasks that involve multi-step or multi-agent systems: context-aware experiment planning [47], evaluation of protocol planning [48], gene editing experiment design [49], reaction development framework [50], metal organic framework generation [51], the scientific discovery pipeline with a team of AI scientists [52]. Most similar to this work, ProtAgents [53] is a multi-agent modeling framework tackling protein-related design and analysis, and LLaMP [54] applies a retrieval-augmented generation-based ReAct agent to simulate inorganic materials by interfacing with literature databases, Wikipedia, and atomistic simulation tools. Despite these advances, no fully adaptive and autonomous systems exist for biochemical MD or protein simulations. See Ramos *et al* [55] for an informative review on the design, assessment, and applications of scientific agents.

Here we present MDCrow, an LLM-agent capable of autonomously completing MD workflows (figure 1). We measure MDCrow's performance across 25 common tasks with varying difficulty and compare performance of different LLM models. We then analyze the effect of prompting style on performance and assess performance against task complexity, based on required number of subtasks; finally, we compare MDCrow with an LLM equipped with a python interpreter with the required packages installed, rather than using a custom built environment. We find that MDCrow with gpt-4o or llama3-405b is able to perform nearly all of our assessed tasks and is relatively insensitive to how precise the instructions are given to it. See figure 1(D) for an overview of the main results.

Following this work, other recent LLM-based MD automation frameworks include NAMD-Agent [56], which prioritizes fast setup protein systems via web automation CHARMM-GUI, and DynaMate [57], which extends LLM automation to protein-ligand complexes through a multi-agent framework utilizing AmberTools and GROMACS. MDCrow focuses on protein simulations and uniquely contributes in the following ways: systematic benchmarking across open-source and frontier LLMs, assessing prompt-style robustness and capability of automating multi-step non-linear tasks of varying complexity, and checkpoint-based chat with memory for pausing and resuming across sessions.

2. Methods

2.1. MDCrow toolset

MDCrow consists of an environment with tools that emit observations and an LLM that selects appropriate actions (tools + input arguments). MDCrow is built with Langchain [58] and a ReAct style prompt [34]. The tools mostly consist of analysis and simulation methods; we use OpenMM [22] and MDTraj [21] packages, but in principle, our findings generalize to any such packages. See SI section 6 for an example of MDCrow adapting to work with GROMACS. The LLM agent operates at LLM temperature of 0.1, which is relatively low stochasticity and reflects realistic LLM agent behavior if one chooses to use LLMs to automate MD workflows.

MDCrow's tools can be categorized in four groups: information retrieval, PDB & protein, simulation, and analysis (see figure 1(B)).

Information retrieval tools These tools enable MDCrow to build context and answer simple questions posed by the user. Most of the tools serve as wrappers for UniProt API functionalities [59], allowing access to data such as 3D structures, binding sites, and kinetic properties of proteins. Additionally, we include a LiteratureSearch tool, which uses PaperQA [45] to answer questions and retrieve information




Figure 1. (A) MDCrow workflow. Starting with a user prompt and initialized with a set of MD tools, MDCrow follows a chain-of-thought process until it completes all tasks in the prompt. The final output includes a response, along with all resulting analyses and files. (B) The tool distribution categorized into four types: information retrieval, PDB and protein handling, simulation, and analysis. A few examples from each category are shown. (C) Two example prompts that MDCrow is tested on. The first is the simplest prompt, containing only 1 subtask. The most complex task requires 10 subtasks. (D) Average subtask completion across all 25 prompts as task complexity (the number of subtasks per prompt) increases. The top three performing base-LLMs are shown. Among them, **gpt-4o** and **llama3-405b** are consistently stable, staying close to 100% completion even as task complexity increases.

from literature. PaperQA accesses a local database of relevant PDFs, selected specifically for the test prompts, which can be found in SI section 3. This real-time information helps the system provide direct answers to user questions and can also assist the agent in selecting parameters or guiding simulation processes.

PDB & protein tools MDCrow uses these tools to interact directly with PDB files, performing tasks such as cleaning structures with PDBFixer [22], retrieving PDBs for small molecules and proteins, and generating visualizations of PDBs for the user through Molrender [60] or NGLview [61].

Simulation tools The simulation tools use OpenMM [22] and Packmol [19] for preprocessing and simulation. These tools are built to manage dynamic simulation parameters, handle errors related to inadequate parameters or incomplete preprocessing, and address missing forcefield templates efficiently. The agent responds to simulation setup errors through informative error messages, improving overall robustness. Finally, the simulation tools outputs Python scripts that can be modified directly by MDCrow whenever the simulation requires additional steps or parameters.

Analysis tools This group of tools is the largest in the toolset, designed to cover common MD analysis methods, many of which are built on MDTraj [21] functionalities. Examples include computing the root mean squared distance (RMSD) with respect to a reference structure, calculating the radius of gyration, analyzing the secondary structure, and generating plots.

2.2. Chatting

A key challenge in developing an automated MD assistant and supporting user interactions with simulations is ensuring it can manage a large number of files, analyses, and long simulations and runtimes. Although MDCrow has been primarily tested with shorter simulations, it is designed to handle larger workflows as well. Its ability to retrieve and resume previous runs allows users to start a simulation, step away during the long process, and later continue interactions and analyses without needing to stay engaged the entire time. When returning, users can easily build on prior sessions—for example, asking for additional plots or modified conditions (e.g. running again at a lower temperature). An example of this chatting feature is shown in figure 2.

MDCrow creates an LLM-generated summary of the user prompt and agent trace, which is assigned to a unique run identifier provided at the end of the run (but accessible at any time during the session). Each run's files, figures, and path registry are saved in a unique checkpoint folder linked to the run identifier.

When resuming a chat, the LLM loads the summarized context of previous steps and maintains access to the same file corpus, as long as the created files remain accessible. To resume a run, the user simply provides the checkpoint directory and run identifier. MDCrow then loads the corresponding memory summaries and retrieves all associated files, enabling seamless continuation of analyses.

3. Results

3.1. MDCrow performance on various tasks


To assess MDCrow's ability to complete tasks of varying difficulty, we designed 25 prompts with different levels of complexity and documented the number of subtasks (minimum required steps) needed to complete each task. MDCrow was not penalized for taking additional steps, but was penalized for omitting necessary ones. For example, the first prompt in figure 1(C) contains a single subtask, whereas the complex task requires 10 subtasks: downloading the PDB file, performing three simulations, and performing two analyses per simulation. If the agent failed to complete an earlier step, it was penalized for every subsequent step it could not perform due to that failure.

The 25 prompts require between 1 and 10 subtasks, with their distribution shown in figure 3(B). Each prompt was tested across three GPT models (gpt-3.5-turbo-0125, gpt-4-turbo-2024-04-09, gpt-4o-2024-08-06) [62, 63], two Llama models (llama-v3p1-405b-instruct, llama-v3p1-70b-instruct) [64] (accessed via the Fireworks AI API with 8-bit floating point (8FP) quantization [65]), and two Claude models (claude-3-opus-2024-02-29, claude-3-5-sonnet-2024-06-20) [66, 67]. A newer Claude Sonnet model, claude-3-5-sonnet-2024-10-22 was tested in later experiments but was not found to give superior results, so it was not tested on these 25 prompts. All other parameters were held constant across tests, and each version of MDCrow executed a single run per prompt.

Each run was evaluated by experts recording the number of required subtasks the agent completed and using Boolean indicators to indicate accuracy, whether the agent triggered a runtime error, and whether the trajectory contained any hallucinations. Since the agent trajectories for each run are inherently variable, accuracy is defined as the result's consistency with the expected trajectory rather than comparing against a fixed reference.

The percentage of tasks that were deemed to have valid solutions for MDCrow with each base-LLM is shown in figure 3(A). The lowest performing model was gpt-3.5. This is not surprising, as this model had some of the highest hallucination rates (32% of prompt completions contained hallucinations), compared to the absence of documented hallucinations in the higher performing models, gpt-4o and llama3-405b. However, the discrepancy in accuracy rates between models cannot solely be attributed to hallucinations, as gpt-3.5 attempted fewer than half of the required subtasks, whereas the higher-performing models, gpt-4o and llama3-405b, attempted 80%–90% of the required subtask, earning accuracy in answering for 72% and 68% of tasks, respectively (figures 3(C) and (D)).

These results indicate that MDCrow can handle complex MD tasks but is limited by the capabilities of the base model. For gpt-4-turbo, gpt-3.5, and llama3-70b, the number of trajectories with verified results decreases significantly as task complexity increases (figure 3(C)). In contrast, gpt-4o and



11ama3-405b show only a slight decline, demonstrating that MDCrow performs well even for complex tasks when paired with more robust base models.

3.2. MDCrow robustness

We evaluated MDCrow's robustness on complex prompts and different prompt styles. We hypothesized that some models would excel at completing complex tasks, while others would struggle—either forgetting steps or hallucinating—as the number of required subtasks increased. To test this, we created a sequence of 10 prompts that increased in complexity. The first prompt required a single subtask, and each subsequent prompt added an additional subtask (see figure 4(A)). Each prompt was tested twice: once in a natural, conversational style and once with explicitly ordered steps. Example prompts can be seen in figure 4(B).

To quantify robustness, we calculated the coefficient of variation (CV) for the percentage of completed subtasks across tasks. A lower CV indicates greater consistency in task completion and, therefore, higher robustness. The analysis revealed clear differences in robustness across models and prompt types. Overall, gpt-4o and 11ama3-405b demonstrated moderate to high robustness, while the Claude models showed significantly lower robustness. The performance comparison is shown in figure 4(C).




Figure 3. MDCrow performance across large language models. (A) Summary of MDCrow performance dependent on LLM. Percentage of accuracy is determined by whether it gave acceptable final answer or not. While statistically indistinguishable from Claude and Llama models, **gpt-4o** significantly outperforms the rest of GPT models on giving accurate solutions (t -test, $0.004 \leq p\text{-value} \leq 0.046$). (B) The distribution of number of subtasks in each task across 25 prompts. The prompts range from 1 to 10 steps, with each step count belonging to at least 2 prompts. (C) Percentages of prompts with accurate solutions with respect to LLM used and number of subtasks per task. The correlation between accuracy and complexity is statistically significant for all LLMs (Spearman correlation, $3.9 \times 10^{-7} \leq p\text{-value} \leq 1.1 \times 10^{-2}$). (D) Percentage of the subtasks that the agent completed for each base LLM per task.

We expected that the percentage of subtasks completed by each model would decrease as task complexity increased. However, with **gpt-4o** and **llama3-405b** as base models, MDCrow demonstrated a strong relationship between the number of required and completed subtasks (figure 4(D)) for both prompt types, indicating consistent performance regardless of task complexity or prompt style. The three included Claude models demonstrated less impressive performance. **claude-3-opus** followed the linear trend very loosely, becoming more erratic as task complexity increased. As the tasks required more subtasks, the model consistently misses nuances in the instructions and makes logical errors. Both **claude-3.5-sonnet** models gave poor performance on these tasks, often producing the same error (see SI section 1).

3.3. MDCrow comparison

We also compared MDCrow to two baselines: a ReAct [34] agent with only a Python REPL tool and a zero-shot LLM. MDCrow and the baselines were tested on the same 25 prompts as previously mentioned, all using **gpt-4o**. We use different system prompts to accommodate each framework, guiding the LLM to utilize common packages with MDCrow, and these prompts can be found in SI section 2.

The single-query LLM is asked to complete the prompt by writing the code for all subtasks, not unlike what standalone ChatGPT would be asked to do. We then execute the code ourselves and evaluate the outcomes accordingly. ReAct with Python REPL can write, execute, and react to code using a chain-of-thought framework. We find that MDCrow outperforms the two baselines significantly, as shown in figure 5(A), in both subtask completion and accuracy. Both baseline methods struggled with code syntax errors and incorrect handling of PDB files, highlighting how human guidance and integrated tools improve reliability. There is not a significant difference between the two baselines, indicating that the ability to react to failures did not significantly boost performance of these coding-systems.

In figure 5(B), we observe that the performance of all three methods generally declines as task complexity increases. In fact, both baseline methods drop to zero after just three steps, with performance fluctuating erratically at higher complexities. This reflects the critical importance of proper file processing and simulation setup for achieving optimal LLM performance in MD tasks. In contrast, MDCrow demonstrates greater robustness and reliability in handling complex tasks, with its pre-defined toolset allowing accurate file processing and simulation setup, as well as its ability to dynamically adjust to errors.




Figure 4. (A) The number of subtasks in each task, categorized by type. Task 1 begins with a single pre-simulation subtask (*Download a PDB file*) and each subsequent task adds a single subtask, adding to a total of 10 tasks with a maximum of 10 subtasks. (B) Example of ‘Natural’ and ‘Ordered’ prompt style on a three-step prompt. (C) The robustness of MDCrow built on each model with both prompt types, measured by coefficient of variation (CV). Lower CV is interpreted as greater consistency. gpt-4o and llama3-405b are the more robust models, as the Claude models have higher CVs. (D) Comparison of subtask completion across models and prompt types. In the 9-subtask prompt, gpt-4o encountered an error after only one step and gave up without trying to fix it. In general, gpt-4o and llama3-405b have relatively robust performance with increasing complexity for both prompt types. claude-3-opus struggles with more complex tasks, making more logical errors for complex tasks. The two claude-3.5-sonnet models showed fairly poor performance across this experiment.




Figure 5. Performance across LLM Frameworks using the same 25-prompt set: MDCrow, direct LLM with no tools (single-query), and ReAct agent with only Python REPL tool. All use gpt-4o. (A) Performance among LLM frameworks measured by whether accuracy and average percentage of subtasks they complete for each of 25 task prompts. MDCrow is significantly better at giving accurate solutions than direct LLM (t -test, $p = 1 \times 10^{-3}$) and ReAct (t -test, $p = 4 \times 10^{-4}$). MDCrow completes significantly more subtasks on average compared to direct LLM (t -test, $p = 1 \times 10^{-6}$) and ReAct (t -test, $p = 6 \times 10^{-6}$). (B) Percentage of tasks completed with the respect to LLM framework used and the number of subtasks required for each task. The correlation between accuracy and number of subtasks required is statistically significant, $p = 1 \times 10^{-3}$ for direct LLM and $p = 1 \times 10^{-4}$ MDCrow. The p value for ReAct is $p = 7 \times 10^{-2}$.

4. Discussion

Although LLMs’ scientific abilities are growing [68–70], they cannot yet independently complete MD workflows, even with a ReAct framework and Python interpreter. However, with frontier LLMs, chain-of-thought, and an expert-curated toolset, MDCrow successfully handles a broad range of MD tasks. It performs 80% better than gpt-4o in ReAct workflows at completing subtasks, which is expected given the technical demands of MD workflows, such as managing complex files and recovering from simulation errors.

In some cases, particularly for complex tasks beyond its explicit toolset, MDCrow's performance may improve with human guidance. The system's chatting feature allows users to continue previous conversations, clarify misunderstandings, and guide MDCrow step-by-step through difficult tasks. This adaptability helps MDCrow recover from failures, refine its approach based on user intent, and handle more complex workflows. Thus, with more advanced LLM models, targeted feedback, and the addition of specialized tools, MDCrow could tackle an even broader range of tasks. Examples include annealing simulations, specifying appropriate box sizes, and using GROMACS instead of OpenMM, all of which can be found in SI. We did not do a full evaluation of MDCrow's capabilities through this chatting feature in this work.

For all LLMs, task accuracy and subtask completion are affected by task complexity. Interestingly, while gpt-4o can handle multiple steps with the lowest variance, llama3-405b is a compelling second best, as an open-source model. Other models, such as gpt-3.5 and claude-3.5-sonnet, struggle with hallucinations or inability to follow multistep instructions. Performance on these models, however, is improved with explicit prompting or model-specific optimization (especially for claude-3.5-sonnet).

These tasks were focused on routine applications of MD with short simulation runtimes, limited to proteins, common solvents, and force fields included in the OpenMM package. We did not explore small-molecule force fields, especially related to ligand binding. We restricted this study to relatively straight-forward packages and examples, and in the future, MDCrow could be expanded to address more niche applications. Future work could also explore multi-modal approaches [71, 72] for tasks like convergence analysis or plot interpretations. The current framework relies on human-created tools and evaluation, but as LLM-agent systems become more autonomous [73], careful evaluation and benchmarking will be essential.

5. Conclusion

Running and analyzing MD simulations is non-trivial and hard to automate. Here, we explored using LLM agents to accomplish this. We built MDCrow, an LLM-agent and an environment consisting of over 40 common tools built for MD simulation and analysis. We found MDCrow could complete 72% of the tasks with the optimal settings (gpt-4o). llama-405B was able to complete 68%, providing a compelling open-source model. The best models were relatively robust to how the instructions are given, although weaker models struggle with unstructured instructions. Simply using an LLM with a python interpreter and required packages installed had a 28% accuracy. The performance of MDCrow was relatively stable as well, though dependent on the model. Correct assessment of these complex scientific workflows is challenging, and had to be done by hand. Chatting with the simulations, via extended conversations, is even more compelling, but is harder to assess.

Overall, this work demonstrates that LLM agents can meaningfully automate components of MD workflows. With careful setup, open-source models perform comparably to proprietary ones, indicating that capable scientific automation is possible with accessible models. As LLM capabilities advance and benchmarks are developed, systems like MDCrow could evolve to even more complex and autonomous tasks.

Acknowledgments

Research reported in this work was supported by the National Institute of General Medical Sciences of the National Institutes of Health under Award Number R35GM137966, National Science Foundation under Grant Number of 1751471, LLNL FLASK Project, Robert L and Mary L Sproull Fellowship gift and U.S. Department of Energy, Grant No. DE-SC0023354.

Work at FutureHouse is supported by the generosity of Eric and Wendy Schmidt. We thank the Center for Integrated Research Computing (CIRC) at University of Rochester for providing computational resources and technical support.

Data availability statement

The data that support the findings of this study are openly available at the following URL/DOI: <https://doi.org/10.5281/zenodo.18193478> [74].

Supplementary data 1 available at <https://doi.org/10.1088/2632-2153/ae4b07/data1>.

ORCID iDs

Quintina Campbell  0009-0006-8712-6416

Sam Cox  0000-0002-4441-9327

Jorge Medina  0000-0003-2174-5557

Brittany Watterson  0000-0002-6873-2553

Andrew D White  0000-0002-6647-3965

References

- [1] Sinha S, Tam B and Wang S M 2022 Applications of molecular dynamics simulation in protein study *Membranes* **12** 844
- [2] Karplus M and McCammon J A 2002 Molecular dynamics simulations of biomolecules *Nat. Struct. Biol.* **9** 646–52
- [3] Hollingsworth S A and Dror R O 2018 Molecular dynamics simulation for all *Neuron* **99** 1129–43
- [4] Velankar S, Burley S K, Kurisu G, Hoch J C and Markley J L 2021 The protein data bank archive *Structural Proteomics: High-Throughput Methods* (Springer) pp 3–21
- [5] Brooks B R et al 2009 CHARMM: the biomolecular simulation program *J. Comput. Chem.* **30** 1545–614
- [6] Ponder J W and Case D A 2003 Force fields for protein simulations *Adv. Protein Chem.* **66** 27–85
- [7] Baumgartner M P and Zhang H 2020 Building admiral, an automated molecular dynamics and analysis platform *ACS Med. Chem. Lett.* **11** 2331–5
- [8] Yekeen A A, Durojaye O A, Idris M O, Muritala H F and Arise R O 2023 CHAPERONg: a tool for automated GROMACS-based molecular dynamics simulations and trajectory analyses *Comput. Struct. Biotechnol. J.* **21** 4849–58
- [9] Groen D, Bhati A P, Suter J, Hetherington J, Zasada S J and Coveney P V 2016 FabSim: facilitating computational research through automation on large-scale and distributed e-infrastructures *Comput. Phys. Commun.* **207** 375–85
- [10] Gygli G and Pleiss J 2020 Simulation foundry: automated and F.A.I.R. molecular modeling *J. Chem. Inf. Model.* **60** 1922–7
- [11] Singh H, Raja A, Prakash A and Medhi B 2023 Gmx_qk: an automated protein protein-ligand complex simulation workflow bridged to MM PBSA, based on Gromacs and zenity-dependent GUI for beginners in MD simulation study *J. Chem. Inf. Model.* **63** 2603–8
- [12] Maia E H B, Medaglia L R, Silva A M D and Taranto A G 2020 Molecular architect: a user-friendly workflow for virtual screening *ACS Omega* **5** 6628–40
- [13] Ganguly A, Tsai H-C, Fernández-Pendás M, Lee T-S, Giese T J and York D M 2022 AMBER drug discovery boost tools: automated workflow for production free-energy simulation setup and analysis (ProFESSA) *J. Chem. Inf. Model.* **62** 6069–83
- [14] Suruzhon M, Senapathi T, Bodnarchuk M S, Viner R, Wall I D, Barnett C B, Naidoo K J and Essex J W 2020 ProtoCaller: robust automation of binding free energy calculations *J. Chem. Inf. Model.* **60** 1917–21
- [15] Rego C R C, Schaarschmidt J, Schlöder T, Penaloza-Amion M, Bag S, Neumann T, Strunk T and Wenzel W 2022 SimStack: an intuitive workflow framework *Front. Mater.* **9** 877597
- [16] Hayashi Y, Shiomi J, Morikawa J and Yoshida R 2022 RadonPy: automated physical property calculation using all-atom classical molecular dynamics simulations for polymer informatics *npj Comput. Mater.* **8** 222
- [17] Martins L C, Cino E A and Ferreira R S 2021 PyAutoFEP: an automated free energy perturbation workflow for GROMACS integrating enhanced sampling methods *J. Chem. Theory Comput.* **17** 4262–73
- [18] Suplatov D, Sharapova Y and Švedas V 2020 EasyAmber: a comprehensive toolbox to automate the molecular dynamics simulation of proteins *J. Bioinform. Comput. Biol.* **18** 2040011
- [19] Martínez L, Andrade R, Birgin E G and Mario Martínez J 2009 PACKMOL: a package for building initial configurations for molecular dynamics simulations *J. Comput. Chem.* **30** 2157–64
- [20] Michaud-Agrawal N, Denning E J, Woolf T B and Beckstein O 2011 MDAnalysis: a toolkit for the analysis of molecular dynamics simulations *J. Comput. Chem.* **32** 2319–27
- [21] McGibbon R T, Beauchamp K A, Harrigan M P, Klein C, Swails J M, Hernández C X, Schwantes C R, Wang L-P, Lane T J and Pande V S 2015 MDTraj: a modern open library for the analysis of molecular dynamics trajectories *Biophys. J.* **109** 1528–32
- [22] Eastman P et al 2017 OpenMM 7: rapid development of high performance algorithms for molecular dynamics *PLoS Comput. Biol.* **13** e1005659
- [23] Abraham M J, Murtola T, Schulz R, Páll S, Smith J C, Hess B and Lindahl E 2015 GROMACS: high performance molecular simulations through multi-level parallelism from laptops to supercomputers *SoftwareX* **1** 19–25
- [24] Thompson A P et al 2022 LAMMPS - a flexible simulation tool for particle-based materials modeling at the atomic, meso and continuum scales *Comput. Phys. Commun.* **271** 108171
- [25] Goret G, Aoun B and Pellegrini E 2017 MDANSE: an interactive analysis environment for molecular dynamics simulations *J. Chem. Inf. Model.* **57** 1–5
- [26] Vieira Ribeiro J ao, Bernardi R C, Rudack T, Schulten K and Tajkhorshid E 2018 QwikMD - gateway for easy simulation with VMD and NAMD *Biophys. J.* **114** 673a–674a
- [27] Rusu V H, Horta V A C, Horta B A C, Lins R D and Baron R 2014 MDWiZ: a platform for the automated translation of molecular dynamics simulations *J. Mol. Graph. Model.* **48** 80–86
- [28] Biarnés X, Pietrucci F, Marinelli F and Laio A 2012 METAGUI: a VMD interface for analyzing metadynamics and molecular dynamics simulations *Comput. Phys. Commun.* **183** 203–11
- [29] Humphrey W, Dalke A and Schulten K 1996 VMD: visual molecular dynamics *J. Mol. Graph.* **14** 33–38
- [30] Sellis D, Vlachakis D and Vlassi M 2009 Gromita: a fully integrated graphical user interface to Gromacs 4 *Bioinform. Biol. Insights* **3** BBI–S3207
- [31] Martínez-Rosell G, Giorgino T and De Fabritiis G 2017 PlayMolecule ProteinPrepare: a web application for protein preparation for molecular dynamics simulations *J. Chem. Inf. Model.* **57** 1511–6
- [32] Schick T, Dwivedi-Yu J, Dessì R, Raileanu R, Lomeli M, Zettlemoyer L, Cancedda N and Scialom T 2023 Toolformer: language models can teach themselves to use tools (arXiv:2302.04761)
- [33] Karpas E et al 2022 MRKL systems: a modular, neuro-symbolic architecture that combines large language models, external knowledge sources and discrete reasoning (arXiv:2205.00445)

- [34] Yao S, Zhao J, Yu D, Du N, Shafran I, Narasimhan K and Cao Y 2022 ReAct: synergizing reasoning and acting in language models (arXiv:2210.03629)
- [35] Narayanan S *et al* 2024 Aviary: training language agents on challenging scientific tasks (arXiv:2412.21154)
- [36] Luo R, Sun L, Xia Y, Qin T, Zhang S, Poon H and Liu T-Y 2022 BioGPT: generative pre-trained transformer for biomedical text generation and mining *Brief. Bioinform.* **23** bbac409
- [37] Singhal K *et al* 2023 Large language models encode clinical knowledge *Nature* **620** 172–80
- [38] Gao S, Fang A, Huang Y, Giunchiglia V, Noori A, Schwarz J R, Ektefaie Y, Kondic J and Zitnik M 2024 Empowering biomedical discovery with AI agents *Cell* **187** 6125–51
- [39] Bran A M, Cox S, Schilter O, Baldassari C, White A D and Schwaller P 2024 Augmenting large language models with chemistry tools *Nat. Mach. Intell.* **6** 525–35
- [40] Boiko D A, MacKnight R, Kline B and Gomes G 2023 Autonomous chemical research with large language models *Nature* **624** 570–8
- [41] Villarreal-Haro J L, Gardier R, Canales-Rodriguez E J, Gomez E F, Girard G, Thiran J-P and Rafael-Patino J 2023 CACTUS: a computational framework for generating realistic white matter microstructure substrates.
- [42] Jablonka K M *et al* 2023 14 examples of how LLMs can transform materials science and chemistry: a reflection on a large language model hackathon *Digit. Discovery* **2** 1233–50
- [43] Su Y, Wang X, Ye Y, Xie Y, Xu Y, Jiang Y and Wang C 2024 Automation and machine learning augmented by large language models in catalysis study *Chem. Sci.* **15** 12200–33
- [44] Kim S, Jung Y and Schrier J 2024 Large language models for inorganic synthesis predictions *J. Am. Chem. Soc.* **146** 19654–9
- [45] Skarlinski M D, Cox S, Laurent J M, Braza J D, Hinks M, Hammerling M J, Ponnampati M, Rodrigues S G and White A D 2024 Language agents achieve superhuman synthesis of scientific knowledge (arXiv:2409.13740)
- [46] Lee W, Kang Y, Bae T and Kim J 2024 Harnessing large language model to collect and analyze metal-organic framework property dataset (arXiv:2404.13053)
- [47] Prince M H, Chan H, Vriza A, Zhou T, Sastry V K, Dearing M T, Harder R J, Vasudevan R K and Cherukara M J 2023 Opportunities for retrieval and tool augmented large language models in scientific facilities
- [48] O'Donoghue O, Shtedritski A, Ginger J, Abboud R, Ghareeb A E, Booth J and Rodrigues S G 2023 Bioplanner: automatic evaluation of llms on protocol planning in biology
- [49] Huang K, Qu Y, Cousins H, Johnson W A, Yin Di, Shah M, Zhou D, Altman R, Wang M and Cong L 2024 CRISPR-GPT: an LLM agent for automated design of gene-editing experiments
- [50] Ruan Y *et al* 2024 Accelerated end-to-end chemical synthesis development with large language models *ChemRxiv*
- [51] Kang Y and Kim J 2024 ChatMOF: an artificial intelligence system for predicting and generating metal-organic frameworks using large language models *Nat. Commun.* **15** 4705
- [52] Xie T, Ma R, Wang Q, Ye X, Liu F, Tai Y, Zhang Z and Yi Z 2024 Anywhere: a multi-agent framework for reliable and diverse foreground-conditioned image inpainting.
- [53] Ghafarollahi A and Buehler M J 2024 ProtAg ents: protein discovery via large language model multi-agent collaborations combining physics and machine learning.
- [54] Chiang Y, Chou C-H and Riebesell J 2024 LLaMP: large language model made powerful for high-fidelity materials knowledge retrieval and distillation (arXiv:2401.17244)
- [55] Ramos M C, Collison C J and White A D 2025 A review of large language models and autonomous agents in chemistry *Chem. Sci.* **16** 2514–72
- [56] Chandrasekhar A and Farimani A B 2025 Automating MD simulations for proteins using large language models: NAMD-Agent (arXiv:2507.07887)
- [57] Guilbert S, Masschelein C, Goumaz J, Naida B and Schwaller P 2025 DynaMate: an autonomous agent for protein-ligand molecular dynamics simulations (arXiv:2512.10034)
- [58] Chase H 2022 LangChain
- [59] The UniProt Consortium 2022 UniProt: the universal Protein Knowledgebase in 2023 *Nucl. Acids Res.* **51** D523–31
- [60] Developers M 2019 Molrender (available at: <https://github.com/molstar/molrender>) (Accessed 10 February 2025)
- [61] Nguyen H, Case D A and Rose A S 2018 NGLview—interactive molecular graphics for Jupyter notebooks *Bioinformatics* **34** 1241–2
- [62] Ouyang L *et al* 2022 Training language models to follow instructions with human feedback *Advances in Neural Information Processing Systems* vol 35 pp 27730–44
- [63] Achiam J *et al* 2023 GPT-4 technical report (arXiv:2303.08774)
- [64] Dubey A *et al* 2024 The llama 3 herd of models (arXiv:2407.21783)
- [65] Fireworks AI, Inc (Accessed 24 January 2025)
- [66] [www-cdn.anthropic.com](http://www-cdn.anthropic.com/fed9cc193a14b84131812372d8d5857f8f304c52/Model_Card_Claude_3_Addendum.pdf) (available at: www-cdn.anthropic.com/fed9cc193a14b84131812372d8d5857f8f304c52/Model_Card_Claude_3_Addendum.pdf) (Accessed 10 January 2025)
- [67] Anthropic 2024 The Claude 3 model family: Opus, Sonnet, Haiku
- [68] Jaech A *et al* 2024 OpenAI o1 system card (arXiv:2412.16720)
- [69] Hurst A *et al* 2024 GPT-4o system card (arXiv:2410.21276)
- [70] Laurent J M, Janizek J D, Ruza M, Hinks M M, Hammerling M J, Narayanan S, Ponnampati M, White A D and Rodrigues S G 2024 LAB-Bench: measuring capabilities of language models for biology research (arXiv:2407.10362)
- [71] Wang C, Luo W, Chen Q, Mai H, Guo J, Dong S, Xuan X, Li Z, Ma L and Gao S 2024 MLLM-Tool: a multimodal large language model for tool agent learning
- [72] Gao D, Ji L, Zhou L, Lin K Q, Chen J, Fan Z and Shou M Z 2023 AssistGPT: a general multi-modal assistant that can plan, execute, inspect, and learn
- [73] Wang G, Xie Y, Jiang Y, Mandekar A, Xiao C, Zhu Y, Fan L and Anandkumar A 2023 Voyager: an open-ended embodied agent with large language models (arXiv:2305.16291)
- [74] Campbell Q, Cox S, Medina J and White A D 2026 Code repository for MDCrow (<https://doi.org/10.5281/zenodo.18193478>)