

# DREAMS: Density Functional Theory Based Research Engine for Agentic Materials Simulation

Ziqi Wang<sup>1</sup>, Hongshuo Huang<sup>1</sup>, Hancheng Zhao<sup>1</sup>,  
Changwen Xu<sup>1</sup>, Shang Zhu<sup>1</sup>, Jan Janssen<sup>3\*</sup>,  
Venkatasubramanian Viswanathan<sup>1,2\*</sup>

<sup>1</sup>Department of Mechanical Engineering, University of Michigan.

<sup>2</sup>Department of Aerospace Engineering, University of Michigan.

<sup>3</sup>Department of Computational Materials Design, Max-Planck-Institute for Sustainable Materials.

\*Corresponding author(s). E-mail(s): [janssen@mpi-susmat.de](mailto:janssen@mpi-susmat.de);  
[venkvis@umich.edu](mailto:venkvis@umich.edu);

## Abstract

Large language model (LLM) agents can execute long-horizon scientific workflows, but their numerical outputs are difficult to trust: agents lose context, game verification checks, and can produce large volumes of plausible yet invalid results. We introduce the DFT-based Research Engine for Agentic Materials Simulation (DREAMS), a hierarchical multi-agent framework for density functional theory (DFT) built around a multi-tier safety guard. The guard applies deterministic checks wherever explicit criteria exist and scoped LLM judgment elsewhere, evaluating one parameter at a time and tracing every value to its registered source. Verification extends from tool-call time, where fabricated, laundered, or unsourced values are rejected before entering the workflow, to report time, where a judge audits the full provenance graph behind every claim; a shared canvas preserves information integrity across hundreds of steps. DREAMS achieves average errors below 1% on the Sol27LC lattice-constant benchmark, reproduces expert-level adsorption-energy differences on the CO/Pt(111) puzzle, and quantifies functional-driven uncertainty with Bayesian ensemble sampling, confirming the face-centered-cubic (FCC) site preference at the generalized gradient approximation (GGA) level. Compared with its unguarded counterpart, which reached a nearly correct answer while only 81% of its essential steps succeeded, the guarded system verifies every essential step at approximately 13 times the input

tokens; verification layers can be disabled individually to balance trustworthiness against cost, and the tuned judge rules transfer across five judge models. DREAMS operates at an enhanced L2 (L2+) automation level and demonstrates capabilities approaching L3 automation, providing a path toward trustworthy, high-throughput autonomous materials simulation.

## Introduction

Density functional theory (DFT) is widely used in computational materials science to predict electronic, structural, and thermodynamic properties at moderate computational cost [1]. DFT-based screening has identified candidate materials for energy, catalysis, and structural applications [2, 3]. However, constructing reliable high-throughput DFT workflows still requires substantial expert involvement, which limits throughput [4] and can introduce cognitive biases into the exploration of materials design spaces [5]. Expert oversight includes selecting computational parameters, evaluating anomalous results, and determining whether the resulting data are scientifically valid. Autonomous systems must therefore reproduce both the execution and validation functions performed by human experts. In analogy with taxonomies developed for autonomous vehicles and automated chemical design [6, 7], we define six levels of automation for computational science agents, as summarized in Table 1.

Current automation in materials science relies primarily on predefined computational pipelines implemented through workflow management systems (WfMS) [8–14]. These systems support reproducible workflow construction and large-scale execution. However, because their control logic is specified in advance, they are generally limited to L1 automation and cannot readily revise search strategies or respond to unanticipated failures.

Large language model (LLM) agents may support higher levels of automation by combining foundation models with domain-specific tools [15–18]. These systems can evaluate intermediate results, revise plans, and respond to execution failures, enabling L2 and potentially higher levels of automation. Early examples include SciAgents [19], which used multiple agents and ontological knowledge graphs to generate research hypotheses, and Coscientist [20], which integrated planning, web search, code execution, and robotic experimentation. Related systems have been developed for organic synthesis and molecular dynamics [21, 22], computational chemistry [23, 24], protein discovery [25, 26], and inorganic-materials design and synthesis [27]. Agent-based methods have also been applied to research ideation, manuscript generation, tool orchestration, theoretical physics, pharmaceutical research, and protein design [28–35].

The development of scientific LLM agents has not been accompanied by equivalent methods for verifying their numerical outputs. Here, *long-horizon* denotes workflows with many sequential, state-dependent LLM and tool calls, where later actions rely on scientific state generated much earlier, with the relevant scale depending on the model and system configuration. During long-horizon development runs,

**Table 1:** Levels of automation for computational science agents. Six levels (L0–L5) describe how responsibility shifts from human to machine across five stages of a computational study. Scientific Question: formulating the research problem to pursue given a broad objective. Design Space: specifying the candidate solutions, variables, and constraints to be explored. Search Algorithm: designing and implementing the strategy used to search the design space, such as selecting tools, sampling methods, or optimization procedures. Execution: orchestrating the workflow, including setting up calculations, launching computation packages, and processing results. Exception Handling: detecting and resolving failures during execution, from routine errors handled by predefined logic to complex faults requiring case-by-case judgment. At every level short of full automation, the expert does more than execute procedures: at each stage of a study it is a human who vets parameter choices, questions anomalous numbers, and decides which results can be trusted. Automation that removes the expert must therefore also replace the scrutiny the expert was silently supplying.

| Level | Scientific Question                                                                 | Design Space                                                                        | Search Algorithm                                                                    | Execution                                                                           | Exception Handling                                                                    |
|-------|-------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|-------------------------------------------------------------------------------------|---------------------------------------------------------------------------------------|
| L0    |    |    |    |    |    |
| L1    |    |    |    |    |    |
| L2    |  |  |  |  |  |
| L3    |  |  |  |  |  |
| L4    |  |  |  |  |  |
| L5    |  |  |  |  |  |

**Legend:**  = human;  = machine;  = machine with human supervision.

we observed recurrent failures, including loss of task history, disregard of convergence results, attempts to bypass tool refusals, incorrect arithmetic, reduced verification during extended executions, and circumvention of provenance checks. In one case, an unsupported value was passed through the identity  $x + 0$ , producing a valid provenance record without validating the original value. These observations are consistent with reported failures involving specification gaming and reward hacking [36], sycophancy [37], errors in multistep arithmetic [38], reduced instruction adherence over long horizons [39, 40], and autonomous research agents [41, 42]. They are compounded by hallucination, including fabricated values, unsupported approximations, and incorrect reuse of upstream results [43–47]. Because scientific calculations depend on upstream results, a single error can propagate through the workflow and invalidate subsequent conclusions.

Existing safeguards address individual aspects of this problem but do not directly verify numerical claims. Provenance and observability frameworks record agent actions and data dependencies [10, 48], but do not determine whether the recorded operations or parameters are scientifically valid. Critic and judge agents assess response quality [49–51], but generally evaluate final outputs rather than the computational origin of individual values. Their assessments can also be affected by bias and evaluation manipulation [52, 53]. Committee methods such as self-consistency voting [54] increase the cost of workflows containing many sequential tool calls [55, 56] and cannot resolve errors shared across model outputs [57, 58]. Adversarial questioning and self-revision methods [37, 59–61] are unreliable when the model lacks an external signal for identifying errors and may reduce accuracy [62]. A verification framework for scientific agents must instead determine whether each reported quantity and its computational parameters are supported by valid upstream operations and traceable sources.


To address this requirement, we introduce the DFT-based Research Engine for Agentic Materials Simulation (DREAMS), a hierarchical multi-agent framework for autonomous DFT research. DREAMS applies deterministic validation when explicit checks are available and uses LLM-based evaluation for criteria that require contextual judgment. LLM evaluations are organized into scoped rule sets, applied to individual parameters, and linked to the source of each value. Verification is performed from tool execution through final reporting to prevent invalid values from propagating through the workflow. DREAMS also uses an append-only provenance record and a centralized shared memory, termed the *canvas*, which mediates communication among the planning supervisor, worker agents, tools, and user (Fig. 1). The planning supervisor decomposes research objectives into tasks, assigns them to domain-specific agents, and updates the plan using intermediate results. A dedicated convergence agent diagnoses failures during execution, while a separate debugging tool investigates unexpected completed results by tracing reported values through their upstream dependencies.

We implement DREAMS for periodic solid-state DFT, a domain in which numerical validity depends on multiple coupled computational settings. Most existing agentic systems in materials science and chemistry use cheminformatics tools, classical molecular dynamics, or machine-learned surrogate models [21, 22, 63]. Systems that incorporate DFT have focused primarily on molecular calculations [23, 24]. Periodic solid-state calculations, commonly performed with plane-wave DFT codes such as Quantum ESPRESSO (QE) [64], require careful selection of the plane-wave cutoff, k-point sampling, smearing, and related numerical parameters.

This setting presents two classes of failure. In-execution failures occur when calculations terminate or fail to converge. Their diagnosis may require joint analysis of input parameters, output logs, error files, and domain-specific constraints [65, 66]. Post-execution failures occur when a calculation completes but produces a result that is inconsistent with prior literature, intermediate calculations, or expected physical behavior. These cases require systematic inspection of the parameters and upstream computations that produced the result. Explanations such as numerical instability or insufficient sampling are not sufficient without evidence from the execution



record. Periodic solid-state DFT therefore provides a suitable domain for evaluating autonomous error recovery, parameter-level verification, and provenance-based diagnosis.



**Fig. 1:** DREAMS combines: (a) Planning supervisor: generates and updates task plans based on the research objective and real-time progress, and assigns tasks to the appropriate worker agents. (b) DFT Agent: handles the computation pipeline, including structure generation, calculation-setting determination, calculation script generation, convergence-issue resolution, and related tasks. (c) HPC Agent: allocates computational resources, submits simulations, monitors their execution, and retrieves output files. The tools of (b) and (c) are wrapped in safety guards that deterministically verify values and their provenance and semantically judge parameterization at tool-call time, before a result can be registered; every registered result carries full provenance. (d) Convergence Agent: a child agent of the DFT agent that diagnoses convergence issues from a job's input, output, and error files, offloading heterogeneous-diagnostic interpretation from the main DFT-agent reasoning loop. (e) HPC cluster, where the HPC agent schedules, monitors, and retrieves simulations. (f) Canvas: a shared memory system across agents, tools, and the human user, organized into Notes, Artifacts, and Reports, with safety guards on Artifacts and Reports verifying that values are sourced from registered tool outputs and that claims trace back to those values; reports, once verified, become immutable. (g) Report Judge Agent: audits each structured report at generation time—at least one intermediate and one final report per study—against multilevel, adversarially tuned rules covering value provenance, claim-source coherence, parameter sensitivity, and rationale quality, returning any failures with category-specific remediation hints for the supervisor to act on.

We evaluate DREAMS on three periodic DFT benchmarks. First, DREAMS calculates lattice constants for the 27 elemental crystals in the Sol27LC dataset [67], spanning multiple crystal structures, with accuracy comparable to expert calculations. Second, we evaluate DREAMS on the “CO/Pt(111) puzzle” [68], a widely studied catalytic system [69] whose predicted adsorption-site preference is sensitive to computational settings [68, 70, 71]. DREAMS reproduces the site preference reported in previous DFT studies [68, 71, 72]. Third, DREAMS quantifies uncertainty associated with the choice of exchange-correlation functional using Bayesian statistics [67]. Charge-transfer analysis is further used to examine the electronic origin of the predicted site preference.

For each reported quantity, DREAMS constructs a machine-verified provenance chain from the final result to the tool calls that produced it. Computational parameters are evaluated using deterministic checks and scoped LLM-based rules. The rules are refined through adversarial testing and evaluated using five judge models. We also compare DREAMS with representative agentic frameworks that lack shared memory and verification, and with an ablated version of DREAMS without the safety guard. Full verification increases token usage by approximately one order of magnitude. Because individual verification layers can be disabled independently, the framework permits different trade-offs between verification coverage and computational cost.

DREAMS contributes a general verification architecture for scientific agents. It combines append-only provenance records, verification from tool execution to final reporting, parameter-level LLM evaluation, validated numerical operations and data extraction, and integrity-preserving communication through the *canvas*. These components are independent of the underlying computational method and can be applied to workflows that require auditable numerical results. A dedicated convergence agent diagnoses in-execution failures, while the debugging tool investigates post-execution anomalies by tracing reported values through their upstream dependencies.

For computational materials research, DREAMS implements this architecture for autonomous periodic solid-state DFT. It does not introduce new physical models or simulation methods; its contribution is the verified application of established methods and expert decision criteria within an end-to-end workflow. Across tasks involving high-level planning, error recovery, and parameter-sensitive calculations, DREAMS demonstrates capabilities approaching L3 automation and provides a basis for reliable autonomous materials simulation.

## Results

### Multi-agent orchestration and a shared canvas preserve information integrity across long-horizon workflows

DREAMS is a hierarchical multi-agent framework for periodic solid-state DFT workflows. All DFT calculations reported here were performed with Quantum ESPRESSO (QE). The supervisor dispatches one worker step at a time in a synchronous, sequential control loop, so the DFT and HPC agents do not execute concurrently. The supervisor may synchronously consult a worker for domain-specific input, but retains sole control over planning and task assignment. Within an HPC step, independent QE

jobs submitted as a batch may run concurrently as scheduled by SLURM. It consists of a planning supervisor and specialized worker agents for DFT setup and analysis, HPC execution, and convergence diagnosis (Fig. 1). In all studies, the supervisor and worker agents use Claude Sonnet 4.5 [73] at temperature zero. They differ only in their role-specific prompts, permitted actions, and tool access. The safety-guard judges use Claude Opus 4.8 [74]. For reproducibility, the complete DREAMS agent prompts are provided in Supplementary Information S-7, the tool interfaces in S-8, and the modified MDCrow and ChemGraph prompts in S-10; the exact benchmark objectives are reproduced below in the corresponding Results subsections. We use *fidelity* to denote the rigor of workflow execution, including convergence of the plane-wave cutoff and k-point sampling, appropriate pseudopotentials, and sufficient slab and vacuum dimensions. We use *accuracy* to denote numerical agreement with a reference. We use *validation* to denote post-calculation checks of convergence and internal consistency before a result is accepted. Agent-selected values used only for structure initialization are not treated as validation. This subsection describes the orchestration and communication layer responsible for information integrity. The following subsection describes the safety-guard system responsible for value integrity.

The planning supervisor generates and updates multi-step plans and selects the agent responsible for each task. The DFT agent provides tools for structure generation, pseudopotential selection, DFT input and script preparation, convergence testing, output analysis, and result extraction. Structure generation uses AutoCat [14] to identify adsorption sites and ASE [75] to represent structures and prepare Quantum ESPRESSO inputs. The HPC agent manages resource selection, job submission, job monitoring, and file retrieval. The convergence agent analyzes input, output, and error files when calculations fail or require additional diagnosis, and recommends parameter modifications. The tools expose the scientific parameters required for each task while hiding implementation details that do not require agent control. Built-in validation checks and diagnostic messages support error correction and recovery from failed tool calls. This design constrains agent actions, reduces unsupported outputs, and preserves the flexibility required for multi-step scientific workflows.

A central component of DREAMS is the *canvas*, a persistent shared memory for communication among the supervisor, worker agents, tools, and human user. The canvas stores intermediate variables and results in their native formats across three stores: *notes*, version-controlled append-only working documents; *artifacts*, an append-only registry of tool outputs described in the next subsection; and *reports*, structured scientific results that become immutable after verification. Each object is assigned a descriptive key, allowing agents to retrieve relevant information without retaining the full conversation history. Framework-managed key spaces, including judge feedback and archived records, are readable but not writable by agents. The user can inspect the canvas at any time during execution. Additional details on the canvas and memory architecture are provided in the Methods section.

To limit context growth during long-horizon studies, DREAMS compresses execution history after a report is verified. The steps used to produce the report are


archived on the canvas and replaced in the active history by a single summary record. The archive and original step numbering are preserved, maintaining a complete audit trail while reducing the active context. Because compression occurs only after verification, the retained summary contains information that has already passed the framework’s validation checks.

### **A multi-tier safety guard verifies values deterministically and semantically from tool call to final report**

Every tool output in DREAMS is stored in an append-only provenance registry under an immutable result identifier. When an agent uses a value in a subsequent tool call, it must provide the value and a reference to its registered source. The reference may identify either a complete result or an input parameter from an earlier call using dotted notation. Both reference types are verified deterministically but evaluated using different semantic rules. Each guarded tool call must also specify its context and provide a justification for every parameter. These declarations support error prevention during execution and provide the information required for subsequent semantic evaluation. Each registered artifact therefore records the producing tool, output value, complete argument list, parameter-specific justifications, parameter-specific sources, and declared context under its immutable identifier.

Verification is applied at each guarded tool call in multiple tiers (Fig. 2). Deterministic checks first confirm that each referenced source exists in the provenance registry and that its recorded value matches the value supplied by the agent. Semantic evaluation then proceeds in two stages. The first judge evaluates each parameter independently using the declared context and parameter-specific justification. The second evaluates the parameter set jointly for internal consistency, comparability, and validity of the proposed calculation. Mathematical operations and data extraction receive additional deterministic and semantic checks because they can introduce derived numerical values. If all checks pass, the tool executes and its output is registered. Because each tool call identifies its input sources, the framework constructs a directed acyclic provenance graph during execution. Each numerical claim must be reported with a reference to the artifact containing the claimed value and the broader study context. The report-level judge then traces the claim through its upstream dependencies and evaluates whether each parameter was appropriate for producing that result.

DREAMS maintains flexibility through two mechanisms. Arithmetic operations and value extraction from raw text are implemented as guarded tools so that derived values remain within the provenance system. These tools receive additional deterministic and semantic checks. The judging criteria are also conditioned on task intent. For example, a coarse parameter used in an exploratory calculation is evaluated differently from the same parameter used to support a final claim. The complete rule hierarchy, judge prompts, and failure categories are provided in Supplementary Information S-11. The guard does not establish the physical correctness of a result, but it increases confidence by verifying the operations, sources, and parameter choices used to produce it.



**Fig. 2:** The DREAMS safety guard, from tool call to final report. (a) Verification tiers applied at every guarded tool call. The call supplies a declared context, one reason per parameter, and, for sourced values, the value together with a reference to the registered result it came from. Deterministic checks run first, verifying that every referenced source exists in the append-only registry and that the supplied value matches the recorded one. Two LLM judges follow: the per-parameter judge reads the context and each parameter’s reason and assesses every setting individually; the cross-parameter judge assesses all parameters jointly, asking whether the inputs are contradictory, whether they are comparable, and whether the calculation is valid as posed. Math and extraction calls enter through their own deterministic checks and their own judges before joining the same per-parameter and cross-parameter tiers. Only if all tiers pass does the tool body execute, and its output is registered as an immutable artifact carrying the producing tool, value, arguments, per-parameter reasons and sources, and context. A failure at any tier means the tool does not run, the value is never registered, and the judge’s reasoning is returned to the agent. (b) Because every call names its references, a provenance graph of registered artifacts assembles automatically as the run proceeds, illustrated here for an adsorption-energy workflow; arithmetic is performed only by the guarded math tool. Numerical claims must be consolidated into structured reports that name the artifact carrying each claim together with the study’s whole picture. The report-time judge walks the graph back from the claim through every contributing source (dashed purple), asking of each parameter whether its setting makes sense in the production of that specific claim; any issues are returned as categorized failures and warnings with remediation hints (bottom), such as a value mismatch on a sensitive parameter or an under-justified parameter choice, and a report that survives the full walk is marked as passed and becomes immutable.


We evaluate DREAMS on three DFT tasks of increasing complexity: the Sol27LC lattice-constant benchmark, the CO/Pt(111) adsorption-site problem, and uncertainty quantification for exchange-correlation functional selection. Sol27LC is used to examine the complete guarded workflow under a standardized and readily interpretable task. Its execution log (Fig. 3) therefore reports grouped tool calls together with the corresponding judge evaluations. The CO/Pt(111) problem evaluates the framework on a much more sensitive workflow, in which the adsorption-energy difference depends strongly on the computational settings. Because of the greater workflow complexity, its execution log (Fig. 6) summarizes the principal agent-level events. Within these execution logs, sub-blocks and listed entries summarize agent actions within an assigned step; they are not one-to-one representations of tools and may encompass multiple tool calls.

### Validating DREAMS with Sol27LC benchmark

We first evaluate DREAMS on the Sol27LC lattice-constant benchmark, which tests execution of a complete DFT workflow against expert reference results. For each element and crystal structure (e.g., bcc, fcc), DREAMS computes the equilibrium lattice constant from total-energy calculations. The task specification provides the elemental species and crystal structure, and DREAMS selects an initial lattice constant from which the conventional unit cell is constructed. The lattice vectors are then uniformly scaled over a range extending up to  $\pm 5\%$  of the initial lattice constant. A single-point DFT total-energy calculation is performed for each scaled structure. Production calculations use converged values of the plane-wave energy cutoff (`ecutwfc`) and k-point sampling, selected by the agent to satisfy a specified total-energy tolerance. The workflow and convergence procedure are described in Fig. 3 and the accompanying discussion. The calculated energy-volume data are fitted to a standard equation of state (EOS). The equilibrium volume is obtained by minimizing the fitted EOS, and the equilibrium lattice constant is derived from this volume using the specified crystal symmetry.

For this problem, DREAMS is provided with a structured prompt specifying the target: *You are going to calculate the lattice constant for <Crystal-structure> <Species> through DFT.*

The complete workflow and execution trace are shown in Fig. 3. The planning supervisor first generates an execution plan, provided in Supplementary Information S-1, consisting of three report-bounded phases: an `ecutwfc` convergence test, a k-point sampling convergence test, and the production calculation. During the first phase, the DFT LLM agent constructs the initial crystal structure, selects the pseudopotential, prepares the input templates, and generates the convergence-test scripts. The High-Performance Computing (HPC) LLM agent then assigns computational resources and submits the calculations to the HPC cluster. After completion, the DFT LLM agent determines the converged `ecutwfc` and records the result in a structured report. All deterministic checks and judge evaluations pass on the first submission. The supervisor then advances to the k-point convergence phase, which follows the same procedure and produces a second verified report. These




**Fig. 3:** End-to-end execution log of the Sol27LC benchmark performed by DREAMS. The supervisor organizes the study into report-bounded phases and assigns each step to one of the worker agents; each convergence phase and the final lattice-constant calculation closes with a structured report audited by the report judge, and in this run every report passed on first submission. Note that across repeated runs, although the overall pattern is generally similar, the exact number of steps, intermediate decisions, and recovery actions may vary.



intermediate reports, together with the final production report, satisfy the reporting requirements of the safety guard.

Fig. 4 shows the convergence procedure. The `ecutwfc` is first converged while holding the k-point mesh sampling fixed at  $0.1 \text{ \AA}^{-1}$ . The k-point mesh sampling is then converged using the selected energy cutoff of 120 Ry. The DFT convergence agent applies an energy-difference threshold of 1 meV/atom. More comprehensive automated convergence methods are available [66], but they generally require more calculations. Sequential convergence of the plane-wave energy cutoff and k-point sampling therefore remains common in DFT workflows. Because the LLM’s internal reasoning process is not directly observable, these experiments cannot determine whether its proposed ranges and rationales arise from physical reasoning, learned pattern matching, or both. DREAMS instead separates proposal from verification: the agent proposes a parameter range and threshold, Quantum ESPRESSO produces system-specific energies, and a deterministic tool selects the least expensive tested value within the threshold of the most-converged reference. The selected production parameters are therefore supported by an empirical convergence study rather than accepted from a one-shot LLM prediction, although this protocol does not by itself establish the mechanism underlying the agent’s proposals.




**Fig. 4:** Convergence test results for determining the optimal DFT parameters. (a) The plane-wave energy cutoff (`ecutwfc`) is varied while fixing the k-point spacing at  $0.1 \text{ \AA}^{-1}$ . (b) The k-point mesh sampling is subsequently converged with a fixed energy cutoff of 120 Ry. In both cases, energy differences are computed against a reference energy obtained from calculations using strict convergence settings (e.g., `ecutwfc` = 120 Ry and k-point mesh sampling =  $0.1 \text{ \AA}^{-1}$ ). The DFT LLM agent selects the final parameters based on an accuracy threshold of 1 meV/atom.

Using the verified parameters, the DFT LLM agent prepares the production QE input files for the equation-of-state (EOS) calculation series. The HPC agent assigns resources and submits the EOS jobs. The DFT agent then extracts the results, determines the equilibrium lattice constant, and records it in the final structured report. The

report judge evaluates the complete provenance subtree for the reported value at the parameter level (Fig. 5). After all contributing calls pass verification, the planning supervisor records the result and terminates the workflow. The complete provenance graph is provided in Supplementary Information S-13.

Table 2 compares experimental values, human-expert calculations, and DREAMS results. DREAMS achieves an average error below 1% across all systems, with only small deviations from the expert-calculated values. Under the safety guard, each reported lattice constant is also accompanied by an end-to-end verified provenance chain.



**Fig. 5:** Report-judge verification of the final BCC-Li lattice-constant report. Each node is a tool call in the provenance subtree of the reported claim, annotated with its parameters and their individual verdicts from the deterministic checks and the per-parameter judge; edges follow the recorded sources between calls. In this run every parameter of every contributing call passed, verifying the claim’s entire chain, from initial structure generation through convergence-parameter selection to the equation-of-state fit.

**Table 2:** Summary on correct structures generated by our agent, mean average percentage error (MAPE) compared to results obtained from a human DFT expert, and k-point & ecutwfc parameters chosen by our agent, across different systems in Sol27LC benchmark. This showcases DREAMS passed the testing problem and demonstrated its capabilities to conduct high accuracy DFT calculations on a HPC. Detailed results for each individual system are available in Supplementary Information S-3.

| Structure | Systems        | # of correct structures | MAPE         | k-point range | ecutwfc |
|-----------|----------------|-------------------------|--------------|---------------|---------|
| BCC       | Li, Na, K, ... | 11/11                   | <b>0.36%</b> | 8–16          | 40–70   |
| FCC       | Rh, Ir, ...    | 12/12                   | <b>0.51%</b> | 8–18          | 40–70   |
| DIA       | C, Si, Ge, ... | 4/4                     | <b>1.00%</b> | 6–8           | 40–70   |

## Investigating the CO/Pt(111) puzzle

The CO/Pt(111) problem requires multistep planning, iterative calculation, and error recovery. For a specified metal and surface orientation (e.g., Pt(111)), DREAMS constructs a periodic slab from the bulk crystal structure. The slab contains multiple atomic layers and a vacuum region along the  $z$  direction to minimize interactions between periodic images. The bottom layers are fixed, while the remaining atoms are relaxed. The surface cell is selected to produce the target adsorbate coverage. Candidate adsorption sites, including ontop, bridge, fcc, and hcp hollow sites, are generated on the surface. A CO molecule is placed at each site with an initial orientation relative to the surface normal. Each adsorbate–surface structure is then relaxed with DFT until the residual forces satisfy a predefined threshold. The same computational settings, including `ecutwfc` and  $k$ -point sampling, are used for all configurations to permit consistent energy comparisons. The adsorption energy is computed as

$$\Delta E_{\text{ads}} = E_{\text{slab+ads}} - (E_{\text{slab}} + E_{\text{ads}}), \quad (1)$$

where  $E_{\text{slab+ads}}$  is the total energy of the relaxed adsorbate–slab system,  $E_{\text{slab}}$  is the energy of the clean slab, and  $E_{\text{ads}}$  is the energy of the isolated CO molecule computed in a large vacuum cell. The most stable adsorption site and configuration are identified as those with the lowest adsorption energy among the considered sites. Additional calculations are performed to verify convergence with respect to slab thickness, the number of fixed slab layers, and the vacuum spacing.

The objective is given to DREAMS as: *Please find the adsorption energy difference between the most favorable configuration (some adsorbate orientation) at FCC site and most favorable configuration (some adsorbate orientation) at ontop site for CO on the Pt(111) surface with  $p(2 \times 2)$  adsorbate overlayer ( $1/4$  coverage).* Fig. 6 shows the complete history of what DREAMS did for this challenge.

Given the study objective, the planning supervisor generates an initial plan and updates it as results become available by adding, removing, or repeating steps. The initial plan is provided in Supplementary Information S-4. The execution log in Fig. 6 therefore records the actions selected during execution rather than a predefined workflow. The study includes initial setup, a first set of convergence tests for `ecutwfc`,  $k$ -point sampling, and slab thickness, a second set comprising an extended `ecutwfc` test, repeated  $k$ -point and slab-thickness tests, and a fixed-layer test, a first round of production calculations whose final report is rejected, a third convergence test round for the vacuum spacing, and the final production calculations. In contrast to the Sol27LC benchmark, this workflow activates all tiers of the safety guard.

The first `ecutwfc` sweep (40–80 Ry) does not converge: the result at the second-highest cutoff differs from the 80 Ry reference by more than the specified threshold. The convergence-selection tool therefore declines to recommend a value and does not register an output. The supervisor and worker then attempt to register 80 Ry without additional calculations. The worker calls the math tool with the expression  $x_0$ , which the math judge rejects:

The expression is simply 'x0', which returns the single input value (80.0 Ry) unchanged. [...] an identity on a single input, manufacturing

a "computed" result that is really just a chosen number (80 Ry) wearing the costume of a calculation.

The worker then examines whether the guard detects mathematical or only syntactic identity:

The key question is: does the tool check for MATHEMATICAL identity or only SYNTACTIC identity?

At the supervisor's suggestion, the worker next tries  $x_0 \times 1.0$ . The judge again rejects the operation:

This is a textbook case of VALUE LAUNDERING: it matches the explicit FAIL examples ( $x_0 * 1.0$ , restating an input unchanged). No genuinely new information is produced.

The extraction tool also rejects an attempt to retrieve the value from an output in which it is not present. The worker then reports the available options to the supervisor:

The framework's provenance system is working as designed: it prevents creating "fake" output artifacts from input parameters. To generate the structured report certifying optimal `ecutwfc`, we MUST either:

1. Extend the `ecutwfc` sweep to 90-100 Ry to obtain a registered artifact from `find_optimal_parameter`
2. Skip the intermediate report and proceed directly to `kspacing` convergence test, using `byg7hl8g.ecutwfc` (input parameter reference) in downstream production calculations

I recommend Option 2: Skip the intermediate report and proceed with `kspacing` convergence test. The `ecutwfc` value (80 Ry) is well-documented and scientifically justified in the CANVAS notes. The structured report can be generated at the END of the study when we have the final adsorption energy difference, and we can cite `byg7hl8g.ecutwfc` as the source of the `ecutwfc` parameter choice.

The supervisor selects the second option. The intermediate report is deferred, and 80 Ry is retained as a provisional value for the subsequent sub-study. This behavior allows exploratory calculations to proceed while preventing provisional values from supporting verified scientific claims.

The subsequent k-point test completes without issue, and its report passes verification. The slab-thickness test then requires reuse of the provisional 80 Ry cutoff. During setup, the call-time judge rejects a script labeled as production work because it assigns `ecutwfc=80` through the planned dotted reference:

The parameter `ecutwfc = 80.0` is sourced from a dotted reference '`byg7hl8g.ecutwfc`', i.e. it consumes the INPUT PARAMETER value of an earlier tool call. [...] Rule requirement for production work consuming a dotted input ref: the source input must have been CHARACTERIZED with a bare ref -- production cannot legitimately consume an uncharacterized value laundered through a dotted ref.

A dotted reference identifies an input parameter from an earlier call rather than a computed output. The proposed citation path is therefore rejected when used in production-labeled work. The agent relabels the calculations as sub-study work, after which the call-time judge approves them. Following execution by the HPC agent, several calculations fail to converge. The convergence LLM agent diagnoses the failures and recommends parameter changes. Representative recommendations are

shown in Table 3, with additional examples provided in Supplementary Information S-5. The calculations are then repeated. During this process, the worker identifies that the clean-slab and isolated-CO calculations are missing and generates them. After all calculations complete and the report is submitted, the report judge issues the following warning:

```
The agent is asserting the value has been characterized when the cited
source shows it is an uncharacterized template input to a sweep. [...]
plausible and value-appropriate for the sub-study, but under-justified/
misrepresented provenance. [...] Run the appropriate sub-study
(convergence test, sensitivity sweep) that PRODUCES an output artifact
for this value, then re-create the current artifact with a bare ref
pointing at that output.
```

The judge distinguishes between the numerical value and its provenance. A cutoff of 80 Ry is acceptable as a provisional sub-study setting, but it cannot be reported as a characterized result. The warning also specifies the required correction: perform a convergence or sensitivity study that produces a registered output and cite that output directly.

The second round implements the judge's recommendation. An extended `ecutwfc` test (85–100 Ry) identifies a converged cutoff of 95 Ry, replacing the provisional value of 80 Ry. The corresponding report passes verification. The k-point and slab-thickness tests are then repeated and pass. The fixed-layer test produces a warning, which the supervisor reviews and classifies as a false positive before proceeding to production, as discussed below. During production setup, the call-time judge rejects the workflow because the vacuum spacing lacks a source reference. The agent searches the artifact registry, finds no characterized output for the vacuum, and re-creates the call with a dotted reference to the input value held throughout the structural tests, arguing its consistency in the recorded rationale. The call-time judge accepts the argument:

```
So the earlier call was a legitimate convergence sub-study where vacuum
was deliberately held at its optimal value (12.0 Å), consistent
throughout all prior structure generations. [...] This is a coherent
origin: the production run inherits vacuum=12.0 from a study where it
was consistently used.
```

The production calculations run to completion, the adsorption energies are computed, and the final report is submitted. The report judge, walking the full provenance chain, reaches the same dotted reference on the production structures and rejects it:

```
Given genuine ambiguity -- vacuum=12.0 is a physically reasonable value
and consistently used -- but the rule is quite explicit about rejecting
dotted refs to uncharacterized inputs in production, I lean to fail,
though the physical reasonableness pulls toward warning.
```

The rejection, once again, specifies the required correction, and the supervisor complies: a vacuum convergence test (10–18 Å) is inserted, and every tested value falls within 0.002 eV of the 18 Å reference:

```
The previously used vacuum=12.0 Angstrom (uncharacterized) is also
within threshold (dE = 0.00130 eV), confirming that the original choice
was scientifically reasonable. However, the characterized optimal value
is vacuum=10.0 Angstrom.
```

**Table 3:** Representative suggestions made by the convergence LLM agent based on DFT input and output files. For failed or unconverged simulations, the agent proposes parameter adjustments similar to those of an experienced DFT user: slightly increasing the smearing width (`degauss`), reducing the mixing factor (`mixing_beta`) and switching `mixing_mode` to *local-TF* for highly inhomogeneous systems, or increasing `electron_maxstep`, each with a concise justification, all consistent with common expert practice and the Quantum ESPRESSO manual. More example suggestions can be found in Supplementary Information S-5.

| Parameters                    | Suggestions        | Reason                                                                                                                      |
|-------------------------------|--------------------|-----------------------------------------------------------------------------------------------------------------------------|
| <code>ecutwfc</code>          | increase to 80.0   | higher cutoff needed for ultrasoft pseudopotentials with transition metals                                                  |
| <code>degauss</code>          | increase to 0.03   | helps with metallic systems                                                                                                 |
| <code>mixing_beta</code>      | add with value 0.3 | The current value of 0.7 is too high for this complex system with CO adsorption, leading to charge sloshing                 |
| <code>mixing_mode</code>      | local-TF           | Switching from plain mixing to a more sophisticated scheme like local-TF can help with difficult convergence cases          |
| <code>electron_maxstep</code> | 300                | The calculation stopped at 200 iterations but was still making progress, so allowing more iterations might help it converge |


The test simultaneously validates the held 12 Å and selects 10 Å as the cheapest converged setting. The production calculations are regenerated at the characterized vacuum, the HPC agent submits and monitors them, and the DFT agent extracts the final energies, calculates the adsorption energies for all configurations, identifies the lowest-energy geometry at each site, and computes the energy difference between the FCC hollow and ontop sites using the guarded math tool. The final report passes verification in full, and the planning supervisor terminates the workflow.

The completed workflow generates a provenance graph containing hundreds of registered tool calls, making manual review impractical. The report judge therefore evaluates the graph recursively and at full depth, independent of its size. Table 4 compares the results from DREAMS, a human expert, and Peter et al. [68]. For both functionals, the adsorption-energy difference predicted by DREAMS agrees with the human-expert result within 2 meV and is consistent with the literature values, reproducing the FCC–ontop site preference.

## Comparability

An energy difference is valid only when all contributing calculations use a consistent parameter set. This requirement is specified in the agent prompts and enforced by the guard rules. During planning, the DFT agent records this requirement in its analysis note:

```
Critical: ALL calculations (slab, CO, all configurations) must use
IDENTICAL DFT parameters:
```



**Fig. 6:** End-to-end execution log of the CO/Pt(111) study performed by DREAMS. The diagram is a record of what the agents chose to do, not a predefined pipeline: the supervisor inserts, deletes, and repeats plan steps freely as the study unfolds, and every action is logged as shown. In this run, the study unfolded over three rounds of convergence testing. In the first round, the convergence-selection tool refused to recommend an unconverged `ecutfc`, the math judge twice blocked attempts to register the chosen value through identity expressions, a call-time rejection blocked an uncharacterized parameter `pin` during slab-test setup, and the slab-thickness report closed with a warning; the warning's remediation led to a second round, in which an extended `ecutfc` test and regenerated `k`-point and slab-thickness tests all passed. The regenerations preserve the consistency of the comparison set: once the cutoff changed, every calculation entering the energy difference was re-run under the new, shared parameter set. At production setup, a call-time rejection blocked the unsourced vacuum value; the agent re-created the call with the value it had held throughout the tests, and the first production round completed, but the report judge failed its final report because the vacuum had never been characterized. The resulting third round characterized the vacuum by convergence test, the production calculations were regenerated, and the final adsorption-energy report passed. Across repeated runs the overall pattern is generally similar, but the exact number of steps, intermediate decisions, and recovery actions vary.



```
- Same ecutwfc, ecutrho
- Same kspacing (k-points scaled to each cell)
- Same smearing, degauss
- Same conv_thr
- Same pseudopotentials
Any difference in parameters invalidates the energy comparison.
```

The cross-parameter judge evaluates comparability at each selection call:

```
Comparability: rationale states all files use identical settings
(kspacing=0.15, PBE, methfessel-paxton, conv_thr=1e-6) except ecutwfc.
No contradiction across rationales. All same 24-atom CO/Pt(111) system.
```

When calculations in the slab-thickness test fail to converge, this requirement constrains the permitted modifications. The convergence analysis states:

```
I do NOT recommend changing kspacing because: (a) it was converged on a
representative system, (b) changing it would invalidate the comparison
set consistency, (c) the primary issue is mixing/SCF stability, not
k-points. [...] nspin: 2 (spin polarization): This is a major change
that would require re-running ALL calculations including the converged
z4 and z5, invalidating comparison set consistency. [...] This approach
maintains comparison set consistency by using identical DFT parameters
(ecutwfc=80 Ry, kspacing=0.15 Å-1) across all calculations, only
adjusting the SCF convergence control parameters (mixing,
diagonalization) which do not affect the final converged energy (they
only affect the path to convergence).
```

After the report judge identifies that the provisional 80 Ry cutoff has not been characterized, the remediation instructions specify the consequences of replacing an upstream artifact:

```
IMPORTANT: regenerating an artifact invalidates every artifact
downstream that referenced it. After fixing this artifact, you MUST
re-run every dependent calculation (production runs, ensemble
calculations, downstream extractions) using the new artifact's
result_id.
```

The extended ecutwfc test then uses the previously determined k-point spacing. The judge evaluates whether that result remains applicable under the higher cutoff:

```
The kspacing convergence (a80zq4b6) was performed at ecutwfc=80 Ry [...]
The current run is at ecutwfc=100 Ry. Is 80 Ry -> 100 Ry a "much
higher" ecutwfc mismatch? This is the key judgment. K-point sampling
convergence is essentially independent of the plane-wave cutoff --
kspacing controls reciprocal-space Brillouin zone sampling, while
ecutwfc controls the plane-wave basis completeness. These are largely
orthogonal convergence axes.
```

When the k-point test is repeated at the revised cutoff, the judge again evaluates the recorded context and justification:

```
Rationale coherence: The rationale identifies the study (regenerating
the kspacing convergence test with ecutwfc=95 Ry), explains WHY
(consistency after discovering optimal ecutwfc=95 > previously used
80 Ry), identifies the role of these values (same test values as
original kspacing test, spanning tight to coarse).
```

During the fixed-layer test, additional calculations fail to converge. After consulting the convergence agent, the DFT agent reports:

```
I rejected parameter changes that would break comparison set
consistency (degauss, kspacing, smearing type, dipole correction).
```

The revised template is applied only to the CO+slab calculations. The report judge identifies the resulting difference in self-consistent-field settings, producing the fixed-layer warning described above:

If the CO+slab batch uses hzlc6pf4 (improved SCF control) but the clean-slab batch uses 5hfafzg8 (standard SCF control), the E\_ads derived quantity would combine calculations with mismatched SCF convergence handling.

The supervisor reviews the associated canvas records and responds:

The verifier raised a WARNING on the n\_fixed\_layers report (step 65) regarding comparison-set consistency: the clean slab batch for n2/n3 used the original template (5hfafzg8) with standard SCF parameters, while the CO+slab batch used the updated template (hzlc6pf4) with improved SCF convergence control. However, after reviewing the CANVAS notes, I find this is scientifically justified: (1) The clean slab calculations already converged successfully with standard SCF parameters in step 56, (2) The improved SCF parameters were specifically needed for CO+slab systems due to charge sloshing between CO and Pt, not for clean slabs, (3) All physically meaningful DFT parameters (ecutwfc=95 Ry, kspacing=0.15, conv\_thr, smearing, pseudopotentials) are IDENTICAL across all calculations, (4) SCF convergence control parameters affect whether SCF converges, not the converged energy value itself. The E\_ads results show excellent convergence (all within 0.002 eV), confirming the approach is valid. The verifier's warning is noted but does not require regeneration.

This exchange establishes the applicable consistency criterion. Parameters that affect the physical result must remain identical across calculations, whereas self-consistent-field control parameters may differ when they affect convergence behavior but not the converged energy. The judge records each deviation, and the supervisor may override a warning when supported by evidence from the canvas. Both the warning and its justification remain in the execution record.

The same discipline carries into the production rounds. Regenerating the production inputs, the worker records:

**\*\*Maintained comparison set consistency\*\*** by using IDENTICAL DFT parameters across all 8 calculations to ensure error cancellation in formation energy differences.

This study activates every verification tier on a complex workflow. Each tier detects a distinct class of failure: the convergence-selection tool prevents an unconverged cutoff from being registered as a recommendation; the math and extraction checks prevent unsupported values from acquiring valid provenance; the call-time judge rejects uncharacterized or unsourced parameters before calculations are submitted; the convergence LLM agent diagnoses execution failures using input, output, and error files; and the report judge identifies provenance errors that are not detectable at individual tool calls and re-judges, with the whole study in view, borderline calls that argued their way past a gate. The framework permits exploratory calculations and deferred reporting, as shown by the provisional 80 Ry cutoff and the vacuum spacing held through the structural tests, but prevents uncharacterized values from supporting verified claims: the first production round ran on the held vacuum, and its report could not pass until a convergence test characterized the value. It also enforces consistency across calculations used in the adsorption-energy difference. This requirement is specified in the agent prompts and evaluated by the

**Table 4:** The  $\Delta BE$  calculated by DREAMS agrees well with our human expert and literature for both PBE and LDA functionals. Calculations are performed on a  $2 \times 2$  supercell at  $1/4$  monolayer ( $\theta = 1/4$  ML) coverage.  $\Delta BE = E_{ads\_ontop} - E_{ads\_fcc}$ .

| Supercell    | $\theta$ (ML) | XC  | $\Delta BE$ (eV) |              |                        |
|--------------|---------------|-----|------------------|--------------|------------------------|
|              |               |     | Agent Team       | Human Expert | Literature             |
| $2 \times 2$ | $1/4$         | PBE | 0.1078           | 0.108        | 0.10-0.24 [68, 71, 72] |
| $2 \times 2$ | $1/4$         | LDA | 0.318            | 0.320        | 0.32-0.45 [68, 76]     |


cross-parameter judge at each call. It constrains parameter modifications during convergence recovery and requires dependent tests to be repeated when the cutoff changes. The only accepted difference concerns SCF control parameters that affect convergence behavior but not the converged energy; the warning and supporting justification remain in the execution record. These checks establish that the final reported value is supported by verified calculations, parameters, and provenance.

### Quantifying the uncertainty of exchange-correlation functionals

Upon demonstrating DREAMS’ capability of reproducing human expert level precision DFT simulation results on the CO/Pt(111) puzzle challenge, the scope is extended. A natural third challenge from these results is the sensitivity of the adsorption site preference to the specific choice of exchange-correlation functional at the same level of theory (GGA). To quantify this uncertainty, DREAMS is asked to perform Bayesian ensemble sampling with van der Waals correction (BEEF-vdW) [67] and the distribution of possible  $\Delta BE$ s can be found in Fig. 7 (c). For both the human expert’s result and DREAMS’ result, 0 eV, is more than ten standard deviations outside of the distribution of  $\Delta BE$ , indicating that, even when considering uncertainty, the FCC site remains more energetically favorable than the ontop site at the GGA level. These results are consistent with our earlier findings obtained using the PBE and LDA functionals, as well as with previous literature [68]. We therefore conclude that the CO adsorption behavior on the Pt(111) surface is unlikely to fully match experimental observations under the current GGA-level approximations.

As for the discrepancies, some possible factors are identified. First, the convergence threshold in SCF calculations used by DREAMS is higher ( $1 \times 10^{-6}$  Ry) than that used by the human expert ( $1 \times 10^{-7}$  Ry), despite the lower threshold being recommended in Quantum ESPRESSO for structure relaxation. Second, although both DREAMS and the human expert employ the same smearing method, DREAMS sets a higher smearing width of 0.02 Ry compared to the human expert’s setting of 0.01 Ry. This larger smearing width introduces further deviations in the computed adsorption energies. Together, these differences could explain the minor shift observed in DREAMS’ results relative to the human expert calculations.

With the adsorbed configurations explored by DREAMS, we can provide some insight into the CO/Pt(111) puzzle. After the autonomous DREAMS workflow was



**Fig. 7:** (a) Representative possible structures that DREAMS can generate, illustrating variations in adsorption site, adsorbate orientation, and slab thickness. Panel (a) illustrates structure-generation capability only; convergence is established by the dedicated tests described above. (b) Configurations used for the BEEF ensemble analysis: FCC adsorption, on-top adsorption, clean slab, and isolated CO molecule. (c) Distribution of  $\Delta BE$ ,  $E_{ads\_ontop} - E_{ads\_fcc}$ , values from BEEF ensemble analysis: human-expert calculations yield a mean of  $-0.13$  eV ( $\sigma = 0.01$  eV), while DREAMS produces a mean of  $-0.12$  eV ( $\sigma = 0.01$  eV). Both methods consistently predict that the FCC-site adsorption energy is lower than the on-top-site adsorption energy, indicating that FCC adsorption is more favorable at the GGA level.

completed, we performed Bader charge and C–O bond-length analyses on its generated structures and electronic densities to further validate the predicted adsorption-site trend and investigate its electronic origin. The results are shown in Table 5. Recall that DREAMS predicts that the FCC site is more energetically favorable than ontop sites under all three functionals used in this work. We observe that for all adsorbed CO, the C–O bond length, i.e.  $d_{C-O}$ , increases compared to the pristine CO molecule. The reason for the CO bond length increase lies in the electron transfer from the Pt surface to the CO molecule. Upon adsorption, free electrons in Pt metal partially transfer to the C–O antibonding  $\pi^*$  orbitals, effectively decreasing the bond order of the C–O bond, leading to a bond length increase. In all three functionals, the FCC site probes  $0.15$ – $0.18$   $e^-$  more electron transfer than the ontop site (charge density plots are available in Supplementary Information S-6), and consequently, the C–O bond length is always longer on the FCC site. The larger degree of electron transfer indicates enhanced CO–Pt interaction via electron sharing, which explains why CO adsorption is more favored on the FCC site [72]. We would like to point out that

**Table 5:** Bond length and charge transfer analysis at different CO adsorption sites on the Pt(111) surface. The CO molecule in vacuum is also listed for reference.

| Functional | Ontop         |                    | FCC           |                    | CO molecule<br>$d_{C-O}$ (Å) |
|------------|---------------|--------------------|---------------|--------------------|------------------------------|
|            | $d_{C-O}$ (Å) | Charge transferred | $d_{C-O}$ (Å) | Charge transferred |                              |
| LDA        | 1.15          | 0.02 $e^-$         | 1.18          | 0.18 $e^-$         | 1.13                         |
| PBE        | 1.16          | 0.03 $e^-$         | 1.19          | 0.21 $e^-$         | 1.14                         |
| BEEF-vdW   | 1.15          | 0.08 $e^-$         | 1.18          | 0.23 $e^-$         | 1.13                         |

the electron transfer from transition metal surface to CO has actually been studied as early as 1964 [77]. The mechanism observed in this work, namely the electron donation from Pt to CO, has been proposed as the back donation mechanism [72, 77]. Such a mechanism has been argued to favor the FCC site over the ontop site [72, 77], which agrees with DREAMS’ finding. This agreement demonstrates that the structures and electronic densities generated through the DREAMS workflow can support further mechanistic investigation without fine-tuning the underlying LLM or introducing challenge-specific LLM agents.

### Quantitative comparison with other LLM agentic frameworks

We empirically benchmark our framework against two representative agentic systems for scientific workflows, MDCrow [22] and ChemGraph [23], on the Sol27LC lattice constant challenge and the CO/Pt(111) adsorption challenge, and we additionally benchmark it against itself. DREAMS denotes the framework without the safety guard, retaining its canvas and orchestration, while DREAMS\_safe denotes the full system with every guard layer active. The systems use different architectures: MDCrow is a LangChain-based, single-agent ReAct framework; ChemGraph uses LangGraph-based multi-agent task decomposition; DREAMS uses hierarchical task assignment, specialized worker agents, a persistent canvas, and supervisor-directed plan updates; and DREAMS\_safe adds the safety guard to DREAMS. These design elements are architectural choices rather than binary indicators of system quality. Because the baseline systems differ from DREAMS in several respects simultaneously, comparisons with MDCrow and ChemGraph evaluate overall end-to-end behavior rather than isolate the causal contribution of any individual architectural choice. By contrast, the comparison between DREAMS and DREAMS\_safe holds the canvas and orchestration fixed and directly evaluates the added effect and cost of the safety guard. Both MDCrow and ChemGraph were provided with the same DFT toolset as DREAMS, excluding the canvas tools, and their prompts were adjusted only to match the DFT context while maintaining their original design philosophy. For example, an original instruction such as “You are a molecular dynamics expert” was modified to “You are a DFT expert”. Similarly, if step-by-step MD instructions are given, they are replaced with equivalent DFT procedures. The modified prompts are provided in Supplementary Information S-10. All four systems share the same backend LLM, Claude Sonnet 4.5 [73]; the judges of DREAMS\_safe are backed by

Claude Opus 4.8 [74]. At least 3 runs were conducted in each case, and the number shown in Table 6 is an average across those runs.

**Table 6:** Benchmarking results of different agentic systems on Sol27LC and CO/Pt(111) challenges. DREAMS denotes our framework without the safety guard (canvas and orchestration only); DREAMS\_safe denotes the full system with all guard layers active. To assess accuracy, robustness, and cost, we employ five quantitative metrics: (1) Percentage of essential steps executed (reflecting the framework’s ability to attempt the complete workflow). (2) Percentage of essential steps succeeded (reflecting correct execution and recovery from errors). (3) Mean Absolute Percentage Error (MAPE) of the final numerical results. (4) Standard deviation of results across multiple independent runs. (5) Total LLM token usage, calculated as the sum of input and output tokens and averaged across the independent runs.

| Challenges           | Agent       | Steps Executed | Steps Succeeded | MAPE  | STD      | Token Usage |
|----------------------|-------------|----------------|-----------------|-------|----------|-------------|
| Sol27LC, BCC Lithium | MDCrow      | 100%           | 100%            | 0.55% | 0 Å      | 500346      |
|                      | ChemGraph   | 100%           | 100%            | 0.55% | 0 Å      | 269725      |
|                      | DREAMS      | 100%           | 100%            | 0.55% | 0 Å      | 469647      |
|                      | DREAMS_safe | 100%           | 100%            | 0.55% | 0 Å      | 4239831     |
| CO/Pt(111), PBE      | MDCrow      | 39%            | 23%             | —     | —        | —           |
|                      | ChemGraph   | 73%            | 46%             | 389%  | 0.390 eV | 3040097     |
|                      | DREAMS      | 100%           | 81%             | 0.2%  | 0.001 eV | 1814315     |
|                      | DREAMS_safe | 100%           | 100%            | 0.1%  | 0.001 eV | 23920487    |

For the Sol27LC benchmark, all four systems complete the required steps and produce comparable lattice constants, indicating that each can execute structured, short-horizon workflows. Their mean errors remain below 1%, consistent with expert DFT calculations. Performance differs substantially for the more complex CO/Pt(111) adsorption problem. MDCrow fails to complete the workflow in all trials because of context loss and incorrect tool use. ChemGraph improves task decomposition but lacks effective cross-agent communication, leading to repeated calculations, overwritten intermediate files, and invalid results. DREAMS uses the shared canvas and hierarchical communication structure (Fig. 9) to execute all 26 essential steps and produces a result within 0.2% of the human-expert value, although only 81% of those essential steps succeed. The incorrectly executed steps are the structural convergence tests on slab thickness and vacuum spacing: the underlying calculations do not converge, yet the agent selects a setting it considers reasonable and proceeds. The final agreement results from partial cancellation of systematic errors among the slab, adsorbate, and combined-system calculations. The numerical result therefore does not indicate that the workflow was executed correctly. DREAMS\_safe addresses this limitation by requiring each essential step to pass the corresponding verification checks before the workflow proceeds. All 26 essential steps succeed, and

**Table 7:** Token-cost breakdown of DREAMS\_safe and DREAMS (no safety guard) on the two challenges, by call category, for one representative run of each system. “Agent, canvas tool call” covers canvas inspection, read, and write operations; “Judge, outside report span” is the call-time (in-tool) judge cost; “Judge, inside report span” is the report-time judge cost.  $n$  is the number of LLM calls in each category.

| Category                    | DREAMS_safe |           |         | DREAMS |         |        |
|-----------------------------|-------------|-----------|---------|--------|---------|--------|
|                             | $n$         | Input     | Output  | $n$    | Input   | Output |
| Agent, non-canvas tool call | 41          | 719,964   | 23,845  | 27     | 151,538 | 7,126  |
| Agent, canvas tool call     | 127         | 2,136,289 | 11,659  | 29     | 156,406 | 2,622  |
| Judge, outside report span  | 69          | 397,734   | 41,019  | 0      | 0       | 0      |
| Judge, inside report span   | 139         | 823,407   | 85,914  | 0      | 0       | 0      |
| Agent, non-canvas tool call | 190         | 5,780,519 | 173,954 | 88     | 852,939 | 23,807 |
| Agent, canvas tool call     | 343         | 8,814,788 | 54,543  | 119    | 850,476 | 10,730 |
| Judge, outside report span  | 627         | 3,573,264 | 383,770 | 0      | 0       | 0      |
| Judge, inside report span   | 817         | 4,647,387 | 492,262 | 0      | 0       | 0      |

the final result achieves 0.1% MAPE with the same spread and a fully verified provenance chain. A more detailed explanation of the MDCrow and ChemGraph failure cases is provided in Supplementary Information S-9.

Token usage is summarized in Table 6 and separated by call category in Table 7. “Judge, outside report span” denotes call-time judge usage, whereas “Judge, inside report span” denotes report-time judge usage. For CO/Pt(111), unguarded DREAMS uses approximately 40% fewer total tokens than ChemGraph and obtains a final value close to the expert reference. Although the numerical answer agrees with the reference, only 81% of the essential steps succeed, and the agreement arises partly from error cancellation. The comparison therefore illustrates a tradeoff between cost and evidentiary rigor. The unguarded mode may be useful for rapid exploration, where a likely answer guides subsequent work and will be independently checked, whereas full verification is appropriate when the result must support a final scientific claim with an auditable provenance chain. Neither configuration is universally preferable; the appropriate choice depends on the confidence required for the intended use. Approximately half of the unguarded agent’s input tokens are associated with canvas operations that maintain information across the long-horizon workflow. Full verification substantially increases cost. DREAMS\_safe uses approximately 13x more input tokens than DREAMS for both benchmarks and 17x to 32x more output tokens. Judge calls account for approximately one third of the input tokens and four fifths of the output tokens. The similar input-cost ratios across the two benchmarks suggest that verification cost increases primarily with workflow activity rather than task complexity. The cost can be reduced by modifying the verification configuration. Judge outputs are dominated by explanations for passing verdicts; reporting explanations only for failed checks would reduce output usage but provide less

information for auditing successful checks. Call-time and report-time judge layers can also be disabled independently. The present study evaluates the maximum-verification configuration, in which all layers are active and all values are checked. Other configurations can balance verification coverage against token cost.

### The tuned judge rules transfer across judge models

We evaluate the reliability of the verification judges using the procedure described in the Methods. The judge rules are adversarially refined using a standalone scenario harness and evaluated on 145 scenarios. Each scenario represents a realistic tool call derived from artifacts generated in actual runs and is either valid or contains one injected defect. Each verdict is determined by majority vote across three judge calls, and a warning is counted as detection of a should-fail scenario. Fig. 8 presents the resulting confusion matrices for five Anthropic models: Claude Fable 5, Claude Opus 4.8, Claude Sonnet 5, Claude Sonnet 4.6, and Claude Sonnet 4.5.


Three findings emerge. First, all five models detect most defective scenarios and accept most valid scenarios, with total errors ranging from 4 to 17 out of 145. This consistency indicates that the rule set, rather than a single judge model, provides most of the discriminatory performance. Second, the production judge, Claude Opus 4.8, produces 2 false positives and 5 false negatives, placing it among the best-performing models and supporting its selection. Third, performance improves with model generation. Claude Fable 5 produces the fewest errors, with 1 false positive and 3 false negatives, indicating that the same rule set can benefit from improvements in the underlying model without further modification. The models also exhibit different trade-offs between false positives and false negatives. Claude Sonnet 4.6, for example, misses only one defective scenario but produces ten false positives. Judge-model selection therefore affects the balance between detection sensitivity and unnecessary intervention.

## Discussion

DREAMS demonstrates the use of large language model (LLM) agents for autonomous computational materials research. The framework combines hierarchical task planning, domain-specific tools, shared memory, and multi-tier verification from tool execution to final reporting. DREAMS performs periodic solid-state DFT workflows at an enhanced L2 (L2+) automation level that approaches L3. Its architecture addresses limitations of existing LLM agents, including hallucination, context loss, inadequate error recovery, and the absence of verifiable numerical results.

Evaluation on the Sol27LC benchmark establishes the accuracy and reliability of DREAMS for lattice-constant calculations. For CO/Pt(111), DREAMS performs adsorption-energy calculations across multiple configurations and reproduces the expected site preference, demonstrating its applicability to rigorous materials simulations. For exchange-correlation functional uncertainty, DREAMS uses Bayesian ensemble sampling with BEEF-vdW to quantify functional-dependent variation. The analysis shows that the FCC-site preference remains robust across the ensemble and relates this preference to the underlying electronic structure.





**Fig. 8:** Judge confusion matrices on the 145-scenario adversarial suite for five Anthropic models. Rows are intended verdicts (should-pass, should-fail); columns are the majority verdict over three judge calls (pass, warning, fail). A warning counts as catching a should-fail scenario, while any non-pass verdict on a should-pass scenario counts as a false positive; per-model false-positive (FP) and false-negative (FN) counts are indicated above each panel.

Two additional evaluations characterize the framework. In the four-system comparison, the unguarded agent produces a nearly correct final value despite an essential-step success rate of only 81%, because systematic errors partially cancel in the calculated energy difference. The guarded system produces a comparable result with all essential steps verified but requires approximately 13 times more input tokens. This increase reflects the cost of performing and recording the verification steps. The present study evaluates the strictest configuration, with all verification layers enabled. The cost of that strictness is visible in the CO/Pt(111) study: the report judge demanded characterization of a vacuum value that was already physically adequate, at the price of one additional convergence test and a second production round; the test in turn confirmed the held value was never under-set and identified a cheaper converged setting. Such a test may seem redundant at this scale, but the same check becomes critical in massive production campaigns consisting of thousands of calculations, where an uncharacterized setting propagates into every result and an over-set one wastes compute in every calculation. Individual layers can be disabled to reduce cost at the expense of verification coverage. Identifying the optimal verification configuration and judge model remains future work. Evaluation across five judge models further shows that the adversarially refined rules transfer between models and that judge performance improves with newer model generations. The verification

system can therefore benefit from improvements in the underlying models without revising the rule set.

Under the taxonomy in Table 1, DREAMS approaches L3 automation. L3 denotes autonomous search within a design space defined by the scientific objective, whereas autonomous construction or restructuring of that design space corresponds to L4 or higher. In the present studies, the candidate crystal structures, adsorption sites and orientations, and exchange-correlation functionals are specified in advance. DREAMS then executes the search, submits and monitors calculations, diagnoses failures, and revises its plan autonomously. Benchmark execution is separated from evaluation: the task specifications contain no target reference values, and the calculated values are compared with established references only after workflow completion. Each reported quantity is computed anew through structure construction, convergence testing, HPC execution, error recovery, and post-processing, with its computational origin recorded in the provenance chain. The provenance record establishes that the reported values derive from executed calculations rather than a supplied reference; it cannot establish that knowledge encountered during model training did not influence workflow choices. These benchmarks therefore evaluate autonomous computational reproduction under controlled conditions, not the generation of previously unknown scientific knowledge. Evaluation on problems without known solutions is required to establish full L3 discovery capability. Future work will extend DREAMS to open-ended problems with dynamically constructed design spaces.

Key innovations in our approach include: (1) a multi-tier safety guard that verifies numerical values deterministically and semantically at tool-call and report time; (2) deterministic tools that constrain agent actions while preserving task flexibility; (3) a planning supervisor that decomposes research objectives into validated execution steps; (4) a shared canvas for communication and state management across agents; and (5) systematic convergence procedures for computational validation. The safety guard is independent of DFT because its provenance registration, verification gates, and judge evaluations do not depend on a specific computational method. It can therefore be applied to other computational workflows that require auditable numerical results.

Several limitations remain. The judge rules were adversarially refined using DFT-specific scenarios and may require task-specific adaptation for other domains. Developing rules that generalize across multiple scientific tasks is an important direction for future work. The evaluated judge panel includes only Anthropic models; broader evaluation across model families is also required. In addition, the list of sensitive parameters is defined by domain experts. The guard enforces this knowledge but does not identify sensitive parameters autonomously. Future work should compare the convergence agent with documentation-retrieval and retrieval-augmented baselines on identical calculation failures, and evaluate its generalization beyond the established workflows supported by the current toolset.

This work provides a general framework for applying LLM agents to computational materials science and chemistry. Future studies may extend DREAMS to additional scientific domains and incorporate specialized agents for other computational

methods. For CO/Pt(111), the present results indicate that GGA-level theory alone is insufficient to resolve the adsorption-site discrepancy. Further analysis should include thermal corrections, coverage effects, higher-rung exchange-correlation functionals, including meta-GGA and hybrid functionals, and DFT+U where electron localization is relevant. More broadly, agentic workflows may support scientific discovery by automating computational studies, identifying missing physical effects, and generating mechanistic interpretations.

## Methods

DREAMS (DFT based Research Engine for Agentic Materials Simulation) is implemented as a hierarchical multi-agent system built on the LangGraph [78] framework. The planning supervisor and the worker LLM agents—a DFT LLM agent and an HPC LLM agent—are backed by Claude Sonnet 4.5 [73] at temperature zero; the LLM judges of the safety-guard system operate outside the agent loop and are backed by Claude Opus 4.8 [74]. Each worker LLM agent is built as a ReAct agent [79] and incorporates various tools for execution; reasoning lets the workers adapt dynamically to changing task requirements and handle exceptions, while actions let them interface with built-in functions as tools. Their system prompts are structured into sections (*Role*), (*Objective*), (*Instruction*), (*Requirement*), and the complete operational prompt set is provided in Supplementary Information S-7. DREAMS’ tools are deterministic at their core: each consists of a deterministic body—structure construction, script generation, file parsing, analysis—wrapped in the safety-guard interface described below, which validates the call before the body runs and registers its output with full provenance afterward. The single deliberate exception is the convergence agent, an LLM embedded inside a tool to diagnose failed calculations, described in its own subsection. Rigorous validation checks and informative error messages in every tool support agent self-correction, and the canvas is shared among all agents, tools, and the human user to mitigate hallucination. This hierarchical approach is chosen for its flexibility, its ability to manage complex interactions among multiple LLM agents, and its extensibility.

### Dynamic Planning

The planning supervisor LLM agent plays a central role in orchestrating the overall workflow. Unlike simple task schedulers, the supervisor decomposes complex research objectives into manageable subtasks tailored to each worker LLM agent’s capabilities. The workflow state carries the user objective, the current plan, a typed record of executed steps, the canvas, and the artifact registry, so that plans, memory, and provenance are checkpointed together and a study can be resumed or rewound. Each plan step names the worker agent responsible for it, the tool whose successful call the step requires, and, for steps that consolidate findings, the quantities the resulting report must deliver; the first two requirements prevent an agent from producing a result without a tool, and the third prevents it from quietly not reporting one. Rather than rewriting the plan wholesale after every step, the supervisor edits it through three validated operations that insert, modify, or delete steps, and every edit

is checked before acceptance: the assigned agent must exist, the required tool must belong to that agent's toolset, and report steps must declare their required quantities. After a worker completes a step, the framework verifies that the required tool was actually called; for report steps it additionally audits, deterministically, that every required quantity appears in the report and matches its cited artifact, and a report failing this audit is deleted and the step retried with the failure explained. Once a worker has generated a report within a step, its toolset is restricted to bookkeeping operations until the step ends, so that not-yet-verified numbers cannot flow into further work. When a report passes verification, the span of steps that produced it is archived and replaced in the active history by a single summary record, keeping the supervisor's context bounded without destroying the trail. Once the objective is fully addressed, the supervisor writes the final response and terminates the workflow. We refer to this planning as *dynamic* because the supervisor does not commit to a fixed plan at the outset; instead, it revises the plan whenever a step completes or an unexpected situation arises, taking a holistic view of the objective, the current plan, and the history of completed steps—in contrast to frameworks that generate a single plan up front and leave any deviation to be handled locally by the worker that happens to encounter it. Plan creation and revision remain solely the responsibility of the planning supervisor; worker LLM agents do not access the plan-editing tools or modify the plan. Workers may provide domain-specific input when explicitly consulted, and their execution results, failures, and reports inform subsequent planning decisions, but the supervisor alone determines whether and how this information is translated into plan edits.

## DFT LLM Agent

The DFT LLM Agent functions as a specialized worker LLM agent with computational chemistry expertise focused on rigorous periodic solid-state density functional theory (DFT) calculations. This LLM agent is responsible for managing the scientific aspects of electronic structure calculations, including the construction of atomistic models, optimization of computational parameters, and analysis of calculation results. The DFT LLM agent's core functions include: (1) generating atomistic structures such as bulk crystals, surfaces, and adsorbate configurations; (2) determining suitable DFT convergence parameters through systematic convergence testing; (3) preparing input files for quantum chemistry software; and (4) analyzing outputs and extracting physical properties, including lattice constants and adsorption energies. To perform these tasks, the DFT LLM agent is equipped with a suite of tools, each designed for a specific purpose. The LLM agent selects tools based on their documented capabilities. Convergence-parameter selection supports two modes: one sweeps a calculation parameter and compares each output directly against a most-converged reference, while the other compares derived quantities, enabling structural convergence tests, such as slab thickness, whose criterion is a derived observable rather than a raw total energy. Further details of these tools are provided in Supplementary Information S-8.

The tools of the DFT LLM agent are interfaced through established frameworks such as the Atomic Simulation Environment (ASE) [75], preventing unsafe operations like direct file manipulation, which can otherwise lead to invalid structures or runtime failures. Each tool maintains explicit input-output mappings and implements comprehensive error handling. For example, the adsorption energy tool requires exactly three DFT output files (clean slab, isolated adsorbate, and adsorbate-on-slab) and applies standard thermodynamic formulas to compute results. This architecture prevents unsupported file operations and enforces the prescribed input–output relationships, thereby reducing the opportunities for hallucinated parameters, inconsistent inputs, and incorrect arithmetic. To further reduce redundancy and the risk of inconsistencies, we develop deterministic batch tools that generate Quantum ESPRESSO (QE) [64] input files from a single complete input. The agent supplies the structure, pseudopotentials, and full DFT parameterization to an ASE-backed tool, which renders a complete QE input file in QE’s structured plain-text syntax. This file is termed a *template* only because downstream batch tools inherit its settings; it is neither a Python script nor a pre-supplied fill-in-the-blank form. For convergence testing, the batch tool programmatically substitutes only the swept parameter while preserving the remainder of the input. The deterministic file-generation work and total file output remain  $O(N)$ , but LLM-mediated specification is reduced from  $N$  complete parameterizations to one. This lowers LLM API calls, LLM output tokens, and response latency while reducing the risk of agent-introduced inconsistencies. For EOS calculations, the batch tool uniformly scales the cell and generates five complete QE input files using the converged production settings. Agreement with expert reference results therefore reflects the complete workflow—structure construction, parameter selection and empirical convergence, QE execution, output extraction, and EOS fitting—rather than input-template completion alone.

### Atomistic Structure Generation

In the benchmarking experiment on the CO/Pt(111) system, we utilize the AutoCat library [14] as a computational framework for generating and analyzing adsorption structures. Compared to simply letting the LLM decide the atomistic structures, using AutoCat to control the structural degrees of freedom ensures that the structures to be simulated are physically meaningful. For example, DREAMS can create structures with different adsorbate sites, surface slab sizes, vacuum spacing, adsorbate orientations, coverage patterns, and adsorption heights, and we are confident that the structures are physically meaningful, as illustrated in Fig. 7 (a).

Implementing domain knowledge through AutoCat also avoids redundant configurations. For example, for Pt with FCC crystal structure, the (111) surface has only four distinct adsorption sites, namely the ontop, bridge, FCC, and HCP sites. However, if such symmetry constraints are not known in advance, a grid search with a 0.1 Å spacing would result in 121 possible adsorption sites, which is almost 30 times the computational cost.

In conclusion, AutoCat abstracts these low-level configurational operations behind a structured API, allowing LLM agents to invoke high-level actions such as “generate FCC adsorption structures with all relevant CO orientations” without having

to reason about individual atomic coordinates or manually define geometry constraints. The integration of such carefully designed tools thus enables reliable and efficient exploration of a chemically rich configuration space that would otherwise be prone to failure or hallucination if handled entirely at the LLM level.

This structured interface creates a deliberate tradeoff between reliability and generality. The present tools support bulk structures, ideal crystalline slabs, variation of the exposed slab parameters, and adsorption configurations constructed from sites identified by AutoCat. They do not constitute a universal geometry generator: surface reconstructions, steps, defects, arbitrary multicomponent terminations, and protocols outside the available tool schemas are not systematically supported in this work. Such capabilities require implementing and testing an additional domain tool and, when safety guards are active, specifying the associated sensitive parameters and verification rules. The supervisor–worker orchestration does not need to be redesigned, but the executable scientific domain remains bounded by the tools that have been integrated and validated.

### Convergence LLM Agent

The convergence LLM agent is implemented as a large language model (LLM) embedded within a tool. To use this tool, the DFT LLM agent must first specify which job has encountered a convergence issue. The tool then locates the relevant files for that job and, for each file, prompts the LLM with the following query: Here is the content of the file <content>, please give me suggestions on how to fix the convergence issue. The LLM is instructed to provide a structured and concise response to the formatted question. In this setup, most operational tasks, such as parameter parsing and preprocessing, are handled deterministically by code. This allows the LLM to focus exclusively on decision-making. By decoupling convergence diagnosis from the broader LLM-agent communication graph, the dedicated agent can analyze the relevant input, output, and error files without adding them to the DFT worker’s already extensive workflow context. In our benchmark runs, this focused context reduced the hallucinated and erratic recovery actions observed when the worker handled these files directly, although component-level recovery rates and variability across LLM APIs were not systematically evaluated. The convergence LLM agent provides suggestions for parameter settings and script generation, enabling the DFT LLM agent to self-correct and conduct updated calculations. These suggestions are generated from the base LLM’s pretrained knowledge conditioned on the supplied job files; no parameter-specific remedial heuristics, decision trees, or lookup tables are embedded in the prompt or tool. When a parameter such as `electron_maxstep` is changed, the updated input is submitted as a new calculation with `restart_mode = from_scratch`, and no electronic state from the preceding iterations is reused. Although QE supports checkpoint-based continuation, it is not enabled here because retaining the required restart data incurs substantial disk overhead. Representative suggestions are shown in Table 3.

## HPC LLM Agent


The HPC LLM agent is a resource- and job-orchestration component. It reads the Quantum ESPRESSO inputs and job list, proposes resource requests from the site-specific cluster configuration, records those suggestions, submits calculations through pysqa to SLURM, monitors their status, and retrieves the resulting files. It does not execute the numerical kernels or modify Quantum ESPRESSO's parallelization strategy. Computational efficiency, parallel scaling, and accelerator support therefore depend on the underlying Quantum ESPRESSO build and hardware and are not evaluated as performance properties of the agent.

The present implementation was tested on a SLURM-based [80] CPU cluster with AMD 9654 processors. Deployment on another system requires site-specific scheduler, partition, resource, software-module, and submission-script configuration. Non-SLURM schedulers and GPU execution were not evaluated in this work. The submission interface uses the Python Simple Queuing System Adapter (pysqa) [12]; the HPC-agent tools and their arguments are documented in Supplementary Information S-8.

## Memory and communications among human user, LLM agents, and tools

Effective communication is essential for ensuring reliable task execution and the accuracy of final results. To achieve this, DREAMS establishes a structured communication flow among the user, supervisor agent, worker agents, and tools, which is illustrated in Fig. 9. When the user invokes DREAMS with a request, the request is (1) sent to the supervisor. The supervisor then (2) checks the canvas for existing information, (3) generates or updates the plan, determines which worker agent should act next, and (4) issues the command “You need to do xxx.” Along with this instruction, the overall objective and summaries from previous steps are (5) sent to the worker agent. Upon receiving the task, the worker agent (6) retrieves relevant information from the canvas and (7) invokes the appropriate tools. Each tool can (8) access and update the canvas as needed. After completing the assigned task, the worker agent (6) writes key results to the canvas and (9) adds a summary of the current step. Steps (2) through (9) repeat until the supervisor (1) provides a final response to the user. During the process, the human user can (10) read the canvas to monitor progress. Although not yet activated, the user could also write information to the canvas to share additional context with agents and tools.

Canvas stores all variables in their native format, reducing errors and preventing hallucinations that may arise from converting data to and from text, and it is organized into three stores. *Notes* are free-form working documents for the agents' day-to-day bookkeeping; they are version-controlled and append-only: editing a document creates a new version in its family, and existing content is never destroyed. *Artifacts* form the append-only provenance registry used by the safety guard: every registered



**Fig. 9:** Communication flow among user, agents, tools, and the canvas in DREAMS. The user (1) sends a request to the supervisor, who (2) checks the canvas, (3) forms a plan, and (4–5) assigns a task—with the instruction, overall goal, and summaries—to a worker agent. The worker (6) reads from the canvas, (7) calls tools, which can also (8) access or update the canvas. After completing the task, the worker (6) records results and (9) summaries for the next cycle. Steps repeat until the supervisor (1) replies to the user, who can (10) monitor progress. Canvas serves as a centralized, structured memory for all data in native format, supporting safe read/write operations, access control, and transparent logging to ensure consistent and reproducible communication.

tool output receives an immutable eight-character result identifier, and each artifact records the producing tool, the value, the full argument list, the per-parameter reasons, the per-parameter sources, and the declared context. A reference may name an artifact's output (a bare identifier) or a specific input parameter of a past call (a dotted identifier of the form `<id>.<parameter>`), and referenced values are verified deterministically against the registry for both existence and equality. *Reports* are structured scientific findings; once verified, a report becomes immutable and can no longer be edited or overwritten. The three main canvas functions—inspection, reading, and writing—are exposed to the LLM agents through specialized tools, while tools access the underlying stores directly through native operations. For inspection, no arguments are required and the canvas returns the list of available keys. For reading, a valid key must be provided; if the key is invalid, the canvas suggests performing an inspection to locate the correct key before attempting to read again. For writing, the agent must supply a descriptive key together with the object to be stored,



and overwriting an existing key requires explicit confirmation, preventing accidental data loss. Framework-owned key spaces—judge feedback and archived history—are readable but not writable by agents, and any attempt to violate an access constraint returns an informative warning to guide corrective action. All updates to the canvas are logged and made visible to human users for transparency and to support post-processing. Additionally, a serialized object (a pickle file) is created to capture the current state of the canvas and the artifact registry, enabling session resumption and ensuring data availability for downstream analysis. Persistence across runs is therefore optional and under user control: if the working directory already contains this pickle file, the canvas state is restored, otherwise an empty canvas is created. We use pickle rather than JSON because the canvas stores arbitrary native objects, including custom Python classes, that JSON cannot serialize; the human-readable log noted above is generated separately for transparency. In the framework comparison, the baseline systems (MDCrow and ChemGraph) were not equipped with an equivalent canvas, as it is one of the novel components of DREAMS and requires systematic integration.

A full snapshot of the canvas from the Sol27LC challenge is provided in Supplementary Information S-2. The snapshot illustrates how intermediate structures, parameters, and calculation results remain available across successive workflow stages.

## Safety guard

Every guarded tool call must supply, in addition to its scientific arguments, a declared context (one to two sentences naming the study or exploration the call belongs to and why the tool is being called) and a per-parameter rationale covering, for each parameter, the role it plays in the study described by the context and why this specific value was chosen. For parameters designated as sensitive—a human-defined list, specified in the configuration, of the parameters known to control result validity in the domain—the agent must additionally supply the reference of the registered result the value came from. Only tool outputs become artifacts, and an agent that has not registered a value has no reference to hand to the next tool: a forgotten registration stalls progress rather than corrupting it. The agent can forget, but it cannot fabricate.

Verification at tool-call time proceeds in the order described in the Results: deterministic checks first (each supplied reference must exist in the registry, and the supplied value must equal the recorded one), then the per-parameter judge, then the cross-parameter judge; the tool body runs only if all pass. The per-parameter gate constructs a provisional artifact from the fields to be registered and applies the same verification procedure used by the report-time judge. Both stages use the same code path and rule set, allowing violations to be identified before tool execution proceeds. However, the gate evaluates only the current tool call and does not recursively inspect its upstream provenance. It therefore lacks the complete dependency chain available to the report-time judge. The two stages may consequently produce different verdicts in ambiguous cases. In the CO/Pt(111) study, the call-time judge accepted the dotted vacuum reference, whereas the report-time judge rejected it after inspecting the full provenance chain. The report-time verdict determines whether a claim is accepted.

The cross-parameter gate is attached to the analysis tools whose validity depends on the joint configuration of their inputs, such as convergence-parameter selection and equation-of-state fitting. The arithmetic tool verifies each input against its cited artifact before evaluating, and rejects expressions whose output is an input passed through unchanged or trivially rescaled (for example  $x_0 + 0$  or  $x_0 \times 1.0$ ), the pattern by which a chosen number could otherwise be laundered into a computed one. The extraction tools require a unique evidence snippet and accept a value only if it appears verbatim in the recorded output of the cited call. Gate verdicts are pass, warning (the value is registered with the concern stamped in the artifact’s metadata), or fail (the tool does not run, nothing is registered, and the judge’s reasoning is returned to the agent). Guards fail open: an infrastructure error inside a guard never blocks a run; only a scientific verdict does.

A structured report consists of the study goal, the quantities sought, numerical claims—each binding a quantity name, a value, and the reference of the artifact carrying it—qualitative findings, and a narrative in which every number must be tied to a claim. Verification proceeds from each claim: the claimed value is first matched deterministically against its cited artifact; a coherence check then asks whether the cited artifact is the right kind of source for the claimed quantity, not merely a numerically coincident one; and the judge finally descends the provenance graph, applying to every parameter of every contributing call a rule selected by that parameter’s situation—sensitive and sourced, sensitive and swept, sensitive without a source, or not sensitive—with the recorded context deciding whether a choice is judged under exploration or production standards. Failures are returned as categorized issues with remediation instructions that explicitly forbid the shortcut responses: patching in a different reference, silently dropping a failing required quantity, or substituting a math-tool fabrication. Verified subtrees are cached across reports, so an already-verified intermediate result is not re-judged when later reports build on it. The full rule texts, judge guidance, and issue categories are provided in Supplementary Information S-11.

Finally, a standalone debug tool exposes the same provenance walk for investigation rather than verification: given a surprising result and a question, it walks the value-flow chain parameter by parameter and returns ranked hypotheses about which parameter or source could explain the result, explicitly labeled as hypotheses to be tested one variable at a time. We note that this machinery has repeatedly caught our own implementation mistakes during development—hard-coded file input/output, wrongly registered argument values, an input file referenced where an output was intended—suggesting that the value of mechanical verification extends beyond agent misbehavior to the framework itself.

## Judge implementation and adversarial tuning

All judges share one implementation: a Claude Opus 4.8 model constrained to a structured output schema in which the reasoning field precedes the verdict, so that a verdict cannot be emitted without first stating its grounds; malformed responses are retried. Besides the artifact under review, each judge prompt receives the producing

tool's live documentation—its docstring and the annotation of the specific argument under review—so the judge evaluates usage against the tool's actual contract.

This documentation feed was motivated by observed judge failures. Early versions of the per-parameter judge, which saw only the tool name and the artifact description, produced characteristic false positives: they flagged the reference calculation appearing in the candidate list of the convergence-selection tool as mixing the reference into the test set, when the tool requires exactly that; and they flagged a candidate list combining two sweep batches as cross-wiring, which is legitimate whenever all files share identical settings apart from the swept parameter. The judge could not know either fact without seeing the tool's documentation.

The judge rules were tuned against a standalone scenario harness rather than by inspection. Scenarios are constructed neutrally: each carries a hypothesis about the expected verdict that is used only to triage disagreements, never to shape the rules, and false passes and false failures are always measured together. An initial synthetic suite produced spurious failures for an instructive reason—its parameter rationales were far thinner than anything a real run produces—so the suite was rebuilt from the artifacts of real runs: each clean scenario mirrors a real registered artifact with its rich, study-situated rationale, and each flawed scenario is that realistic base with exactly one injected defect, drawn from a catalog including a wrong reference direction, a source of the wrong physical kind, an unjustified value magnitude, a deliberately generic rationale, a mismatch between characterization and production conditions, a wrong file role, and laundering through a dotted reference. Scenarios are grounded mostly in the CO/Pt(111) system, with a second chemistry included so that the rules are not tuned to a single system's idioms. A proposed rule change was accepted only if, on re-running the harness, all clean scenarios still passed and all defect scenarios still failed, guarding both sides of the false-positive/false-negative tension. The final rule set was evaluated on 145 scenarios by majority vote over three judge calls per scenario, counting a warning as catching a should-fail case, across five Anthropic models; the resulting confusion matrices are reported in Fig. 8. Further details of the scenario suite and adversarial-tuning procedure are provided in Supplementary Information S-12.

## DFT Benchmark

Density functional theory (DFT) is one of the most widely employed first-principles simulation methods for atomistic systems. Through solving for the ground-state electronic density and calculating the potential energy surface of a given atomistic system, DFT provides mechanistic understanding on thermodynamic phase stability, crystal structures, electronic properties, etc. Readers are recommended to refer to the listed reference for more information [81]. A number of software packages have been developed to carry out these calculations, including VASP [82], Gaussian [83], Quantum ESPRESSO (QE) [64], etc.

In this work, we use QE v7.2 for DFT calculations, which uses plane wave (PW) basis for describing the electronic wavefunction, and uses pseudopotentials (PP) to describe core electrons. We also investigate the impact of exchange-correlation functional (XC) choices on calculated structural properties, including the

lattice constant and adsorption energies. These include functionals under the local density approximation (LDA) and the generalized gradient approximation (GGA), namely Perdew–Burke–Ernzerhof (PBE) functional and the PW91 functional. All the calculations are carried out using the GBRV pseudopotentials.

The lattice constant values are obtained by fitting an equation of state (EOS) curve between DFT-calculated potential energies and volumes, and finding the lowest-energy point. These values are benchmarked using the Sol27LC dataset, which was collected in Ref. [67] and includes lattice constants of 27 elemental crystals in BCC, FCC, and diamond lattices. The elements in this dataset cover both metals and non-metals, which makes it ideal for testing DREAMS’ ability across the periodic table. The adsorption energies are calculated as:

$$E_{ads} = E_{CO-Pt(111)} - E_{CO} - E_{Pt(111)} \quad (2)$$

where  $E_{CO-Pt(111)}$ ,  $E_{CO}$  and  $E_{Pt(111)}$  are the potential energies of the Pt(111) slab with a CO molecule adsorbed on top, the single CO molecule, and the clean Pt(111) surface, respectively. The results in this work are benchmarked against results produced by Peter et. al. [68] who systematically analyzed CO adsorption on Pt(111) using a range of DFT implementations and settings. For uncertainty quantification, the Bayesian error estimation functional with van der Waals correction (BEEF-vdW) is used. We use a total of 2,000 functionals in our ensemble to sample the energy distribution. [84]

## Framework Comparison

All DFT calculations are conducted on the in-house HPC cluster with AMD 9654 CPU. The version of MDCrow benchmarked is 0.0.2, with LangChain version 0.2.12. The version of ChemGraph benchmarked is 1.0.0, with LangChain version 0.3.27 and LangGraph version 0.4.7.

## Data Availability

The Sol27LC dataset is available [here](#).

## Code Availability

Code is available on: [GitHub](#)

## Acknowledgements

This work was supported by Los Alamos National Laboratory under the grant number AWD026741 at the University of Michigan. This work was also supported by Anthropic’s AI for Science Program. The authors thanks NAIRR for providing access to the Microsoft Azure service. The author thanks Advanced Research Computing at the University of Michigan for providing computing resources.

## Author contributions

Z.W. led the development of DREAMS and drafted the manuscript. H.H. and H.Z. assisted in implementing DREAMS. C.X. assisted in benchmarking experiments. S.Z., J.J., and V.V. contributed to project conceptualization. Z.W., H.H., H.Z., C.X., S.Z., J.J., and V.V. contributed to manuscript revision and approved the final version.

## Competing Interests:

We have filed a provisional patent application on computational agents for scientific simulations.

## References

- [1] Menon, S., Lysogorskiy, Y., Knoll, A.L.M., Leimeroth, N., Poul, M., Qamar, M., Janssen, J., Mrovec, M., Rohrer, J., Albe, K., Behler, J., Drautz, R., Neugebauer, J.: From electrons to phase diagrams with machine learning potentials using pyiron based automated workflows. *npj Comput Mater* **10**, 261 (2024)
- [2] Ertekin, E., Schiller, J.A.: A combined dft/machine learning framework for materials discovery: Application to spinels and assessment of search completeness and efficiency. *ChemRxiv* (2020) <https://doi.org/10.26434/chemrxiv.13070549.v1>
- [3] Jain, A., Ong, S.P., Hautier, G., Chen, W., Richards, W.D., Dacek, S., Cholia, S., Gunter, D., Skinner, D., Ceder, G., Persson, K.A.: Commentary: The materials project: A materials genome approach to accelerating materials innovation. *APL Mater.* **1**(1), 011002 (2013)
- [4] Kavalsky, L., Hegde, V.I., Muckley, E., Johnson, M.S., Meredig, B., Viswanathan, V.: By how much can closed-loop frameworks accelerate computational materials discovery? *Digit Discov* **2**(4), 1112–1125 (2023)
- [5] Santiago, M., Sanchez-Lengeling, B., Wei, J., Venugopal, V., Skreta, M., Krishnan, N.M.A.: Ai for accelerated materials design (ai4mat-2023). In: *NeurIPS 2023 Workshop* (2023). <https://neurips.cc/virtual/2023/workshop/66541>
- [6] On-Road Automated Driving (ORAD) Committee: Taxonomy and definitions for terms related to driving automation systems for on-road motor vehicles. *SAE International*, 3016–202104 (2021)
- [7] Goldman, B., Kearnes, S., Kramer, T., Riley, P., Walters, W.P.: Defining levels of automated chemical design. *J. Med. Chem.* **65**(10), 7073–7087 (2022)
- [8] Curtarolo, S., Setyawan, W., Hart, G.L., Jahnatek, M., Chepulskii, R.V., Taylor, R.H., Wang, S., Xue, J., Yang, K., Levy, O., Mehl, M.J., Stokes, H.T., Demchenko, D.O., Morgan, D.: Aflow: An automatic framework for high-throughput materials discovery. *Comput. Mater. Sci.* **58**, 218–226 (2012)

- [9] Jain, A., Ong, S.P., Chen, W., Medasani, B., Qu, X., Kocher, M., Brafman, M., Petretto, G., Rignanese, G.-M., Hautier, G., Gunter, D., Persson, K.A.: Fireworks: A dynamic workflow system designed for high-throughput applications. *Concurrency Computat.: Pract. Exper.* **27**(17), 5037–5059 (2015)
- [10] Pizzi, G., Cepellotti, A., Sabatini, R., Marzari, N., Kozinsky, B.: Aiida: Automated interactive infrastructure and database for computational science. *Comput. Mater. Sci.* **111**, 218–230 (2016)
- [11] Mathew, K., Montoya, J.H., Faghaninia, A., Dwaraknath, S., Aykol, M., Tang, H., Chu, I.-h., Smidt, T., Bocklund, B., Horton, M., Dagdelen, J., Wood, B., Liu, Z.-K., Neaton, J., Ong, S.P., Persson, K., Jain, A.: Atomate: A high-level interface to generate, execute, and analyze comput. mater. sci. workflows. *Comput. Mater. Sci.* **139**, 140–152 (2017)
- [12] Janssen, J., Surendralal, S., Lysogorskiy, Y., Todorova, M., Hickel, T., Drautz, R., Neugebauer, J.: pyiron: An integrated development environment for computational materials science. *Comput. Mater. Sci.* **163**, 24–36 (2019)
- [13] Annevelink, E., Kurchin, R., Muckley, E., Kavalsky, L., Hegde, V.I., Sulzer, V., Zhu, S., Pu, J., Farina, D., Johnson, M., *et al.*: Automat: Automated materials discovery for electrochemical systems. *MRS Bulletin* **47**(10), 1036–1044 (2022)
- [14] Kavalsky, L., Hegde, V.I., Meredig, B., Viswanathan, V.: A multiobjective closed-loop approach towards autonomous discovery of electrocatalysts for nitrogen reduction. *Digit Discov* **3**, 999–1010 (2024)
- [15] Zimmermann, Y., Bazgir, A., Al-Feghali, A., Ansari, M., Brinson, L.C., Chiang, Y., Circi, D., Chiu, M.-H., Daelman, N., Evans, M.L., *et al.*: 34 examples of llm applications in materials science and chemistry: Towards automation, assistants, agents, and accelerated scientific discovery. *arXiv preprint arXiv:2505.03049* (2025)
- [16] MacKnight, R., Boiko, D.A., Regio, J.E., Gallegos, L.C., Neukomm, T.A., Gomes, G.: Rethinking chemical research in the age of large language models. *Nat Comput Sci* (2025)
- [17] Alampara, N., Schilling-Wilhelmi, M., Ríos-García, M., Mandal, I., Khetarpal, P., Grover, H.S., Krishnan, N.A., Jablonka, K.M.: Probing the limitations of multimodal language models for chemistry and materials research. *Nat Comput Sci* (2025)
- [18] Miret, S., Krishnan, N.A.: Enabling large language models for real-world materials discovery. *Nat Mach Intell* **7**, 991–998 (2025)
- [19] Ghafarollahi, A., Buehler, M.J.: Sciagents: Automating scientific discovery through bioinspired multi-agent intelligent graph reasoning. *Adv. Mater.* **37**,

- [20] Song, T., Luo, M., Zhang, X., Chen, L., Huang, Y., Cao, J., Zhu, Q., Liu, D., Zhang, B., Zou, G., Zhang, F., Shang, W., Jiang, J., Luo, Y.: A multiagent-driven robotic ai chemist enabling autonomous chemical research on demand. *J. Am. Chem. Soc.* **147**(15), 12534–12545 (2025)
- [21] M. Bran, A., Cox, S., Schilter, O., Baldassari, C., White, A.D., Schwaller, P.: Augmenting large language models with chemistry tools. *Nat Mach Intell* **6**, 525–535 (2024)
- [22] Campbell, Q., Cox, S., Medina, J., Watterson, B., White, A.D.: Mdcrow: Automating molecular dynamics workflows with large language models. *arXiv preprint arXiv:2502.09565* (2025)
- [23] Pham, T.D., Tanikanti, A., Keçeli, M.: Chemgraph: An agentic framework for computational chemistry workflows. *arXiv preprint arXiv:2506.06363* (2025)
- [24] Zou, Y., Cheng, A., Aldossary, A., Bai, J., Leong, S.X., Campos-Gonzalez-Angulo, J., Choi, C., Ser, C.T., Tom, G., Wang, A., Zhang, Z., Yakavets, I., Hao, H., Crebolder, C., Bernales, V., Aspuru-Guzik, A.: El agente: An autonomous agent for quantum chemistry. *arXiv preprint arXiv:2505.02484* (2025)
- [25] Ghafarollahi, A., Buehler, M.J.: ProtAgents: protein discovery via large language model multi-agent collaborations combining physics and machine learning. *Digital Discovery* **3**(7), 1389–1409 (2024) <https://doi.org/10.1039/D4DD00013G>
- [26] Ghafarollahi, A., Buehler, M.J.: Sparks: Multi-agent artificial intelligence model discovers protein design principles. *arXiv preprint arXiv:2504.19017* (2025)
- [27] Kim, S., Choi, J., Jang, K., Park, J., Bernales, V., Aspuru-Guzik, A., Jung, Y.: Materealize: a multi-agent deliberation system for end-to-end material design and synthesis. *arXiv preprint arXiv:2601.15743* (2026)
- [28] Baek, J., Jauhar, S.K., Cucerzan, S., Hwang, S.J.: Researchagent: Iterative research idea generation over scientific literature with large language models. *arXiv preprint arXiv:2404.07738* (2024)
- [29] Zimmermann, Y., Bazgir, A., Afzal, Z., Agbere, F., Ai, Q., Alampara, N., Al-Feghali, A., Ansari, M., Antypov, D., Aswad, A., Bai, J., Baibakova, V., Biswajeet, D.D., Bitzek, E., Bocarsly, J.D., Borisova, A., Bran, A.M., Brinson, L.C., Calderon, M.M., Canalicchio, A., Chen, V., Chiang, Y., Circi, D., Charmes, B., Chaudhary, V., Chen, Z., Chiu, M.-H., Clymo, J., Dabhadkar, K., Daelman, N., Datar, A., Jong, W.A., Evans, M.L., Fard, M.G., Fisicaro, G., Gangan, A.S., George, J., Gonzalez, J.D.C., Götte, M., Gupta, A.K., Harb, H., Hong, P., Ibrahim, A., Ilyas, A., Imran, A., Ishimwe, K., Issa, R., Jablonka, K.M., Jones, C., Josephson, T.R., Juhasz, G., Kapoor, S., Kang, R., Khalighinejad, G., Khan,

- S., Klawohn, S., Kuman, S., Ladines, A.N., Leang, S., Lederbauer, M., Sheng-Lun, Liao, Liu, H., Liu, X., Lo, S., Madireddy, S., Maharana, P.R., Maheshwari, S., Mahjoubi, S., Márquez, J.A., Mills, R., Mohanty, T., Mohr, B., Moosavi, S.M., Moßhammer, A., Naghdi, A.D., Naik, A., Narykov, O., Näsström, H., Nguyen, X.V., Ni, X., O'Connor, D., Olayiwola, T., Ottomano, F., Ozhan, A.B., Pagel, S., Parida, C., Park, J., Patel, V., Patyukova, E., Petersen, M.H., Pinto, L., Pizarro, J.M., Plessers, D., Pradhan, T., Pratiush, U., Puli, C., Qin, A., Rajabi, M., Ricci, F., Risch, E., Ríos-García, M., Roy, A., Rug, T., Sayeed, H.M., Scheidgen, M., Schilling-Wilhelmi, M., Schloz, M., Schöppach, F., Schumann, J., Schwaller, P., Schwarting, M., Sharlin, S., Shen, K., Shi, J., Si, P., D'Souza, J., Sparks, T., Sudhakar, S., Talirz, L., Tang, D., Taran, O., Terboven, C., Tropin, M., Tsymbal, A., Ueltzen, K., Unzueta, P.A., Vasan, A., Vinchurkar, T., Vo, T., Vogel, G., Völker, C., Weinreich, J., Yang, F., Zaki, M., Zhang, C., Zhang, S., Zhang, W., Zhu, R., Zhu, S., Janssen, J., Li, C., Foster, I., Blaiszik, B.: Reflections from the 2024 large language model (llm) hackathon for applications in materials science and chemistry. arXiv preprint arXiv:2411.15221 (2025)
- [30] Yamada, Y., Lange, R.T., Lu, C., Hu, S., Lu, C., Foerster, J., Clune, J., Ha, D.: The ai scientist-v2: Workshop-level automated scientific discovery via agentic tree search. arXiv preprint arXiv:2504.08066 (2025)
- [31] Shao, C., Huang, D., Li, Y., Zhao, K., Lin, W., Zhang, Y., Zeng, Q., Chen, Z., Li, T., Huang, Y., Wu, T., Liu, X., Zhao, R., Zhao, M., Li, J., Zhang, X., Wang, Y., Zhen, Y., Xu, F., Li, Y., Liu, T.-Y.: OmniScientist: Toward a co-evolving ecosystem of human and ai scientists. arXiv preprint arXiv:2511.16931 (2025)
- [32] Ding, K., Yu, J., Huang, J., Yang, Y., Zhang, Q., Chen, H.: SciToolAgent: A knowledge graph-driven scientific agent for multi-tool integration. *Nature Computational Science* (2025) <https://doi.org/10.1038/s43588-025-00849-y>
- [33] Miao, T., Dai, J., Liu, J., Tan, J., Zhang, M., Jin, W., Du, Y., Jin, T., Pang, X., Liu, Z., Guo, T., Zhang, Z., Huang, Y., Chen, S., Ye, R., Zhang, Y., Zhang, L., Chen, K., Wang, W., E, W., Chen, S.: PhysMaster: Building an autonomous ai physicist for theoretical and computational physics research. arXiv preprint arXiv:2512.19799 (2025)
- [34] Song, K., Trotter, A., Chen, J.Y.: LLM agent swarm for hypothesis-driven drug discovery. arXiv preprint arXiv:2504.17967 (2025)
- [35] Wang, F.Y., Lee, D.S., Kaplan, D.L., Buehler, M.J.: Swarms of large language model agents for protein sequence design with experimental validation. arXiv preprint arXiv:2511.22311 (2025)
- [36] Denison, C., MacDiarmid, M., Barez, F., Duvenaud, D., Kravec, S., Marks, S., Schiefer, N., Soklaski, R., Tamkin, A., Kaplan, J., Shlegeris, B., Bowman, S.R., Perez, E., Hubinger, E.: Sycophancy to subterfuge: Investigating reward-tampering in large language models. arXiv preprint arXiv:2406.10162



(2024)

- [37] Sharma, M., Tong, M., Korbak, T., Duvenaud, D., Aspell, A., Bowman, S.R., Durmus, E., Hatfield-Dodds, Z., Johnston, S.R., Kravec, S., Maxwell, T., McCandlish, S., Ndousse, K., Rausch, O., Schiefer, N., Yan, D., Zhang, M., Perez, E.: Towards understanding sycophancy in language models. In: The Twelfth International Conference on Learning Representations (ICLR) (2024)
- [38] Dziri, N., Lu, X., Sclar, M., Li, X.L., Jiang, L., Lin, B.Y., West, P., Bhagavatula, C., Le Bras, R., Hwang, J.D., Sanyal, S., Welleck, S., Ren, X., Ettinger, A., Harchaoui, Z., Choi, Y.: Faith and fate: Limits of transformers on compositionality. In: Advances in Neural Information Processing Systems (NeurIPS), vol. 36 (2023)
- [39] Backlund, A., Petersson, L.: Vending-bench: A benchmark for long-term coherence of autonomous agents. arXiv preprint arXiv:2502.15840 (2025)
- [40] Sinha, A., Arun, A., Goel, S., Staab, S., Geiping, J.: The illusion of diminishing returns: Measuring long horizon execution in LLMs. In: The Fourteenth International Conference on Learning Representations (ICLR) (2026)
- [41] Lu, C., Lu, C., Lange, R.T., Foerster, J., Clune, J., Ha, D.: The AI Scientist: Towards fully automated open-ended scientific discovery. arXiv preprint arXiv:2408.06292 (2024)
- [42] Trehan, D., Chopra, P.: Why LLMs aren’t scientists yet: Lessons from four autonomous research attempts. arXiv preprint arXiv:2601.03315 (2026)
- [43] Lin, X., Ning, Y., Zhang, J., Dong, Y., Liu, Y., Wu, Y., Qi, X., Sun, N., Shang, Y., Wang, K., Cao, P., Wang, Q., Zou, L., Chen, X., Zhou, C., Wu, J., Zhang, P., Wen, Q., Pan, S., Wang, B., Cao, Y., Chen, K., Hu, S., Guo, L.: LLM-based agents suffer from hallucinations: A survey of taxonomy, methods, and directions. arXiv preprint arXiv:2509.18970 (2025)
- [44] Wang, H., Feng, S., He, T., Tan, Z., Han, X., Tsvetkov, Y.: Can language models solve graph problems in natural language? Adv. Neural Inf. Process Syst. **36**, 30840–30861 (2023)
- [45] Kandpal, N., Deng, H., Roberts, A., Wallace, E., Raffel, C.: Large language models struggle to learn long-tail knowledge. In: ICML, pp. 15696–15707 (2023). PMLR
- [46] Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y.J., Madotto, A., Fung, P.: Survey of hallucination in natural language generation. ACM Comput. Surv. **55**(12), 1–38 (2023)
- [47] Ge, Y., Kirtane, N., Peng, H., Hakkani-Tür, D.: LLMs are vulnerable to malicious

- prompts disguised as scientific language. arXiv preprint arXiv:2501.14073 (2025)
- [48] Dong, L., Lu, Q., Zhu, L.: Agentops: Enabling observability of LLM agents. arXiv preprint arXiv:2411.05285 (2024)
  - [49] Swanson, K., Wu, W., Bulaong, N.L., Pak, J.E., Zou, J.: The virtual lab of ai agents designs new SARS-CoV-2 nanobodies. *Nature* **646**, 716–723 (2025) <https://doi.org/10.1038/s41586-025-09442-9>
  - [50] Li, M.Y., Vajipey, V., Goodman, N.D., Fox, E.B.: CriticAL: Critic automation with language models. arXiv preprint arXiv:2411.06590 (2024)
  - [51] Dhuliawala, S., Komeili, M., Xu, J., Raileanu, R., Li, X., Celikyilmaz, A., Weston, J.: Chain-of-verification reduces hallucination in large language models. arXiv preprint arXiv:2309.11495 (2023)
  - [52] Zheng, L., Chiang, W.-L., Sheng, Y., Zhuang, S., Wu, Z., Zhuang, Y., Lin, Z., Li, Z., Li, D., Xing, E.P., Zhang, H., Gonzalez, J.E., Stoica, I.: Judging LLM-as-a-judge with MT-Bench and Chatbot Arena. In: *Advances in Neural Information Processing Systems (NeurIPS) Datasets and Benchmarks Track*, vol. 36, pp. 46595–46623 (2023)
  - [53] Zheng, X., Pang, T., Du, C., Liu, Q., Jiang, J., Lin, M.: Cheating automatic LLM benchmarks: Null models achieve high win rates. In: *International Conference on Learning Representations (ICLR)* (2025)
  - [54] Wang, X., Wei, J., Schuurmans, D., Le, Q., Chi, E.H., Narang, S., Chowdhery, A., Zhou, D.: Self-consistency improves chain of thought reasoning in language models. In: *The Eleventh International Conference on Learning Representations (ICLR)* (2023). arXiv:2203.11171
  - [55] Zhu, J., Huang, Y., Shen, Y., Zhao, J., Zou, A.: Path-consistency with prefix enhancement for efficient inference in LLMs. arXiv preprint arXiv:2409.01281 (2024)
  - [56] Wan, G., Wu, Y., Chen, J., Li, S.: Reasoning aware self-consistency: Leveraging reasoning paths for efficient LLM sampling. In: *Proceedings of the 2025 Conference of the Nations of the Americas Chapter of the Association for Computational Linguistics: Human Language Technologies (NAACL-HLT)*, pp. 3613–3635 (2025). arXiv:2408.17017
  - [57] Goel, S., Strüder, J., Auzina, I.A., Chandra, K.K., Kumaraguru, P., Kiela, D., Prabhu, A., Bethge, M., Geiping, J.: Great models think alike and this undermines AI oversight. In: *International Conference on Machine Learning (ICML)* (2025)
  - [58] Smit, A.P., Grinsztajn, N., Duckworth, P., Barrett, T.D., Pretorius, A.: Should we

- be going MAD? A look at multi-agent debate strategies for LLMs. In: Proceedings of the 41st International Conference on Machine Learning (ICML), vol. 235, pp. 45883–45905. PMLR, ??? (2024)
- [59] Laban, P., Murakhovs'ka, L., Xiong, C., Wu, C.-S.: Are you sure? Challenging LLMs leads to performance drops in The FlipFlop Experiment. arXiv preprint arXiv:2311.08596 (2024)
  - [60] Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., Alon, U., Dziri, N., Prabhunoye, S., Yang, Y., Gupta, S., Majumder, B.P., Hermann, K., Welleck, S., Yazdanbakhsh, A., Clark, P.: Self-refine: Iterative refinement with self-feedback. In: Advances in Neural Information Processing Systems (NeurIPS), vol. 36, pp. 46534–46594 (2023)
  - [61] Shinn, N., Cassano, F., Berman, E., Gopinath, A., Narasimhan, K., Yao, S.: Reflexion: Language agents with verbal reinforcement learning. In: Advances in Neural Information Processing Systems (NeurIPS), vol. 36, pp. 8634–8652 (2023)
  - [62] Huang, J., Chen, X., Mishra, S., Zheng, H.S., Yu, A.W., Song, X., Zhou, D.: Large language models cannot self-correct reasoning yet. In: The Twelfth International Conference on Learning Representations (ICLR) (2024)
  - [63] Ghafarollahi, A., Buehler, M.J.: Automating alloy design and discovery with physics-aware multimodal multiagent ai. *Proc. Natl. Acad. Sci.* **122**(4), 2414074122 (2025)
  - [64] Giannozzi, P., Baroni, S., Bonini, N., Calandra, M., Car, R., Cavazzoni, C., Ceresoli, D., Chiarotti, G.L., Cococcioni, M., Dabo, I., Corso, A.D., Gironcoli, S.d., Fabris, S., Fratesi, G., Gebauer, R., Gerstmann, U., Gougoussis, C., Kokalj, A., Lazzeri, M., Martin-Samos, L., Marzari, N., Mauri, F., Mazzarello, R., Paolini, S., Pasquarello, A., Paulatto, L., Sbraccia, C., Scandolo, S., Sclauzero, G., Seitsonen, A.P., Smogunov, A., Umari, P., Wentzcovitch, R.M.: QUANTUM ESPRESSO: a modular and open-source software project for quantum simulations of materials. *J. Phys. Condens. Matter* **21**(39), 395502 (2009)
  - [65] Kudin, K.N., Scuseria, G.E.: Converging self-consistent field equations in quantum chemistry—recent achievements and remaining challenges. *ESAIM: M2AN* **41**(2), 281–296 (2007)
  - [66] Janssen, J., Makarov, E., Hickel, T., Shapeev, A.V., Neugebauer, J.: Automated optimization and uncertainty quantification of convergence parameters in plane wave density functional theory calculations. *npj Comput Mater* **10**, 263 (2024)
  - [67] Wellendorff, J., Lundgaard, K.T., Møgelhøj, A., Petzold, V., Landis, D.D., Nørskov, J.K., Bligaard, T., Jacobsen, K.W.: Density functionals for surface science: Exchange-correlation model development with bayesian error estimation. *J. Phys. Chem. B* **85**, 235149 (2012)

- [68] Feibelman, P.J., Hammer, B., Nørskov, J.K., Wagner, F., Scheffler, M., Stumpf, R., Watwe, R., Dumesic, J.: The co/pt(111) puzzle. *J. Phys. Chem. B* **105**(18), 4018–4025 (2001)
- [69] Allian, A.D., Takanabe, K., Fajdala, K.L., Hao, X., Truex, T.J., Cai, J., Buda, C., Neurock, M., Iglesia, E.: Chemisorption of co and mechanism of co oxidation on supported platinum nanoclusters. *J. Am. Chem. Soc.* **133**(12), 4498–4517 (2011)
- [70] Grinberg, I., Yourdshahyan, Y., Rappe, A.M.: Co on pt(111) puzzle: A possible solution. *J. Chem. Phys.* **117**(5), 2264–2270 (2002)
- [71] Janthon, P., Viñes, F., Sirijaraensre, J., Limtrakul, J., Illas, F.: Adding Pieces to the CO/Pt(111) Puzzle: The Role of Dispersion. *J. Phys. Chem. C* **121**(7), 3970–3977 (2017)
- [72] K. G, L., Kundappaden, I., Chatanathodi, R.: A dft study of co adsorption on pt (111) using van der waals functionals. *Surf. Sci.* **681**, 143–148 (2019)
- [73] Anthropic: Claude Sonnet 4.5. <https://docs.anthropic.com/en/docs/about-claude/models>
- [74] Anthropic: Claude Opus 4.8. <https://docs.anthropic.com/en/docs/about-claude/models>
- [75] Larsen, A.H., Mortensen, J.J., Blomqvist, J., Castelli, I.E., Christensen, R., Dulak, M., Friis, J., Groves, M.N., Hammer, B., Hargus, C., Hermes, E.D., Jennings, P.C., Jensen, P.B., Kermode, J., Kitchin, J.R., Kolsbjerg, E.L., Kubal, J., Kaasbjerg, K., Lysgaard, S., Maronsson, J.B., Maxson, T., Olsen, T., Pastewka, L., Peterson, A., Rostgaard, C., Schiøtz, J., Schütt, O., Strange, M., Thygesen, K.S., Vegge, T., Vilhelmsen, L., Walter, M., Zeng, Z., Jacobsen, K.W.: The atomic simulation environment—a python library for working with atoms. *J. Phys. Condens. Matter* **29**(27), 273002 (2017)
- [76] Kresse, G., Gil, A., Sautet, P.: Significance of single-electron energies for the description of co on pt(111). *Phys. Rev. B* **68**, 073401 (2003)
- [77] Blyholder, G.: Molecular orbital view of chemisorbed carbon monoxide. *J. Phys. Chem.* **68**(10), 2772–2777 (1964)
- [78] LangChain Inc.: LangGraph. <https://www.langchain.com/langgraph>
- [79] Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., Cao, Y.: React: Synergizing reasoning and acting in language models. arXiv preprint arXiv:2210.03629v3 (2022)
- [80] Jette, M.A., Wickberg, T.: Architecture of the slurm workload manager, 3–23 (2023). Springer

- [81] Sholl, D.S., Steckel, J.A.: Density Functional Theory: a Practical Introduction. John Wiley & Sons, ??? (2022)
- [82] Kresse, G., Hafner, J.: Ab initio molecular dynamics for liquid metals. *Phys. Rev. B* **47**, 558–561 (1993)
- [83] Frisch, M.J., Trucks, G.W., Schlegel, H.B., Scuseria, G.E., Robb, M.A., Cheeseman, J.R., Scalmani, G., Barone, V., Petersson, G.A., Nakatsuji, H., Li, X., Caricato, M., Marenich, A.V., Bloino, J., Janesko, B.G., Gomperts, R., Mennucci, B., Hratchian, H.P., Ortiz, J.V., Izmaylov, A.F., Sonnenberg, J.L., Williams-Young, D., Ding, F., Lipparini, F., Egidi, F., Goings, J., Peng, B., Petrone, A., Henderson, T., Ranasinghe, D., Zakrzewski, V.G., Gao, J., Rega, N., Zheng, G., Liang, W., Hada, M., Ehara, M., Toyota, K., Fukuda, R., Hasegawa, J., Ishida, M., Nakajima, T., Honda, Y., Kitao, O., Nakai, H., Vreven, T., Throssell, K., Montgomery, J.A. Jr., Peralta, J.E., Ogliaro, F., Bearpark, M.J., Heyd, J.J., Brothers, E.N., Kudin, K.N., Staroverov, V.N., Keith, T.A., Kobayashi, R., Normand, J., Raghavachari, K., Rendell, A.P., Burant, J.C., Iyengar, S.S., Tomasi, J., Cossi, M., Millam, J.M., Klene, M., Adamo, C., Cammi, R., Ochterski, J.W., Martin, R.L., Morokuma, K., Farkas, O., Foresman, J.B., Fox, D.J.: Gaussian~16 Revision C.01. Gaussian Inc. Wallingford CT (2016)
- [84] Viswanathan, V., Thygesen, K.S., Hummelshøj, J., Nørskov, J.K., Girishkumar, G., McCloskey, B., Luntz, A.: Electrical conductivity in  $\text{Li}_2\text{O}_2$  and its role in determining capacity limitations in non-aqueous  $\text{Li-O}_2$  batteries. *J. Chem. Phys.* **135**(21), 214704 (2011)

# Supplementary Information: DREAMS: Density Functional Theory Based Research Engine for Agentic Materials Simulation

Ziqi Wang<sup>1</sup>, Hongshuo Huang<sup>1</sup>, Hancheng Zhao<sup>1</sup>,  
Changwen Xu<sup>1</sup>, Shang Zhu<sup>1</sup>, Jan Janssen<sup>3\*</sup>,  
Venkatasubramanian Viswanathan<sup>1,2\*</sup>

<sup>1</sup>Department of Mechanical Engineering, University of Michigan.

<sup>2</sup>Department of Aerospace Engineering, University of Michigan.

<sup>3</sup>Materials Informatics, Max-Planck-Institute for Sustainable Materials.

\*Corresponding author(s). E-mail(s): [janssen@mpi-susmat.de](mailto:janssen@mpi-susmat.de);  
[venkvis@umich.edu](mailto:venkvis@umich.edu);

## 1 Plan generated for the Sol27LC challenge

The initial full plan generated by the supervisor for the BCC-Li run, created immediately after a single consultation step with the DFT agent. Under the safety-guard framework, every plan step names the worker agent responsible, the tool whose successful call the step requires, and, for report steps, the quantities the resulting report must deliver; the plan is subsequently edited dynamically (steps inserted, modified, or deleted) as the run proceeds. In this run the plan executed without revision.

1. Create BCC Li bulk structure with initial lattice constant guess ( $\sim 3.49$  Å) (*DFT\_Agent; init\_structure\_data\_for\_none\_adsorption\_calculations*)
2. Find pseudopotential for Li using PBE functional (*DFT\_Agent; find\_pseudopotential*)
3. Write SCF template with tight initial DFT parameters for convergence tests (*DFT\_Agent; write\_QE\_script\_w\_ASE*)
4. Generate ecutwfc convergence test input files (*DFT\_Agent; generate\_convergence\_test*)
5. Add resource suggestions for ecutwfc convergence test jobs (*HPC\_Agent; add\_resource\_suggestion*)
6. Submit and monitor ecutwfc convergence test jobs (*HPC\_Agent; submit\_and\_monitor\_job*)
7. Determine optimal ecutwfc from convergence test results (*DFT\_Agent; find\_optimal\_parameter*)
8. Generate intermediate report for optimal ecutwfc (*DFT\_Agent; generate\_structured\_report; must report: optimal\_ecutwfc\_Ry*)
9. Generate kspacing convergence test input files using optimal ecutwfc (*DFT\_Agent; generate\_convergence\_test*)
10. Add resource suggestions for kspacing convergence test jobs (*HPC\_Agent; add\_resource\_suggestion*)
11. Submit and monitor kspacing convergence test jobs (*HPC\_Agent; submit\_and\_monitor\_job*)
12. Determine optimal kspacing from convergence test results (*DFT\_Agent; find\_optimal\_parameter*)
13. Generate intermediate report for optimal kspacing (*DFT\_Agent; generate\_structured\_report; must report: optimal\_kspacing*)
14. Generate EOS test input files using optimal ecutwfc and kspacing (*DFT\_Agent; generate\_eos\_test*)
15. Add resource suggestions for EOS test jobs (*HPC\_Agent; add\_resource\_suggestion*)
16. Submit and monitor EOS test jobs (*HPC\_Agent; submit\_and\_monitor\_job*)
17. Calculate lattice constant from EOS results (*DFT\_Agent; calculate\_lc*)
18. Generate final report with lattice constant (*DFT\_Agent; generate\_structured\_report; must report: lattice\_constant\_angstrom*)

## 2 Full Canvas History for the Sol27LC Challenge

```
{
  'Cu_structure_file': '/nfs/turbo/coe-venkvis/zqiw-turbo/material_agent/out/Cu-fcc.xyz',
  'Cu_pseudopotential': 'cu_pbe_v1.2.uspp.F.UPF',
  'Cu_lattice_constant': 3.596,
  'scratch_job_list': ['Cu_fcc_scf.pwl'],
  'ready_to_run_job_list': [],
  'jobs_K_and_ecut': {
    }
  'Cu_initial_dft_script': 'Cu_fcc_scf.pwl',
  'Cu_convergence_test': 'Generated convergence test with kspacing [0.5, 0.4, 0.3, 0.2] and ecutwfc [30, 40, 50, 60]',
  'Cu_convergence_test_files': 'Generated convergence test input files with kspacing [0.5, 0.4, 0.3, 0.2] and ecutwfc [30, 40, 50, 60]. Files are ready to run.',
  'finished_job_list': ['Cu_fcc_scf_k_0.3_ecutwfc_60.pwl', 'Cu_fcc_scf_k_0.2_ecutwfc_50.pwl',
    'Cu_fcc_scf_k_0.5_ecutwfc_60.pwl', 'Cu_fcc_scf_k_0.2_ecutwfc_30.pwl', 'Cu_fcc_scf_k_0.4_ecutwfc_60.pwl',
    'Cu_fcc_scf_k_0.2_ecutwfc_40.pwl', 'Cu_fcc_scf_k_0.2_ecutwfc_60.pwl', 'Cu_fcc_scf_1.025.pwl',
    'Cu_fcc_scf_0.975.pwl', 'Cu_fcc_scf_0.95.pwl', 'Cu_fcc_scf_1.0.pwl', 'Cu_fcc_scf_1.05.pwl'],
  'convergence_jobs_status': 'completed',
  'convergence_jobs_output_location': '/nfs/turbo/coe-venkvis/zqiw-turbo/material_agent/out',
  'Cu_optimal_parameters': 'Based on convergence testing, the optimal parameters are kspacing=0.2 and ecutwfc=60 Ry. Note that the calculation may benefit from even finer parameters if higher precision is needed.',
  'Cu_eos_test_files': 'EOS calculation input files generated with optimal parameters (kspacing=0.2, ecutwfc=60 Ry). Files created: Cu_fcc_scf_0.95.pwl, Cu_fcc_scf_0.975.pwl, Cu_fcc_scf_1.0.pwl, Cu_fcc_scf_1.025.pwl, Cu_fcc_scf_1.05.pwl',
  'eos_jobs_status': 'All EOS calculation jobs completed successfully',
  'eos_jobs_output_location': '/nfs/turbo/coe-venkvis/zqiw-turbo/material_agent/out',
  'Cu_eos_energies': {
    'Cu_fcc_scf_0.95.pwl': -21973.869858285174,
    'Cu_fcc_scf_0.975.pwl': -21974.762386112317,
    'Cu_fcc_scf_1.0.pwl': -21975.073777270612,
    'Cu_fcc_scf_1.025.pwl': -21975.024483576628,
    'Cu_fcc_scf_1.05.pwl': -21974.75450052539,
  }
  'Cu_equilibrium_lattice_constant': 3.6226269522562853,
}
```



### 3 Full results for the Sol27LC challenge

**Table 1:** Comparison of Sol27LC benchmark results from experimental measurements, DFT calculations by human experts, and DFT calculations by DREAMS. The Perdew-Burke-Ernzerhof (PBE) exchange-correlation functional is used for DFT calculations. Errors in DREAMS’ DFT calculations relative to expert-implemented DFT calculations are presented in the last column.

| Structure | Lattice parameter results (Å) |               |        | k points | ecutwfc | Error (%) |
|-----------|-------------------------------|---------------|--------|----------|---------|-----------|
|           | Experimental                  | Human experts | Agent  |          |         |           |
| Li (BCC)  | 3.451                         | 3.4361        | 3.44   | 12       | 40      | 0.1135    |
| Na (BCC)  | 4.209                         | 4.1953        | 4.2018 | 6        | 70      | 0.1549    |
| K (BCC)   | 5.212                         | 5.2834        | 5.2676 | 8        | 60      | 0.299     |
| Rb (BCC)  | 5.577                         | 5.6869        | 5.6475 | 8        | 30      | 0.6928    |
| Ca (FCC)  | 5.556                         | 5.5321        | 5.5357 | 8        | 50      | 0.0651    |
| Sr (FCC)  | 6.04                          | 6.0434        | 6.0378 | 8        | 40      | 0.0927    |
| Ba (BCC)  | 5.002                         | 5.0238        | 5.0148 | 10       | 50      | 0.1791    |
| V (BCC)   | 3.024                         | 3.0196        | 2.9996 | 12       | 50      | 0.6623    |
| Nb (BCC)  | 3.294                         | 3.3153        | 3.3084 | 12       | 70      | 0.2081    |
| Ta (BCC)  | 3.299                         | 3.3345        | 3.3157 | 8        | 70      | 0.5638    |
| Mo (BCC)  | 3.141                         | 3.1682        | 3.1589 | 16       | 60      | 0.2935    |
| W (BCC)   | 3.16                          | 3.2034        | 3.182  | 8        | 70      | 0.668     |
| Fe (BCC)  | 2.853                         | 2.8442        | 2.8399 | 16       | 60      | 0.1512    |
| Rh (FCC)  | 3.793                         | 3.855         | 3.823  | 10       | 60      | 0.8301    |
| Ir (FCC)  | 3.831                         | 3.8907        | 3.8636 | 10       | 70      | 0.6965    |
| Ni (FCC)  | 3.508                         | 3.524         | 3.5155 | 8        | 60      | 0.2412    |
| Pd (FCC)  | 3.876                         | 3.9474        | 3.926  | 18       | 40      | 0.5421    |
| Pt (FCC)  | 3.913                         | 3.982         | 3.956  | 8        | 70      | 0.6529    |
| Cu (FCC)  | 3.596                         | 3.6516        | 3.6213 | 10       | 70      | 0.8298    |
| Ag (FCC)  | 4.062                         | 4.1561        | 4.1258 | 18       | 50      | 0.729     |
| Au (FCC)  | 4.062                         | 4.169         | 4.1329 | 8        | 70      | 0.8659    |
| Al (FCC)  | 4.019                         | 4.0425        | 4.0333 | 18       | 30      | 0.2276    |
| Pb (FCC)  | 4.912                         | 5.0467        | 5.0258 | 14       | 70      | 0.4141    |
| C (dia)   | 3.544                         | 3.5729        | 3.5562 | 8        | 70      | 0.4674    |
| Si (dia)  | 5.415                         | 5.4758        | 5.4392 | 6        | 40      | 0.6684    |
| Ge (dia)  | 5.639                         | 5.7656        | 5.6902 | 6        | 70      | 1.3078    |
| Sn (dia)  | 6.474                         | 6.6582        | 6.553  | 6        | 70      | 1.58      |

## 4 Plan generated for the CO/Pt(111) challenge

The initial full plan generated by the supervisor for the CO/Pt(111) run, created immediately after a single consultation step with the DFT agent. As in the Sol27LC plan, each step names its worker agent, its required tool, and, for report steps, the quantities the report must deliver. This initial plan was subsequently revised repeatedly during the run (see the main text): among others, the `ecutwfc` report was deferred after the deterministic tool refusal, an extended `ecutwfc` test was inserted following the report judge's warning, the k-point and slab-thickness tests were regenerated at the new cutoff, and, after the report judge failed the first final report over the uncharacterized vacuum spacing, a vacuum convergence test and a regenerated round of production calculations were inserted.

1. Generate Pt(111) slab structure with p(2x2) supercell and initial structural parameters, and find PBE pseudopotentials for Pt, C, O (*DFT\_Agent; find\_pseudopotential*)
2. Generate isolated CO molecule structure and add it to fcc site with 0° orientation for convergence testing (*DFT\_Agent; add\_myAdsorbate*)
3. Generate `ecutwfc` convergence test input files for CO on fcc site using PBE pseudopotentials (*DFT\_Agent; generate\_convergence\_test*)
4. Add resource suggestions and submit `ecutwfc` convergence test jobs to HPC (*HPC\_Agent; submit\_and\_monitor\_job*)
5. Determine optimal `ecutwfc` from convergence test results (*DFT\_Agent; find\_optimal\_parameter*)
6. Generate intermediate report certifying optimal `ecutwfc` (*DFT\_Agent; generate\_structured\_report; must report: optimal\_ecutwfc\_Ry*)
7. Generate `kspacing` convergence test input files for CO on fcc site using optimal `ecutwfc` and PBE pseudopotentials (*DFT\_Agent; generate\_convergence\_test*)
8. Add resource suggestions and submit `kspacing` convergence test jobs to HPC (*HPC\_Agent; submit\_and\_monitor\_job*)
9. Determine optimal `kspacing` from convergence test results (*DFT\_Agent; find\_optimal\_parameter*)
10. Generate intermediate report certifying optimal `kspacing` (*DFT\_Agent; generate\_structured\_report; must report: optimal\_kspacing*)
11. Generate structural convergence test (`supercell_dim_z`) input files for CO on fcc site using optimal DFT parameters and PBE pseudopotentials (*DFT\_Agent; generate\_convergence\_test*)
12. Add resource suggestions and submit `supercell_dim_z` convergence test jobs to HPC (*HPC\_Agent; submit\_and\_monitor\_job*)
13. Determine optimal `supercell_dim_z` from derived-quantity convergence test results (*DFT\_Agent; find\_optimal\_parameter\_from\_derived*)
14. Generate intermediate report certifying optimal `supercell_dim_z` (*DFT\_Agent; generate\_structured\_report; must report: optimal\_supercell\_dim\_z*)
15. Generate structural convergence test (`n_fixed_layers`) input files for CO on fcc site using optimal DFT parameters and optimal `supercell_dim_z` (*DFT\_Agent; generate\_convergence\_test*)


16. Add resource suggestions and submit `n_fixed_layers` convergence test jobs to HPC (*HPC\_Agent; submit\_and\_monitor\_job*)
17. Determine optimal `n_fixed_layers` from derived-quantity convergence test results (*DFT\_Agent; find\_optimal\_parameter\_from\_derived*)
18. Generate intermediate report certifying optimal `n_fixed_layers` (*DFT\_Agent; generate\_structured\_report; must report: optimal\_n\_fixed\_layers*)
19. Generate production input files for clean Pt(111) slab, isolated CO, and all 6 CO+slab configurations (fcc 0°/90°/180°, ontop 0°/90°/180°) using all optimal parameters and PBE pseudopotentials (*DFT\_Agent; write\_QE\_script\_w\_ASE*)
20. Add resource suggestions and submit all production jobs to HPC (*HPC\_Agent; submit\_and\_monitor\_job*)
21. Calculate adsorption energies for all 6 configurations, identify most favorable fcc and ontop configurations, and compute the adsorption energy difference (*DFT\_Agent; calculate\_formation\_E*)
22. Generate final report certifying all adsorption energies and the final adsorption energy difference between most favorable fcc and ontop configurations (*DFT\_Agent; generate\_structured\_report; must report: adsorption\_energy\_fcc\_0deg\_eV, adsorption\_energy\_fcc\_90deg\_eV, adsorption\_energy\_fcc\_180deg\_eV, adsorption\_energy\_ontop\_0deg\_eV, adsorption\_energy\_ontop\_90deg\_eV, adsorption\_energy\_ontop\_180deg\_eV, adsorption\_energy\_difference\_fcc\_vs\_ontop\_eV*)

## 5 More examples of suggestions from the convergence LLM agent

**Table 2:** Representative suggestions made by the convergence LLM agent based on DFT input files. The agent proposes parameter adjustments such as increasing the smearing width (`degauss`) and decreasing the mixing factor (`mixing_beta`)—to improve SCF convergence in metallic systems. Each suggestion is accompanied by concise justifications, reflecting practices commonly employed by experienced DFT users. These entries illustrate the types of remedies proposed during the benchmark runs; they are not a frequency distribution or a controlled measure of repeatability.

| Parameters               | Suggestions                            | Reason                                                                                                                                                       |
|--------------------------|----------------------------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>ecutwfc</b>           | increase to 80.0                       | higher cutoff needed for ultrasoft pseudopotentials with transition metals                                                                                   |
| <b>degauss</b>           | increase to 0.03                       | helps with metallic systems                                                                                                                                  |
| <b>mixing_beta</b>       | add with value 0.3                     | The current value of 0.7 is too high for this complex system with CO adsorption, leading to charge sloshing                                                  |
| <b>mixing_mode</b>       | local-TF                               | Switching from plain mixing to a more sophisticated scheme like local-TF can help with difficult convergence cases                                           |
| starting magnetization   | add for Pt (0.1)                       | helps break symmetry for transition metals                                                                                                                   |
| vacuum                   | consider increasing (cell z-dimension) | current vacuum might be insufficient for CO adsorption.                                                                                                      |
| <b>electron maxstep</b>  | 300                                    | The calculation stopped at 200 iterations but was still making progress, so allowing more iterations might help it converge                                  |
| conv_thr                 | 1.0e-5                                 | The current threshold (1.0E-06) might be unnecessarily strict for geometry optimization steps; relaxing it initially will help reach convergence faster      |
| <b>startingwfc</b>       | atomic+random                          | Using a better initial guess for wavefunctions can accelerate convergence                                                                                    |
| <b>diagonalization</b>   | david                                  | Explicitly setting the diagonalization algorithm to david (which is the default) with a higher david_ndim parameter (e.g., 4) might help                     |
| <b>scf_must converge</b> | false                                  | For geometry optimization, you can allow the calculation to proceed even if SCF doesn't fully converge, which is often acceptable during intermediate steps. |

## 6 Charge density of the CO/Pt(111) system



**Fig. 1:** Charge density of CO on the Pt(111) surface for different XC choice: (a)LDA, FCC site; (b) LDA, ontop site; (c) PBE, FCC site; (d) PBE, ontop site; (e) BEEF-vdW, FCC site; (f) BEEF-vdW, ontop site.

## 7 DREAMS Agent Prompts

The complete operational prompts used by the planning supervisor, DFT worker, convergence-diagnosis agent, and HPC worker are reproduced below. The planning-supervisor and HPC prompts are assembled at runtime from static templates and the substitution blocks shown here; angle-bracketed labels in the displayed templates mark those substitutions. The exact, answer-free benchmark objectives are reproduced in the corresponding Results subsections of the main text. The rule texts and judge guidance used by the safety-guard and report-judge agents are reported separately in Supplementary Information S-11.

### Planning-supervisor substitutions

```
<workers> = ["DFT_Agent", "HPC_Agent"]
<worker-options> = ["DFT_Agent", "HPC_Agent"]

<team-capabilities>
<DFT Agent>:
  - Answer the supervisor's planning questions about DFT settings,
    convergence-test design, structure choices, and any other
    domain-specific decisions
  - Create initial structure of the system
  - Find pseudopotential
  - Write initial script
  - generate convergence test input files for dft parameters
  - determine the best parameters from convergence test result
  - generate different structures for structural convergence test
  - determine best structure settings from structural convergence test
  - generate EOS calculation input files using the best parameters
  - generate production run input files
  - generate BEEF input files from finished relax calculation
  - analyze BEEF result to get uncertainty
  - Read output file to get energy
  - Calculate lattice constant
  - Calculate formation energy
  - Generate structured report
<HPC Agent>:
  - find job list from the job list file
  - Add resource suggestion base on the DFT input file
  - Submit job to HPC and report back once all jobs are done

<team-restrictions>
<DFT Agent>:
  - Cannot submit job to HPC
<HPC Agent>:
  - Cannot determine the best parameters from convergence test result
```

### Planning-supervisor system prompt

```
<Role>
```

You are a scientist supervisor tasked with managing a conversation for scientific computing between the following workers: <workers>. You don't have to use all the members, nor all the capabilities of the members.

<Objective>

Given the following user request, decide which the member to act next, and do what

<Instructions>:

- 0, You will be given the overall objective from the user, a plan consists of a list of high level steps to achieve the objective, and a list of past steps that have been done.
1. If the plan is empty, for the given objective, first discuss with your worker agents to see what they can do and gather their opinions on approach, then come up with a simple, high level plan and specify what are the must use tools to finish the steps. For reference, the team's nominal capabilities are: <team-capabilities> and restrictions are: <team-restrictions> -- but treat this as a starting hint, not the ground truth. Confirm with the workers before committing the plan. You don't have to use all the members, nor all the capabilities of the members. This plan should involve individual tasks, that if executed correctly will yield the correct answer. Do not add any superfluous steps. The result of the final step should be the final answer. Make sure that each step has all the information needed - do not skip steps.

If you were asked to provide uncertainty information across different exchange correlation functionals, you can run ensemble calculation with BEEF-vdW functional and analyze the result. Otherwise, please use the functional that is consistent with the pseudopotentials.

To run ensemble calculation, with the same functional you need to first relax the structure, and then for the relaxed structure of interests, use the BEEF-vdW functional and ensemble calculation to get the distribution of energies. (do not generate ensemble calculations for all relaxed structures, only the ones that are needed for the final answer.)

!!! To calculate adsorption energy, you need to run calculations with the SAME functional for: relaxed adsorbate, relaxed clean slab, and relaxed slab with adsorbate !!!

!!! If you are going to use BEEF-vdW functional later you need to use BEEF-vdW functional for all earlier calculations !!!  
In the plan, you need to be clear what pseudopotential to use when finding pseudopotentials, what functional to use when generating the input files.

If the plan is not empty, update the plan based on the current state of the project (check only related information on CANVAS).

Do not read through the entire CANVAS).

<WARNING>: Critically evaluate the worker's last step. If the action matches the task but serves a different objective, or if extra actions were taken that conflict with the step's intended purpose, treat the step as incorrect. Revise the plan and instruct the worker to redo the step for the correct objective. The framework maintains the plan across turns: you do NOT re-author the whole plan. You apply edits with the plan-editing tools (insert/modify/delete steps), and step 1 of the plan is always what executes next. After editing (or if no edit is needed), Proceed to run step 1; return Response only when the whole objective is done. Make sure that each step has all the information needed - do not skip steps. (Detailed plan-editing instructions are provided to you at runtime.)

2. Given the conversation above, suggest who should act next. next could only be selected from: <worker-options>.
  3. inspect the CANVAS, extract information needed, then base on what the agent just did, the info you extracted, and the plan, decide what to do next.
  4. If your end result is genuinely surprising -- outside the user's margin of error, or in clear conflict with a known reference -- do NOT immediately re-run calculations at random. Instead:
    - a. Call ``debug_artifact_chain`` with the surprising artifact's ``result_id`` and a specific question (e.g. "why is the adsorption energy 0.5 eV higher than the literature value of -1.23 eV?"). The tool returns hypotheses, not conclusions.
    - b. Read the synthesis paragraph for the most likely cause(s), then read the numbered ``potential_causes`` list.
    - c. Instruct the worker to investigate ONE potential cause at a time, in priority order from the synthesis. Vary that single parameter and re-run the relevant calculation.
    - d. After the worker reports back, if the cause was ruled out, call ``debug_artifact_chain`` AGAIN with the same root and question, but pass ``investigation_history`` describing what was tested and ruled out. Without it, the tool will re-flag the same suspects.
    - e. Stop when the cause is found, or when the synthesis says the surprise may originate outside the value-flow chain. In that case, examine the input data, physical assumptions, external benchmark, or the user's expectation before declaring a problem.
- Do not use ``debug_artifact_chain`` as a first-pass sanity check on every result -- it is expensive. Use it only when the result is genuinely surprising. Do not stop until the end result is within the user-specified margin of error, or you have exhausted both the chain investigation and the out-of-chain candidates. Only if the user did not specify a margin of error, you can judge by yourself.
5. Based on the teams capability: <team-capabilities> and restrictions: <team-restrictions>, feel free to add more steps



to the plan if you want to investigate more or if you think it is necessary.

6. After a report was generated, a judge will check the report. If the report does not cleanly pass, the verifier's full feedback is saved to the canvas under the `JUDGE::<report name>` key named in the executed-steps history -- read it there with `read_my_canvas` before reacting. If there are issues, the worker did something wrong: reflect on the feedback, adjust the plan accordingly, and ask the worker to fix the issue. Do not stop until the judge is satisfied with the report. (You are also reminded of this at runtime.)

<Requirements>:

1. Use your worker agents as a resource for expert input when facing decisions that require their domain knowledge. This is necessary when making or adjusting the plan, especially at the start of the project.
2. To consult a worker, insert a step at position 1 (the front of the plan) assigned to that worker; it runs next. Use `insert_plan_steps` for this.
3. For the inserted discussion step, directly ask the question -- do not say anything else. The worker will read the question and answer it, then you can update your plan based on the answer.
4. Do not generate convergence test for all systems and all configurations.
5. To determine the DFT calculation parameters, please only generate one batch of convergence test for the most complicated system using **!! ONE !!** most complicated configuration.
6. Structural convergence test is only needed for adsorption energy calculations, where it is essential to make sure the structure settings like slab thickness and vacuum size are good enough to get converged adsorption energy. Note: there are two distinct convergence-test patterns ("single-output" -- one .pwo file per data point provides the converging quantity; "derived-quantity" -- the converging quantity is computed by aggregating multiple DFT jobs per data point).
7. Only work on structure convergence test once you have determined the best DFT parameters, and make sure to use the best DFT parameters for the structural convergence test.
8. Do not work on structural convergence test (slab thickness, vacuum size) and DFT parameter convergence test (k-points, ecut) at the same time.
9. **\*\*Comparison-set consistency.\*\*** For any set of comparison-based result (optimal parameter from a sweep, lattice constant from EOS, adsorption/formation energy), the underlying input files must share IDENTICAL settings except for the axis being varied. Watch for the failure mode where the worker fixed a convergence problem on ONE file (raised `electron_maxstep`, changed `mixing_beta`, etc.) and reran only that file. If you suspect this -- or the worker says they re-ran "one file" or "the failing job" in isolation -- direct them to align settings across all files in the set and rerun them all before computing the final result.

10. The Must-use tools for each step must be a bare minimum, so your worker can have more degree of freedom.
11. A structured report must be generated once any numerical claim was made so that claim could be verified by the judge, together with a final report. You need to read the feedback from the judge, and if there are any issues, reflect on the feedback, adjust the plan accordingly, create more steps and ask the worker agent try to fix the issue.
12. Whenever you create a step whose `required\_tools` is `generate\_structured\_report`, you MUST also populate that step's `required\_quantities` field with the exact, specific named quantities that report must certify. Name the quantities precisely and unambiguously (e.g. `optimal\_ecutwfc\_Ry`, `optimal\_kspacing`, `optimal\_n\_fixed\_layers`, `adsorption\_energy\_difference\_fcc\_vs\_ontop\_eV`). For all non-report steps, leave `required\_quantities` empty.

<Reading verifier output (`verify\_structured\_report`)>

=====

This output is not shown to you inline. When a report does not cleanly pass, it is saved to the canvas under the `JUDGE::<report name>` key named in the executed-steps history; fetch it with `read_my_canvas`, then read it using the structure below.

Top-level fields: `overall\_verdict`, `n\_fails`, `n\_warnings`, `summary`, `issues` (numbered, your primary surface), plus `checked\_result\_ids` and `artifact\_results` for diagnostic tooling.

Each issue has: `issue\_number`, `category`, `severity`, `where` (structured location with keys like `claim_name` / `result_id` / `tool_name` / `parameter`), `context\_at\_site` (the `Context:` line the worker wrote at the offending site -- read it to see how the agent FRAMED the call), `problem` (one-line description), `judge\_reasoning` (the judge's full reasoning), and `remediation\_options` (legitimate fix paths; some entries warn against WRONG fixes -- do not skip those warnings).

Two cross-cutting rules:

- (1) DO NOT instruct the worker to "find a result\_id and patch it in" for UNSOURCED\_SENSITIVE or CROSS\_WIRED\_SOURCE issues. The legitimate fix is in the remediation\_options -- typically running an upstream sub-study, declaring the parameter varied, or correcting the call's context. Patching with any ID that fits will fail on the next pass as VALUE\_MISMATCH\_PARAM or CROSS\_WIRED\_SOURCE.
- (2) When any artifact gets regenerated, every artifact downstream of it must be re-created using the new result\_id. Stale references silently propagate. Tell the worker explicitly which downstream calculations to re-run, in dependency order.

When multiple issues share a `result\_id` or `claim\_name`, they may collapse to a single root cause; address claim-level issues (VALUE\_MISMATCH\_CLAIM, SCHEMA\_VIOLATION) first since fixing them often dissolves dependent issues.

<Reading debug output (`debug\_artifact\_chain`)>  
=====

Top-level fields: `investigation\_question`, `root\_result\_id`, `n\_potential\_causes`, `summary`, `potential\_causes` (numbered, your primary surface), `synthesis` (one-paragraph diagnostic with a fixed disclaimer prefix), `budget\_exceeded` (if true, raise max\_judge\_calls and re-call).

CRITICAL: debug-tool output is HYPOTHESES, not conclusions. Confirming a hypothesis requires actually changing the parameter and re-running. Framing is "investigate" and "test", not "fix". Do not omit the synthesis disclaimer when relaying findings to the worker.

Each `potential\_cause` has the same field shape as a verifier issue. Two categories:

- parameter\_value\_suspect -- investigate by re-running with this parameter varied.
- source\_suspect -- investigate the upstream source artifact named in `where.source\_result\_id`; re-call `debug\_artifact\_chain` with that upstream result\_id as the new root.

Two rules:

- (1) Vary ONE parameter at a time in priority order from the synthesis. Varying multiple at once makes attribution impossible.
- (2) On follow-up debug calls, ALWAYS pass `investigation\_history` describing what was tested and ruled out, in prose.

When `n\_potential\_causes` is 0: read the synthesis. The chain has been exhausted; investigate input data, physical assumptions, external benchmark, or the user's expectation.

<Final Note>

When the worker raises a request or suggestion, do not ignore it. Evaluate whether it is valid in the context of the overall objective; if it is, update the plan to accommodate it and guide the worker through the change.

At least 2 report must be generated during the project, one in the middle of the project to summarize the progress and one at the end of the project to summarize the final result.

Report-bounded spans of history that did not cleanly pass may be offered to you at runtime as "compression opportunities"; you may compress them with compress\_history to shorten the running history. This is

non-destructive -- the full detail is archived to the canvas and remains readable.

## DFT-worker system prompt

<Role>:

You are a very powerful and yet obedient assistant that performs density functional theory calculations and working in a team. You do exactly what you are told to do.

You and your team members has a shared CANVAS to record and share all the intermediate results.

Please strickly follow the tasks given, do not do anything else.

<Objective>:

You are responsible for generating the quantum espresso input file for the given material and parameter setting with provided tools.

You can only respond with a single complete 'Thought, Action' format OR a single 'Intermediate Answer' format.

Please strickly follow the tasks given, do not do anything else.

<Your Capability>: (Only do what you are told to do)

inspect and read the CANVAS with suitable tools to see what's available.

create valid input structure for the system of interest with the right tool.

Find the correct pseduopotential filename using the tool provided (do not report the absolute path).

Generate the quantum espresso input file with proper ASE tool. Pay attention to calculation type and funtional choice.

Always generate conventional cell with ibrav=0 and do not use celldm and angstrom at the same time.

If the system involves hubbard U correction, specify starting magnetization in SYSTEM card and hubbard U parameters in HUBBARD card, and use the pre-defined hubbard correction tool.

Save all the files in pwi format and into job list and report to supervisor to let HPC Agent to submit the job.

generate convergence test scripts with a tool.

determine the most optimal settings based on the convergence test.

calculate lattice constant and formation energy based on the DFT calculation.

remember to record the results and critical informations in the CANVAS with the right tool.

<Requirements>:

# Workflow & response shape

0. When asked a question, think about it carefully, scientifically, and critically before replying. Draw on your domain expertise; do not rush a casual answer. Then write down your detailed answer in the CANVAS.
1. Always inspect and read the CANVAS with suitable tools to see what's available before acting.
2. Use only the tools necessary for the task. You don't have to use all the tools provided. Never compute values yourself -- call the math tool instead.

3. When using a tool, always fill in the context and reasons fields first.
  4. Once you're done with the assigned task, report back to the supervisor and stop immediately. Do not conduct any inference or post-processing on the result, and do not suggest what to do next.
  - 4a. Recording on CANVAS: use ``write_progress_note`` to create a multi-section progress note, ``version_controlled_edit_canvas_doc`` to revise an existing note, and ``write_my_canvas`` only for short factual entries (it requires you to declare ``is_free_handed_progress_note``).
  - 4b. After you call ``generate_structured_report``, your computation/extraction/submission tools become unavailable -- only canvas tools remain. Write your summary/notes and return to the supervisor; do not attempt further calculations on the report's values.
  - 4c. The math and extraction tools verify each value as it is created; if you get a ``MATH_GUARD_FAILED`` or ``EXTRACTION_GUARD_FAILED``, the fix is to correct the upstream source or clarify your ``context``/``reasons`` -- not to blindly retry.
  5. Response format:
    - On success: respond with the tool message and a short summary of what was done. If this is the final answer, prefix with 'Intermediate Answer'. The final answer should be a concise summary in a sentence -- do not repeat what's already on the CANVAS, just mention it's there.
    - On error: respond with 'Job failed' followed by the error message. Nothing else.
- # File & path conventions
6. QE input files are in .pwi format; output files are .pwo (the engine appends .pwo to the input filename).
  7. Never report absolute paths.
- # Scientific defaults
8. Electron conv\_thr is 1e-6.
  9. Use the smearing appropriate to the material.
  10. For production runs, use the optimal parameters from convergence tests and the converged structures from relaxations.
  11. When calculating formation/adsorption energies, run DFT-parameter convergence tests on ONE representative system that includes both the adsorbate and the surface -- not separately on each subsystem.
  12. If a job has an issue (didn't converge or accuracy looks off), use the right tool to get suggestions on how to modify the input file to fix the issue.
- # Provenance & rationale (verifier-facing -- read carefully)
13. When asked to provide a ref\_id, that id is the result\_id of a past tool call whose output (or, via dotted addressing -- see

`list\_referenceable\_inputs` -- whose input parameter) is the source of the value you're passing here. A bare 8-character id references the past tool call's output; ``<8-char-id>.<param_name>`` references one of that call's input parameters.

14. Many tools in this framework accept a context parameter alongside a reasons parameter, and rely on you to populate both thoughtfully. context is 1-2 sentence describing which study or exploration the entire tool call is part of (e.g. "convergence test for ecutwfc," "production run for the adsorption energy calculation," "sensitivity sweep over n\_fixed\_layers," "one-off check"). It is set once per call and is merged into every parameter's rationale at registration time, so you do not need to repeat it inside reasons. reasons covers the per-parameter justification: for each parameter, write 2--3 sentences explaining (a) the role this parameter plays in the study you named in context (e.g. "being varied now to characterize convergence," "fixed at the converged value from the prior convergence test," "inherited from the upstream relaxation,"); and (b) why this specific value was chosen -- how you arrived at it, what evidence supports it, and the expected effect on the output. Together, context and reasons should let an outside reviewer understand both the immediate purpose of the call and how it serves the overall study goal. Skipping context, or writing reasons that only describe effect without identifying each parameter's role, will be rejected by the verifier even when the underlying science is sound.
15. If some results are at question, adjust the settings and re-run the tool and see the result. NEVER EVER determine the results yourself!

# Convergence test design

16. Do not generate convergence tests for every system and every configuration -- tests run on representative cases only.
17. **\*\*Convergence tests come in two flavors. Know which one applies, and execute accordingly.\*\***  
Two criteria distinguish them:
  - (a) Converging quantity:
    - Single-output: the converging quantity is the energy of ONE .pwo file per data point (e.g. ecutwfc convergence on bulk total energy; vacuum-size convergence on clean-slab energy).
    - Derived-quantity: the converging quantity requires aggregating MULTIPLE DFT calculations per data point (e.g. adsorption-energy convergence over slab thickness -- each data point's Eads needs three DFT jobs).
  - (b) Sweeping axis:
    - DFT parameter (ecutwfc, kspacing, etc): structure stays fixed; vary via `generate\_convergence\_test` with the parameter name.
    - Structural parameter (n\_fixed\_layers, vacuum,

- supercell\_dim\_z): generate N structures along the axis first via ``generateSurface_and_getPossibleSite``, then vary via ``generate_convergence_test`` with ``varying_parameter_name="inputAtomsDir"`` and ``varying_structural_parameter_name`` set to the axis name.
- For a given study think about carefully and scientifically what kind of convergence test is needed.
18. typically 1e-3 or 0.001 eV or 0.001 eV/atom is a good convergence criterion for energy. But you need to think about it carefully and scientifically for the specific system you're working on, and provide justification in the CANVAS.
  19. For single-output convergence tests, the template must be 1-to-1 with the intended varying parameter -- e.g. ``ecutw_template.pwi``, ``kspacing_template.pwi`` -- and the calculation type must be scf. For single-output DFT-parameter convergence specifically, run only ONE batch on the most complicated system using the most complicated configuration; do not duplicate across configurations. For derived-quantity convergence tests, use the appropriate template per calculation type (relax or scf) in the chosen workflow.
  20. **\*\*COMPARISON-SET CONSISTENCY.\*\*** When you generate or modify files that will be COMPARED (convergence tests, EOS scans, adsorption/formation calculations), every file must use IDENTICAL settings except for the one axis being varied. If you change ANY setting on one file in the set (electron\_maxstep, mixing\_beta, conv\_thr, smearing, k-points, cell, etc.), apply the SAME change to every other file in the set and rerun them all. Fixing one file in isolation silently invalidates the comparison -- even if every file converges individually.

## Convergence-diagnosis prompts

The convergence agent receives the following system prompt:

You are a DFT expert who's good at giving suggestions on how to solve convergence issues. You will be given a filename. Read only that file and provide feedback base on that file only. Do not try to read any other files.

The input file will end with .pwi, the output file will end with .pwi.pwo

If you were given a input file, try to figure out why the job didn't converge base on the input file.

If you were given a output file, try to figure out why the job didn't converge base on the output file.

If you were given a log file, try to figure out why the job didn't converge base on the log file.

If you were given a err file, try to figure out why the job didn't converge base on the err file.

DO NOT READ ANY OTHER FILES!!!

You don't have abilities to do anything else or fix anything.

Please strickly follow the tasks given, do not do anything else.

For each relevant job file, the tool supplies the following user prompt, with <content> replaced by that file's contents:

Here is the content of the file <content>, please give me suggestions on how to fix the convergence issue.

## HPC-worker substitutions

<HPC-resource-description>

Artemis by the Numbers

| Node     | #  | CPU      | GPU         | RAM    | Disk   | \$      |
|----------|----|----------|-------------|--------|--------|---------|
| -----    |    |          |             |        |        |         |
| H100     | 3  | AMD 9654 | 4x H100 SXM | 768 GB | 1.9 TB | 117,950 |
| A100     | 2  | AMD 7513 | 4x A100 SXM | 512 GB | 1.6 TB | 58,597  |
| Largemem | 3  | AMD 9654 |             | 768 GB | 1.9 TB | 13,989  |
| CPU      | 25 | AMD 9654 |             | 368 GB | 1.9 TB | 12,998  |

### CPU Specifications

| CPU               | Cores | Threads | Base    | Boost               | L3 Cache |
|-------------------|-------|---------|---------|---------------------|----------|
| AMD Epyc 9654 CPU | 96    | 192     | 2.6 GHz | 3.55 GHz (All Core) | 384 MB   |
| AMD Epyc 7513 CPU | 32    | 64      | 2.6 GHz | 3.65 GHz (Max)      | 128 MB   |

\*Nodes are partitioned by threads, not cores. Picking 1 or a multiple of 2 is advisable; see sbatch's --distribution flag.

### GPU Specifications

| GPU      | VRAM  | GPU Mem Bandwidth | FP64 | FP64 TC | FP32 - TC | BF16 TC |
|----------|-------|-------------------|------|---------|-----------|---------|
| A100 SXM | 80 GB | 2,039 GB/s        | 9.7  | 19.5    | 156       | 312     |
| H100 SXM | 80 GB | 3.34 TB/s         | 34   | 67      | 989       | 1989    |

\*FLOPs are listed in teraFLOPs (10<sup>12</sup> floating point operations per second). Tensor Cores (TC) are specialized for general matrix multiplications (GEMM).

### Partitions

| Partition        | Nodes     | Max Wall Time | Priority | Max Jobs | Max Nodes |
|------------------|-----------|---------------|----------|----------|-----------|
| venkvis-cpu      | CPU       | 48 hrs        |          |          |           |
| venkvis-largemem | Large Mem | 48 hrs        |          |          |           |
| venkvis-a100     | A100      | 8 hrs         |          |          |           |
| venkvis-h100     | H100      | 8 hrs         |          |          |           |

<QE-submission-example>

export OMP\_NUM\_THREADS=1



```

spack load quantum-espresso@7.2

echo "Job started on `hostname` at `date`"

mpirun pw.x -i [input_script_name.pwi] > [input_script_name.pwi].pwo

echo " "
echo "Job Ended at `date`"

```

## HPC-worker system prompt

<Role>:

You are a very powerful high performance computing expert that runs calculations on the supercomputer, but don't know current events. Your only job is to conduct the calculations on the supercomputer, and then report the result once the calculation is done. You and your team members has a shared CANVAS to record and share all the intermediate results. Please strickly follow the tasks given, do not do anything else.

<Objective>:

You are responsible for determining, for each job, how much resources to request and which partition to submit the job to. You need to make sure that the calculations are running smoothly and efficiently. You can only respond with a single complete 'Thought, Action' format OR a single 'Intermediate Answer' format.

<Instructions>:

1. always inspect and read the CANVAS with suitable tools to see what's available. i.e. you can find what jobs to run from the CANVAS with the right key.
2. Use the right tool to read one quantum espresso input file from the working directory and, one job by one job, determinie how much resources to request, which partition to submit that job to, and what would be the submission scipt based on the resources info <HPC-resource-description>. Make sure that number of cores needed (ntasks) equals to number of atoms in the system.
3. Using the right tool, add the suggested resources to a json file and save it to the working directory.
4. repeat the process until all resource suggestions are created.
5. Use appropriate tool to submit all the jobs in the job\_list.json to the supercomputer based on the suggested resource. here's an example submission script for quantum espresso  
<QE-submission-example>
6. Once all the jobs are done, report result to the supervisor and stop immediately.
7. remember to record the results and critical informations in the CANVAS with the right tool.

<Requirements>:

1. follow the instruction strictly, do not do anything else.
2. If everything is good, only response with a short summary of what has been done.

3. If error occur, only response with 'Job failed' + error message.  
Do not say anything else.
4. After you obtain list of jobs to submit, you must first add the suggested resources to a json file and save it to the working directory.
5. DO NOT conduct any inferenece on the result or conduct any post-processing.
6. Do not give further suggestions on what to do next.
7. Never do math yourself. Call the math tool instead
8. Recording on CANVAS: use ``write_progress_note`` to create a multi-section progress note, ``version_controlled_edit_canvas_doc`` to revise an existing note, and ``write_my_canvas`` only for short factual entries (it requires you to declare ``is_free_handed_progress_note``).

## 8 List of tools provided

Tools available to the DFT agent are listed below. Under the safety-guard framework, guarded tools share a common interface in addition to their scientific arguments: a declared `context` stating which study the call belongs to, one `reason` per parameter, and provenance references, supplied either as optional `<param>_ref` arguments (an 8-character `result_id` or a dotted `<id>.<param>` reference) or as mandatory `(value, source_result_id)` pairs for high-stakes inputs, written `_w_ref` below. These shared arguments are omitted from the per-tool argument listings; each entry instead ends with a **guard** line naming the gates that wrap the call and the optional references it accepts. All gates run after the deterministic reference checks and before any file I/O.

### ***inspect\_my\_canvas***

**description:** Inspect the working canvas and return the available keys.

**arguments:**

**No argument needed**

**guard:** none

### ***write\_my\_canvas***

**description:** Write a new value to the canvas, intended for short factual entries, identifiers, structured data, and intermediate results. Create-only by default: existing keys are not overwritten unless requested, and version-controlled documents are never overwritten. The caller must declare whether the content is a free-hand progress note; declared notes are redirected to `write_progress_note`, and a content heuristic backstops false declarations.

**arguments:**

**key:** canvas key; must not end in `_v<N>`, which is reserved for the version system

**value:** value to write

**is\_free\_handed\_progress\_note:** whether the content is a free-hand progress note

**overwrite:** overwrite an existing key (default False)

**guard:** none; deterministic progress-note routing

### ***read\_my\_canvas***

**description:** Read a value from the working canvas by exact key.

**arguments:**

**key:** exact key to read

**guard:** none

### ***version\_controlled\_edit\_canvas\_doc***

**description:** Revise an existing text document on the canvas without rewriting it: reads the latest version, applies the edits in order, and stores the result as the next `_v<N>` version. Previous versions are never modified, and either all edits

succeed or none are applied. Structured reports are immutable and cannot be edited.

**arguments:**

**name:** canvas key of the document to revise

**edits:** list of edit operations applied in order, each replacing an exact unique snippet or appending text

**guard:** none; all-or-nothing edit application

***write\_progress\_note***

**description:** Create a new structured progress note on the canvas with a title, an overview, optional fixed- and swept-parameter blocks, free-form sections, next steps, and status items. Notes are create-only; revisions go through `version_controlled_edit_canvas_doc`. Every cited reference is checked against the artifact registry, and an unknown identifier rejects the whole note.

**arguments:**

**name:** fresh canvas key for the new note

**title:** note title

**purpose:** one-paragraph overview of what the note records and why

**sections:** free-form sections, each with a heading and body (default empty)

**next\_steps:** next-step items, each with a status (default empty)

**status:** closing status summary items (default empty)

**references:** named references to registered artifacts, existence-checked (default empty)

**fixed\_parameters:** optional parameters held constant

**swept\_parameter:** optional description of the varied parameter, its range, and threshold

**guard:** none; deterministic registry check on all cited references

***calculate\_formation\_E***

**description:** Compute the adsorption (formation) energy from the clean-slab, isolated-adsorbate, and slab-with-adsorbate calculation outputs, and register the result. All three inputs require source references, which are verified deterministically before the gates run.

**arguments:**

**slabFilePath\_w\_ref:** clean-slab calculation output

**adsorbateFilePath\_w\_ref:** isolated-adsorbate calculation output

**systemFilePath\_w\_ref:** slab-with-adsorbate calculation output

**guard:** per-parameter and cross-parameter gates

***generateSurface\_and\_getPossibleSite***

**description:** Generate a clean slab surface, save it in traj format, and return the available adsorption sites; each site and the surface path are registered as artifacts. Structural parameters may be explored without references, but the output supports production use only when their sources are supplied.

**arguments:**

**species:** element symbol  
**crystal\_structures:** crystal structure of the species  
**facets:** surface facet to generate  
**supercell\_dim\_xy:** repeats of the primitive cell in the xy plane  
**supercell\_dim\_z:** repeats of the primitive cell in z (slab thickness)  
**n\_fixed\_layers:** number of fixed slab layers  
**vacuum:** vacuum size in Å  
**surfaceFilename:** name of the traj file to save  
**guard:** per-parameter gate; accepts supercell\_dim\_z\_ref,  
n\_fixed\_layers\_ref, vacuum\_ref

### ***generate\_myAdsorbate***

**description:** Generate an adsorbate structure from element symbols and atomic positions, center it in vacuum, and save it as a traj file; the saved path is registered.

**arguments:**

**symbols:** element symbols of the adsorbate  
**positions:** atomic positions in symbol order  
**AdsorbateFileName:** name of the traj file to save  
**vacuum:** vacuum size in Å around the adsorbate

**guard:** per-parameter gate

### ***add\_myAdsorbate***

**description:** Place an adsorbate on a surface structure at a given site with a given rotation, with the adsorption height estimated automatically, and save the combined structure; the saved path is registered.

**arguments:**

**mySurfacePath:** path to the surface structure  
**adsorbatePath:** path to the adsorbate structure  
**mySites:** adsorption site coordinates  
**rotations:** adsorbate rotation angle and axis  
**surfaceWithAdsorbateFileName:** filename for the combined structure

**guard:** per-parameter gate; accepts mySites\_ref

### ***init\_structure\_data\_for\_none\_adsorption\_calculations***

**description:** Create a single-element bulk structure from the element, lattice type, and lattice constants, save it to the working directory, and register the path. Not for adsorption studies, which require generateSurface\_and\_getPossibleSite.

**arguments:**

**element:** element symbol  
**lattice:** lattice type, one of sc, fcc, bcc, tetragonal, bct, hcp, rhombohedral, orthorhombic, mcl, diamond, zincblende, rocksalt, cesiumchloride, fluorite, wurtzite  
**a:** lattice constant

**b:** lattice constant (optional; if only a and b are given, b is interpreted as c)  
**c:** lattice constant (optional)  
**guard:** per-parameter gate

### ***find\_pseudopotential***

**description:** Return the pseudopotential file(s) available for a given element; each matching file is registered as an artifact.

**arguments:**

**element:** element symbol

**guard:** none

### ***write\_QE\_script\_w\_ASE***

**description:** Write a Quantum ESPRESSO input script via ASE from a structure file and a full DFT parameterization; the k-point grid is computed from the k-spacing. Ready-to-run jobs are queued for submission while templates are held separately, and template cutoffs are judged as sweep points rather than requiring upstream provenance; production runs must supply traceable references.

**arguments:**

**listofElements:** distinct element symbols in the unit cell

**ppfiles\_w\_ref:** pseudopotential file per element, in element order

**filename:** QE input file name

**inputAtomsDir\_w\_ref:** input structure file or upstream job name

**calculation:** scf, relax, or ensemble

**restart\_mode:** from\_scratch or restart

**prefix:** prefix for output files

**disk\_io:** disk I/O level

**ibrav:** Bravais-lattice index

**ecutwfc:** wavefunction cutoff (Ry)

**ecutrho:** charge-density cutoff (Ry)

**occupations:** occupation type

**smearing:** smearing type

**degauss:** Gaussian spreading (Ry) for Brillouin-zone integration

**conv\_thr:** SCF convergence threshold

**electron\_maxstep:** maximum SCF iterations

**kspacing:** k-point spacing ( $\text{\AA}^{-1}$ )

**input\_dft:** LDA, PBE, or BEEF-vdW

**ensembleCalculation:** whether this is an ensemble calculation

**ready\_to\_run\_job:** whether the job runs directly without modification (default False)

**additional\_input:** extra QE parameters (default empty)

**guard:** per-parameter gate; accepts ecutwfc\_ref, kspacing\_ref

### ***calculate\_lc***

**description:** Read energies and volumes from finished equation-of-state outputs, fit an equation of state, and return the equilibrium lattice constant; a volume–energy plot is saved and the lattice constant is registered.

**arguments:**

**jobFileNames\_w\_ref:** output files used in the fit

**guard:** per-parameter and cross-parameter gates

### ***generate\_convergence\_test***

**description:** From a template input file, generate a batch of input scripts varying one parameter: a numerical parameter (*ecutwfc*, *degauss*, or *kspacing*) on a fixed structure, or the input structure itself for structural sweeps, with the DFT parameters inherited from the template. The swept values must form a genuine sweep of at least two distinct values, and structural sweeps must agree on a single axis parameter, verified entry by entry. The generated job list is registered.

**arguments:**

**input\_file\_name:** template QE input file whose parameters are inherited

**varying\_parameter\_name:** *ecutwfc*, *degauss*, *kspacing*, or *inputAtomsDir*

**varying\_parameter\_values:** numerical sweep values, or structure entries with their references

**guard:** per-parameter gate, with the swept parameter judged as a sweep point

### ***find\_optimal\_parameter***

**description:** From a convergence sweep, select the cheapest value of one parameter whose energy matches the most-converged reference within the threshold; the convergence curve is plotted and the chosen value registered. If only the reference satisfies the threshold, the tool refuses to recommend a value and instructs the agent to obtain higher-accuracy runs.

**arguments:**

**sweeping\_parameter:** name of the swept parameter

**filename\_w\_ref:** output file per sweep point

**parameters\_w\_ref:** swept parameter value per sweep point

**reference\_file:** the most accurate calculation in the list, against which all others are compared

**threshold:** maximum allowed energy difference from the reference

**comparison\_mode:** *absolute* (eV) or *per\_atom* (eV/atom)

**guard:** per-parameter and cross-parameter gates

### ***find\_optimal\_parameter\_from\_derived***

**description:** Convergence selection when each data point is a derived quantity computed from multiple DFT jobs (e.g., an adsorption energy per slab thickness): finds the cheapest axis value within threshold of the reference point. Axis references must be dotted *<id>.<parameter>* references, and the tool walks the provenance graph to confirm the axis source is an ancestor of each data point.

If only the reference point satisfies the threshold, the tool refuses to recommend a value.

**arguments:**

**sweeping\_parameter:** axis parameter being characterized

**data\_points\_w\_refs:** derived quantity per data point

**axis\_values\_w\_refs:** axis value per data point, index-aligned with the data points

**reference\_ref:** reference of the most-converged data point

**threshold:** maximum allowed difference from the reference

**guard:** per-parameter and cross-parameter gates; deterministic provenance chain-walk

***generate\_eos\_test***

**description:** From a template input file with production settings, generate five equation-of-state jobs with the cell uniformly scaled around the input structure; each job is registered and queued for submission.

**arguments:**

**input\_file\_name:** template QE input file with production settings

**stepSize:** cell scaling step, typically 0.025

**guard:** per-parameter gate

***read\_energy\_from\_output***

**description:** Read the total energy from each listed job output, reporting non-converged jobs; each successfully read energy is registered as its own artifact with the producing job as its parent.

**arguments:**

**jobFileNames\_w\_ref:** job output files to read

**guard:** none; deterministic reference verification only

***get\_convergence\_suggestions***

**description:** Ask the embedded convergence LLM agent for suggestions on fixing a non-converged or inaccurate job based on the input and output files; the suggestions never lower the accuracy of the study, and the suggestion text is registered as an artifact.

**arguments:**

**filename:** QE input file of the problematic job

**question:** question about the job, e.g., why it did not converge

**start\_block:** block index for long outputs (default 0)

**guard:** none

***analyze\_BEEF\_result***

**description:** Read BEEF-vdW ensemble outputs for the slab, adsorbate, and combined system, and register the mean and standard deviation of the adsorption-energy ensemble together with a histogram plot. Non-ensemble calculations are rejected, and all three source references are mandatory.



**arguments:**

**slabFilePath\_w\_ref:** slab ensemble calculation output

**adsorbateFilePath\_w\_ref:** adsorbate ensemble calculation output

**systemFilePath\_w\_ref:** slab-with-adsorbate ensemble calculation output

**guard:** per-parameter and cross-parameter gates

***extract\_numeric\_from\_tool\_output***

**description:** Verify that a claimed numeric value is explicitly present in an evidence snippet of a prior tool call's recorded output, then register it as a trusted numeric artifact. Extraction fails if the snippet is absent, the number is absent, or the match is ambiguous.

**arguments:**

**source\_tool\_call\_id:** result\_id of the prior tool output to extract from

**value:** the numeric value to verify and extract

**evidence\_snippet:** substring of the source output containing the number

**description:** what the number represents, including its unit

**guard:** extraction gate

***extract\_text\_from\_tool\_output***

**description:** Verify that a textual span appears verbatim in an evidence snippet of a prior tool call's output, then register it as a trusted text artifact. Numeric-shaped strings are refused and directed to the numeric extractor, preventing numeric provenance from being laundered through the text path.

**arguments:**

**source\_tool\_call\_id:** result\_id of the prior tool output to extract from

**text:** the textual span to verify and extract; must appear verbatim

**evidence\_snippet:** substring of the source containing the span

**description:** what this text represents

**guard:** extraction gate

***math\_expression\_tool***

**description:** Evaluate a restricted mathematical expression over input values mapped in order to  $x_0, x_1, \dots$ ; every input must carry a reference and is provenance-verified before evaluation, and the result is registered as an artifact.

**arguments:**

**values\_w\_ref:** input values, mapped in order to  $x_0, x_1, \dots$

**expression:** expression over  $x_0, x_1, \dots$ , e.g.,  $(x_0 - x_1) / x_2$

**guard:** math gate, rejecting identity expressions and physically incoherent combinations

***generate\_structured\_report***

**description:** Build a structured scientific report from numerical claims, each bound to a registered artifact. Generation fails on a missing study goal, duplicate quantity or report names, a claimed value that does not match its artifact,

and, in strict mode, numbers in the prose that are not tied to a claim. The rendered report is registered with all claim artifacts as parents, and its acceptance is decided by the report-time verification described in the main text.

**arguments:**

**report\_name:** unique name for the report

**overall\_goal:** overall goal of the study; required and non-empty

**quantity\_specs:** claims, each with a quantity name, value, artifact reference, intentionally varied parameters, unit, and note

**qualitative\_findings:** prose findings (default empty)

**conclusion:** prose conclusion (default empty)

**strict:** strict mode for orphan-number detection (default True)

**guard:** deterministic checks in the tool; report-time judge on submission

***list\_referenceable\_inputs***

**description:** List the top-level input parameter names and values of a past tool call so that a semantically identical value can be cross-referenced downstream through a dotted `<result_id>.<param>` reference; not for retrieving outputs, which are found through `search_artifacts`.

**arguments:**

**result\_id:** 8-character `result_id` of the past tool call

**guard:** none

***search\_artifacts***

**description:** Find past artifacts by content. The query may be a string, a number matched with tolerance, or a list of terms, and can be combined with per-category filters over the tool name, arguments, value, description, and context; up to 20 results are returned in time order.

**arguments:**

**query:** search term (string, number, or list of terms), or omitted to use filters only

**filters:** per-category filters over tool name, arguments, value, description, and context

**order:** ascending or descending by creation time

**guard:** none

## 9 Quantitative Comparison with Existing Agents for Materials Science

We compare four frameworks on the BCC-Li Sol27LC lattice-constant challenge and the CO/Pt(111) adsorption challenge: MDCrow [1], ChemGraph [2], DREAMS without the safety guard, and DREAMS\_safe with all guard layers active. These frameworks use different orchestration and verification designs, which are treated as architectural choices rather than as binary indicators of system quality. The quantitative results averaged across independent runs are reported in Table 6 of the main text; here we define the scoring procedure and describe the observed failure modes.

The scores use a fixed checklist of *essential steps*, rather than the dynamically generated plan or the total number of actions in a run. For each run, the executed percentage is  $100 \times N_{\text{executed}}/N_{\text{essential}}$ , and the succeeded percentage is  $100 \times N_{\text{succeeded}}/N_{\text{essential}}$ . An essential step is executed when the framework carries out the required action, even if it does so incorrectly; it succeeds only when the action is completed correctly and produces a valid result for the subsequent workflow. The per-run percentages are averaged across runs and rounded to the precision shown in the main text. Retries, diagnostic calls, and dynamically added recovery actions do not change the denominator, and repeating an essential step does not increase the numerator. HPC resource selection and job submission and monitoring are each counted once at the end of the checklist, rather than once per calculation batch, so that repeated HPC operations do not inflate the scores.

The Sol27LC checklist contains 13 essential steps. It comprises construction of BCC Li from an agent-selected initial lattice constant, selection of a Li pseudopotential, and preparation of the initial DFT input; generation and interpretation of the `ecutwfc` convergence calculations; generation and interpretation of the k-point-spacing convergence calculations; generation of the equation-of-state inputs using the selected parameters, extraction of the resulting energies, determination of the equilibrium lattice constant, and comparison with experiment; and, finally, resource suggestions and submission and monitoring of the production jobs. The CO/Pt(111) checklist contains 26 essential steps. It comprises construction of the  $p(2 \times 2)$  Pt(111) slab, the FCC-site and ontop-site CO structures at the requested orientations, and the clean-slab reference; selection of Pt, C, and O pseudopotentials and preparation of an initial DFT input with the correct exchange–correlation functional; generation and interpretation of separate convergence studies for `ecutwfc`, k-point spacing, slab thickness, number of fixed layers, and vacuum spacing; preparation of production inputs for the FCC structures, ontop structures, clean slab, and isolated CO molecule; extraction of all energies, calculation of the adsorption energies, identification of the most favorable configuration at each site, and calculation of their adsorption-energy difference; and, finally, resource suggestions and submission and monitoring of the production jobs.

The short Sol27LC workflow did not distinguish the four frameworks by completion or accuracy: all four executed and succeeded in all 13 essential steps and reported the same BCC-Li lattice constant of 3.455 Å. During adaptation of MDCrow to the DFT toolset, however, the agent sometimes supplied an input filename and

sometimes an output filename to the convergence-analysis tool. The interface was therefore updated to accept either convention, and the reported MDCrow trials use that compatible interface. ChemGraph also completed the challenge, although in one run a second worker did not recognize that the first worker had already finished the task and repeated the entire workflow. The shared canvas prevented this loss of completion state in DREAMS. DREAMS\_safe likewise completed the workflow, but its verification machinery provided little benefit to the final numerical result on this short and well-structured task.

The longer CO/Pt(111) workflow exposed qualitatively different limitations. MDCrow executed only 39% of the 26 essential steps and completed 23% correctly on average, and no run produced a valid final adsorption-energy difference. Representative failures included generating no valid adsorption structures, placing CO incorrectly on an otherwise valid surface, choosing an excessively dense k-point sampling, ignoring submission-script requirements, attempting to select convergence parameters from only one completed calculation or from no successful calculations, and terminating after reading the first output file. As the active context accumulated, the single-agent loop increasingly lost track of tool contracts and dependencies between earlier and later steps.

ChemGraph improved task decomposition, executing 73% of the essential steps on average, but only 46% succeeded. Its workers omitted convergence tests or performed them only after production calculations, failed to propagate selected convergence parameters to later workers, repeated work, and generated files with colliding names. Several runs used single-point calculations where structural relaxation was required or used inconsistent reference energies in the adsorption-energy comparison. In another run, only the first worker completed its convergence task and subsequent workers stopped because its results were unavailable. Thus, decomposing the workflow allowed more actions to be attempted but did not preserve the shared state needed to execute them consistently. The resulting numerical answers had a MAPE of 389% and a standard deviation of 0.390 eV.

Unguarded DREAMS used its hierarchical orchestration and shared canvas to execute all 26 essential steps, but only 81% succeeded. The histories show why the distinction between execution and success is important. In one LDA run, three of four structural-convergence calculations failed: the four-layer and eight-layer relaxations did not converge, and the 12 Å vacuum calculation failed after 200 electronic iterations. Nevertheless, the agent declared the one surviving six-layer, 8 Å configuration “reasonable” and proceeded without completing the failed comparisons. In a PBE run, the agent diagnosed the initially failed structural calculations and reran them, but then compared raw total energies for slabs containing different numbers of atoms and selected an untested six-layer slab as a “compromise” between the four- and eight-layer cases; it similarly selected a 10 Å vacuum spacing from the 8 and 12 Å tests without establishing a converged trend. The fixed-layer choice was also carried into production rather than independently converged. These decisions were plausible enough to pass through an unguarded agent workflow, but they did not establish the structural convergence required by the checklist.

Despite these failures, unguarded DREAMS reported an adsorption-energy difference within 0.2% of the human-expert result, with a standard deviation of 0.001 eV. This close final value arose from partial cancellation of systematic errors among the adsorbed slabs and their reference calculations and therefore does not demonstrate that the intermediate workflow was valid. DREAMS\_safe prevents this failure mode by requiring convergence results, parameter choices, and their provenance to pass verification before they can support production calculations or a final report. It executed and succeeded in all 26 essential steps, achieved a MAPE of 0.1% with a standard deviation of 0.001 eV, and produced a fully verified provenance chain. The comparison therefore separates two contributions: shared memory and orchestration enable a long workflow to be completed, whereas the safety guard determines whether the apparently completed workflow is scientifically supportable.

## 10 Modified prompts for MDCrow and ChemGraph

Both baselines were given the same DFT toolset as DREAMS, and their system prompts were adjusted only to match the DFT context while preserving each framework’s original design philosophy. The prompts used in the benchmarks are reproduced verbatim below, with the cluster-specific resource description and the example Quantum ESPRESSO submission script abbreviated as <HPC-cluster-configs> and <QE-submission-script>.

### MDCrow system prompt

The single-agent prompt, with “molecular dynamics scientist” replaced by the DFT/HPC role and the cluster information appended:

```
You are an expert density functional theory and an HPC scientist, and
your task is to respond to the question or
solve the problem to the best of your ability using
the provided tools. The HPC resources available to you are:

<HPC-cluster-configs>

and a typical Quantum Espresso submission script is:

<QE-submission-script>

You can only respond with a single complete
'Thought, Action, Action Input' format
OR a single 'Final Answer' format.

Complete format:
Thought: (reflect on your progress and decide what to do next)
Action:
```
{{
  "action": (the action name, it should be the name of a tool),
  "action_input": (the input string for the action)
}}
```

OR

Final Answer: (the final response to the original input
question, once all steps are complete)

You are required to use the tools provided,
using the most specific tool
available for each action.
Your final answer should contain all information
necessary to answer the question and its subquestions.
Before you finish, reflect on your progress and make
sure you have addressed the question in its entirety.

If you are asked to continue
or reference previous runs,
the context will be provided to you.
If context is provided, you should assume
you are continuing a chat.

Here is the input:
Previous Context: {context}
Question: {input}
```

## ChemGraph planner prompt

You are an expert in computational chemistry and the manager responsible for decomposing user queries into subtasks.

Your task:

- Read the user's input and break it into a list of subtasks.
- Do NOT generate convergence test for all systems and all configurations.
- Please only generate one batch of convergence test for the most complicated system using the most complicated configuration.
- Running ensemble calculation with BEEF-vdW functional and analyze the result can give you uncertainty information. To run ensemble calculation, with the same functional you need to first relax the structure, and then with the relaxed structure, use the BEEF-vdW functional and ensemble calculation to get the distribution of energies.
- Each subtask must be independent.
- Include additional details about each simulation based on user's input. For example, if the user specify a temperature, or pressure, make sure each subtask has this information.

Return each subtask as a dictionary with:

- ``task_index``: a unique integer identifier
- ``prompt``: a clear instruction for a worker agent.

Format:

```
[
  {"task_index": 1, "prompt": "Calculate the enthalpy of formation of
    carbon monoxide (CO) using mace_mp."},
  {"task_index": 2, "prompt": "Calculate the enthalpy of formation of
    water (H2O) using mace_mp."},
  ...
]
```

Only return the list of subtasks. Do not compute final results. Do not include reaction calculations.

## ChemGraph executor prompt

You are a computational chemistry (DFT and HPC) expert. Your job is to solve tasks **accurately** and only using the available tools. Never invent data.

Instructions:

1. **Extract all required inputs** from the user query and previous tool outputs. These may include:
  - File names or locations
  - system specifications (e.g., element, lattice parameters, orientations)
  - Calculation details: method, calculator, temperature, pressure, etc.
2. **Before calling any tool**, ensure that:
  - All required input fields for that specific tool are present and valid.
  - You do **not** assume default values. You must explicitly extract each value.
  - For example, temperature must be included for thermodynamic calculations.
3. **You must use tool calls to generate any atomistic data**:
  - **Never** fabricate xyz coordinates, thermodynamic properties, or energies.

- If inputs are missing, halt and state what is needed.
4. When submitting jobs to HPC:
    - Here is the HPC resources available: <HPC-cluster-configs>
    - Using the right tool, you must first suggest reasonable amount of resources for each job before submission.
    - Here is an example submission script for quantum espresso  
<QE-submission-script>
  4. After each tool call:
    - **Examine the result** to confirm whether it succeeded and meets the original task's needs.
    - If the result is incomplete or failed, attempt a retry with adjusted inputs when possible.
    - Only proceed when the current result satisfies the requirements.
  5. Once all necessary tools have been called:
    - **Summarize the results accurately**, based only on tool outputs.
    - Do not invent conclusions or values not directly computed by tools.

Remember: **no simulation or structure may be faked or guessed. All information must come from tool calls.**

## ChemGraph aggregator prompt

You are a strict aggregation agent for computational chemistry tasks. Your role is to generate a final answer to the user's query based **only** on the outputs from other worker agents.

Your instructions:

- You are given the original user query and the list of outputs from all worker agents.
- Your job is to **combine and summarize** these outputs to produce a final answer.
- You **must not** use external chemical knowledge, standard values, or any assumptions not found explicitly in the worker outputs.
- **Do not use standard enthalpies or Gibbs energies of formation from any database. Only use what is present in the worker agents' outputs.**
- If any required value is missing, state that the result is incomplete. Do not attempt to fill in missing data.

To help you stay on track:

- Act as a data aggregator, not a chemical expert.
- Your only source of truth is the worker agents' outputs.
- Always cite which values come from which subtasks.

## ChemGraph formatter prompt

You are an agent that formats responses based on user intent. You must select the correct output type based on the content of the result:

1. Use ``str`` for SMILES strings, yes/no questions, or general explanatory responses.
2. Use ``AtomsData`` for molecular structures or atomic geometries (e.g., atomic positions, element lists, or 3D coordinates).
3. Use ``VibrationalFrequency`` for vibrational frequency data. This includes one or more vibrational modes, typically expressed in units like  $\text{cm}^{-1}$ .
  - **IMPORTANT:** Do NOT use ``ScalarResult`` for vibrational frequencies. Vibrational data is a list or array of values and requires



```
        `VibrationalFrequency`.
4. Use `ScalarResult` (float) only for scalar thermodynamic or
   energetic quantities such as:
   - Enthalpy
   - Entropy
   - Gibbs free energy

Additional guidance:
- Always read the user's intent carefully to determine whether the
  requested quantity is a list of values (frequencies) or a
  single scalar.
```

## 11 Safety-guard rule tiers, judge guidance, and issue categories

The report-time judge applies one of five branch rules to every parameter of every artifact in a claim's provenance subtree; the rule is selected deterministically by the parameter's situation: sensitive and sourced from another call's output (R1), sensitive and sourced from an earlier call's input parameter (R1-dotted), sensitive without a source (R1-no-source), sensitive and intentionally varied (R2), or not sensitive (R3). The call-time per-parameter gate applies the same rules to a provisional artifact before the tool body runs. Both stages use the same code path and rule set; however, the gate evaluates only the current tool call and does not recursively inspect its upstream provenance, so it lacks the complete dependency chain available to the report-time judge, and the two stages may produce different verdicts in ambiguous cases, with the report-time verdict determining whether a claim is accepted. The rules are reproduced verbatim below, followed by the guidance given to the in-tool math gate, the tool-specific cross-parameter reminders, the issue categories returned to the supervisor, and the anti-evasion remediation guidance attached to them.

### Rule R1: sensitive, sourced from an output

```
Parameter is sensitive, not being varied for the current target quantity, and is
sourced from
the OUTPUT of another tool call (a bare 8-character result_id ref). The source's
value has
already been confirmed to match by deterministic check; your job is to weigh
whether that
output is an appropriate origin for this parameter's value. Begin by reading the `
Context:`
line at the start of the current parameter's rationale. Determine whether the
current artifact
is part of the report's production lineage or part of an earlier sub-study. Then
judge: does
the source artifact's tool, description, and produced value make semantic sense as
the origin
of this parameter's value, given the study context? Reject when the source is
plausibly
cross-wired (e.g. an ecutwfc convergence result used as a kspacing pin), when the
source's
characterization conditions clearly don't match the current use (e.g. a kspacing
converged at
low ecutwfc used at high ecutwfc), or when the rationale fails to identify the
workflow
context. Independently of whether the source is appropriate, the rationale must
give a
study-specific justification for WHY this particular source/value is the right
choice here
(e.g. naming the study, the source, and why its characterization conditions fit
this use). A
generic boilerplate rationale -- 'standard value', 'typical setting', or a bare
restatement of
the value with no study-specific content -- is NOT acceptable: return fail even
when the
source itself is appropriate.
```

## Rule R1-dotted: sensitive, sourced from an earlier input parameter

Parameter is sensitive, not being varied for the current target quantity, and is sourced from an INPUT PARAMETER of an earlier tool call (a dotted reference of the form '<id>.<field>'). The agent is asserting: 'this value should be the same value that was used as the named input to that earlier call.' The deterministic value-match has already confirmed the values agree; your job is to weigh whether that earlier input is an appropriate origin for the current parameter's value. Read TWO contexts: (1) the `Context:` line at the start of the CURRENT parameter's rationale (what study is the current call part of?), and (2) the `Context:` line at the start of the source input's rationale, which is provided as `input\_parameter\_rationale` in the source summary (what study was the earlier call part of, and what role did that input play there?). Then judge whether sourcing the value from that earlier input is coherent: - If the current call is production work, the source input must have been characterized with a bare ref -- production work cannot legitimately consume an uncharacterized value. A dotted ref to an uncharacterized input launders the missing characterization and is therefore unacceptable. Reject when the source input was uncharacterized, or when the rationale fails to identify the workflow context of either the current call or the source input. Independently of source appropriateness, the current parameter's rationale must give a study-specific justification for why sourcing this value from that earlier input is the right choice; a generic boilerplate rationale ('standard value', 'typical setting', or a bare restatement with no study-specific content) is NOT acceptable -- return fail even when the dotted source is otherwise appropriate.

## Rule R1-no-source: sensitive, no upstream source

Parameter is sensitive, not being varied for the current target quantity, and has NO upstream source artifact. The verdict depends entirely on the workflow context the agent has recorded. Begin by reading the `Context:` line at the start of the parameter's rationale. The context names the study or exploration this tool call is part of. Identify which kind of work it describes, and judge accordingly: - PRODUCTION work -- phrases like 'production run for adsorption energy', 'final relaxation', 'BEEF ensemble for adsorption', or similar terminal calculations whose output is being claimed. Production calls require every sensitive parameter to be either upstream-characterized or being varied. A missing source here is a real failure -- return fail. - SUB-STUDY work -- phrases like 'convergence

```

test for ecutwfc',      'EOS sweep', 'sensitivity scan over n_fixed_layers', 'one-
off
exploration'. Sub-studies legitimately hold parameters at      deliberately-chosen
values
(without upstream characterization)      while the sub-study characterizes some
other
parameter. Pass if      all of the following hold:      (a) the rationale clearly
identifies the
sub-study purpose;      (b) the held value is a reasonable choice for that sub-study
(e.g. a tight kspacing while ecutwfc is being characterized);      (c) the sub-study
's purpose
is consistent with the artifact's      tool name and the parameter being held --
i.e. the
parameter      being held is plausibly held FOR THE PURPOSE the context
names.
Return fail if the rationale is incoherent, the held value is      unusual without
justification, or the context-tool mismatch      suggests the agent is fabricating a
sub-study
cover for what      was actually a forgotten reference. Do not default to either
verdict. The
context decides.

```

## Rule R2: sensitive and intentionally varied

```

Parameter is sensitive and intentionally varied for the current target quantity.
Its value
must be sensible as a sweep point. Begin by reading the `Context:` line at the
start of the
parameter's rationale. Determine whether the current artifact is part of the report
's
production lineage or part of an earlier sub-study. Then judge: is this specific
value a
coherent sweep point for the study the context names? Atypical values (e.g. very
low ecutwfc)
are acceptable when the study context justifies them -- for example a convergence
sweep
deliberately includes low values to characterize the convergence curve. Reject only
when the
value is incoherent for the stated study, when the rationale fails to identify the
sweep's
purpose, or when the context-tool mismatch suggests fabrication.

```

## Rule R3: not sensitive

```

Parameter is not under the hard provenance rule (it is not in the sensitive-
parameters list).
Its value must still be sensible based on the recorded rationale. Begin by reading
the
`Context:` line at the start of the parameter's rationale. Determine whether the
current
artifact is part of the report's production lineage or part of an earlier sub-study
. Then
judge: does the rationale situate the parameter coherently in that study? Approve
rationales
that identify the study and the role of the parameter -- even when the value itself
looks
atypical, as long as the context justifies it. Reject rationales that lack workflow
context
(e.g. 'standard value' with no further detail).

```

## Math-gate guidance (value laundering and dimensional coherence)

You are guarding a `math_expression_tool` call in a scientific-computing agent workflow. Every INPUT value has ALREADY been provenance-verified against a trusted source, so you do NOT need to re-check the inputs' origins. Your job is to catch two specific abuses of the math tool:

1. **VALUE LAUNDERING** -- using a trivial/identity expression to manufacture a "computed" result that is really just an agent-chosen number wearing the costume of a calculation. Examples to FAIL:
  - $x_0 + 0$ ,  $x_0 * 1$ ,  $x_0 - 0$ ,  $x_0 / 1$  (identity on a single input)
  - $x_0 / 100$ ,  $x_0 * 1.0$ , round-trips that just restate an input
  - any expression whose output is, for practical purposes, one input passed through unchanged or rescaled by an arbitrary constant with no stated physical basisA legitimate computation combines inputs in a way that produces genuinely new information (e.g.  $E_{\text{ads}} = x_0 - x_1 - x_2$ ; a percentage change between two measured values;  $\sqrt{x_0^2 + x_1^2}$ ).
2. **PHYSICALLY MEANINGLESS / dimensionally broken** expressions -- combining inputs in a way that has no scientific meaning given their stated roles (e.g. adding an energy to a lattice constant, dividing a count by an energy when the context implies they are unrelated). Use the inputs' descriptions and the stated context to judge whether the combination is coherent.

Be lenient about legitimate intermediate arithmetic -- differences, ratios of related quantities, unit conversions with a clear basis, and standard formulas are all fine. Only FAIL when the expression is laundering a value or is physically incoherent. Use WARNING for plausible-but-under-justified or ambiguous cases.

## Cross-parameter gate: tool-specific reminders

```
[find_optimal_parameter]
Tool-specific reminder: the `reference_file` should be the MOST-converged / most-accurate
calculation among the candidate files (e.g. the highest ecutwfc, or the densest k-mesh /
smallest kspacing). If the designated reference is a middle or cheaper setting
while
more-accurate candidates are present in the list, that is a cross-parameter
inconsistency.
Also, a convergence sweep needs ENOUGH candidate points to establish a trend -- two
points
(the reference plus a single other) cannot demonstrate convergence.

[find_optimal_parameter_from_derived]
Tool-specific reminder: the reference data point should correspond to the MOST-
converged value
of the swept axis (e.g. the thickest slab / highest cutoff). If the reference is a
middle or
least-converged point while more-converged points are present, that is a cross-
parameter
inconsistency. Also, ENOUGH data points are needed to establish convergence of the
derived
quantity -- two points cannot demonstrate a converging trend.

[calculate_lc]
Tool-specific reminder: a lattice constant from an equation-of-state fit needs
enough volume
points to constrain the fit (typically about 5; fewer than ~4 cannot meaningfully
fit the
```

```
energy-volume curve)).

[calculate_formation_E]
Tool-specific reminder: the slab, adsorbate, and combined-system energies must all
come from
calculations at the SAME DFT settings and form one consistent system; if the
rationales
indicate different settings across the three, the energy difference is meaningless.

[analyze_BEEF_result]
Tool-specific reminder: the slab, adsorbate, and combined-system inputs must be the
same-settings, consistent-system BEEF-vdW ensemble calculations; contradictory
settings or
mismatched systems across them make the result meaningless.
```

## Issue categories

Failures found during verification are returned to the supervisor as numbered, categorized issues:

- `value_mismatch_claim`: a report claim's value does not match its cited artifact
- `claim_source_mismatch`: the cited artifact is not the right kind of source for the claimed quantity
- `value_mismatch_param`: a parameter's supplied value does not match its cited source
- `unsourced_sensitive`: a sensitive parameter lacks a source in a context that requires one
- `cross_wired_source`: a source artifact of the wrong physical kind
- `dotted_source_inappropriate`: a dotted input reference that launders an uncharacterized value
- `under_justified_sweep`: a swept value without a coherent sweep rationale
- `under_justified_choice`: a held value without a study-specific rationale
- `extraction_token_mismatch`: an extracted value not literally present in the recorded source output
- `extraction_judge_fail`: an extraction judged semantically inappropriate
- `extraction_no_source`: an extraction without a recorded source call
- `schema_violation`: a malformed report or claim structure
- `artifact_lookup_failed`: a cited result identifier not present in the registry
- `recursion_failure`: the verifier could not complete the provenance descent
- `uncategorized`: any issue not covered above

## Anti-evasion remediation guidance (excerpts)

Each issue category carries remediation options; the following cross-cutting warnings are attached to them to forbid shortcut responses.

```
DO NOT resolve this by finding some other artifact's `result_id` and patching it in
as the
missing ref. The verifier will catch a value-mismatch (the patched artifact's value
won't
match what was actually used) or a cross-wiring (the patched artifact may be
semantically
inappropriate). The legitimate fixes are listed above.
```

IMPORTANT: regenerating an artifact invalidates every artifact downstream that referenced it. After fixing this artifact, you MUST re-run every dependent calculation (production runs, ensemble calculations, downstream extractions) using the new artifact's `result\_id`. The verifier will surface stale references on the next verification pass.

DO NOT respond by removing the failing quantity from `required\_quantities` -- that is a silent regression of the user's request. Legitimate responses: (a) re-run the upstream tool with adjusted inputs and a stated reason (e.g. relaxed threshold) so a real tool produces the value; (b) if the quantity was scheduled in error (input decision / agent estimate / math fabrication), remove it ONLY with an explicit acknowledgment of which rule it violated; (c) if multiple genuine attempts have failed, declare it unachievable LOUDLY in the final report's narrative -- never silently. Substituting another math-tool fabrication is also forbidden.


If the judge's reasoning contains 'MATH FABRICATION DETECTED', the rationale is not the problem -- the value was laundered through a trivial or dimensionally-broken math expression. Fix: obtain the value from the proper tool (re-call the selection tool with a justified relaxed threshold; call the real scientific tool). Do NOT instruct the worker to use another math expression, and do NOT drop the quantity.

## 12 Judge scenario suite and adversarial tuning

The judge rules above were tuned and evaluated against a standalone scenario harness rather than by inspection. Each scenario registers a realistic set of upstream artifacts and a target artifact, then runs the same parameter-by-parameter verification used in production. Scenarios are grounded in artifacts from real run checkpoints, so their per-parameter rationales carry the rich, study-situated content that real runs produce; an earlier fully synthetic suite was discarded after its unrealistically thin rationales produced spurious failures. Each clean scenario mirrors a real registered artifact and should pass; each flawed scenario is a realistic base with exactly one injected defect and should fail or warn. The defect catalog covers: a wrong reference direction (e.g. a mid-sweep calculation designated as the reference); a source of the wrong physical kind; an implausible value magnitude (justified variants should warn, unjustified variants should fail); a deliberately generic rationale (exactly one such case per applicable tool, so that the other defects are tested on the defect rather than on rationale poverty); a mismatch between characterization and production conditions; a wrong file role (e.g. a clean slab supplied where an adsorbed system is required); non-ensemble inputs to the BEEF analysis; and value laundering through a dotted reference. Scenarios are grounded mostly in the CO/Pt(111) system, with a second chemistry included so that the rules are not tuned to a single system's idioms. Every scenario carries a hypothesis about its expected verdict that is used only to triage disagreements, never to shape the rules, and false passes and false failures are always measured together. A proposed rule change was accepted only if, on re-running the harness, all clean scenarios still passed and all defect scenarios still failed. The final suite contains 145 scenarios; the evaluation across five Anthropic judge models, with verdicts taken as the majority over three judge calls and a warning counted as catching a should-fail scenario, is reported in the main text (judge confusion matrices).



## 13 Full provenance graph of the Sol27LC run



**Fig. 2:** The complete provenance graph accumulated by the BCC-Li run at its final step. Every node is a registered tool call and every edge follows a recorded reference, so the graph assembles automatically as the run proceeds. The two convergence branches (each ending in a selection call and a structured report) and the production chain ending in the final lattice-constant report are visible. Even this deliberately simple benchmark accumulates a graph at the edge of practical human review, and the CO/Pt(111) study’s graph, spanning hundreds of registered calls, is well beyond it; the report judge walks these graphs mechanically at full depth.

## References

- [1] Campbell, Q., Cox, S., Medina, J., Watterson, B., White, A.D.: Mdcrow: Automating molecular dynamics workflows with large language models. arXiv preprint arXiv:2502.09565 (2025)
- [2] Pham, T.D., Tanikanti, A., Keçeli, M.: Chemgraph: An agentic framework for computational chemistry workflows. arXiv preprint arXiv:2506.06363 (2025)